

1 Explainability in Reinforcement Learning

2 Maxime Alaarabiou

3 Saturday 15th August, 2026

Chapter 1

Clustering Agent Representations

As discussed in the previous chapter, current explainability methods for Reinforcement Learning (RL) still lack intrinsic approaches capable of explaining the internal mechanisms of deep neural policies. To the best of our knowledge, only [Zahavy et al., 2017] have proposed exploiting the Latent Space (LS) learned by neural policies in order to identify the concepts used by the agent. Its main idea is to analyze the Latent Representation (LR) produced by the Neural Network (NN) and organize them into clusters that correspond to high-level concepts.

While the core idea of this paper is closely aligned with our vision of explainability in RL, this approach suffers from two important limitations. First, the clusters are manually constructed. Such a procedure heavily depends on the user’s expertise and neither guarantees reproducibility or scales. Second, no evaluation protocol is proposed to assess the quality of the obtained clusters. Their relevance is essentially determined through visual inspection.

The Clustering Agent Representation (CAR) method proposed in this chapter addresses these two limitations. It automates concept discovery by relying on unsupervised clustering algorithms to identify the Concept Structure (CS) of the LS. It introduces an objective evaluation metric based on the stability of the extracted representations across independently trained policy.

This chapter is organized as follows:

- Section 1.1 lays out the theoretical assumptions underlying the proposed CAR method, which we term the *Representation Convergence Hypothesis* (Hyp. 1).
- Section 1.2 then details the CAR method itself, covering both the extraction of LR using surrogate policies and the automatic detection of CS through clustering.
- Section 1.3 introduces the metrics we propose for evaluation.
- Section 1.4 reports the results obtained by applying CAR and these metrics across several Atari environments.

- Section 1.5 discusses these results, situates CAR within the broader landscape of mechanistic interpretability, and outlines directions for future work.

1.1 Representation Convergence Hypothesis

The CAR method relies on a assumption, referred to as the *Representation Convergence Hypothesis* (Hyp. 1). This hypothesis follows the two assumptions established previously throughout this manuscript.

The *Mechanistic Hypothesis* (Hyp. ??) rejects the notion that deep NN constitute irreducible black boxes. Adopting a mechanistic perspective, we consider that every process results from a sequence of causal computations. Although these computations may be complex, they remain explainable through the analysis of the internal mechanisms.

The *Concept Hypothesis* (Hyp. ??) states that NN achieve generalization by learning concepts. During optimization, the network progressively organizes its LS into representations that capture the information required to solve the task.

Following these hypotheses, it becomes possible to explain the NN’s process by understanding the organization of its LS. This family of approaches is commonly referred to as mechanistic interpretability.

In mechanistic interpretability, universality refers to the consistency with which a concept is recovered across different models or independent extraction processes [Templeton et al., 2024]. We develop the *Representation Convergence Hypothesis* (Hyp. 1) provides a theoretical explanation for this phenomenon. Under this assumption, the universality of a concept is not merely an empirical observation but the expected consequence of a common underlying CS shared by independently learned representations.

Hypothesis 1: Representation Convergence Hypothesis

For a given task, NN tend to converge toward a privileged CS because it provides an efficient representation for solving the task.

It is important to clarify that this hypothesis does not imply convergence at the parameter level. Independently trained models are expected to converge toward different sets of weights. The proposed hypothesis concerns a abstract level of representation. We assume that the conceptual organization induced by their LR converges toward an equivalent CS, even though the numerical implementation of the networks differ.

The remainder of this chapter builds upon this hypothesis. This principle forms the theoretical foundation of the *Clustering Universality* metric, introduced in Subsection 1.3.3.

1.2 Method

Following the work of [Zahavy et al., 2017], we analyze the LS through clustering techniques, placing our approach within the broader field of Deep Clustering. This design choice is motivated by our rejection of the *Superposition Hypothe-*

sis (Hyp. ??) and, more generally, of any view that assumes a linearly organized CS, as discussed in Section ??.

The CAR methodology is composed of two main stages. The first stage trains multiple independent NN, called surrogate policies, which implement the same decision policy. The second stage clusters their LR in order to identify the underlying conceptual organization.

This section is organized as follows. In the first subsection, we justify the use of multiple surrogate policies in order to study the stability of representations. In the second subsection, we detail how these surrogate policies are trained so as to best match the original policy.

1.2.1 Multiple surrogate policies

The CAR method was developed to test the *Representation Convergence Hypothesis* (Hyp. 1). This requires training several NN independently and comparing their LR. Consistent similarities in the resulting CS would provide empirical support for this hypothesis.

In supervised learning, this protocol is straightforward since several models can be independently trained on the same target function. RL introduces an additional difficulty. Because of stochastic exploration, independently trained agents are not guaranteed to converge toward the same policy. Even if they achieve comparable performance, they may solve the task using different strategies. This makes it difficult to determine whether differences in the LS result from different representations of the same strategy or from differences in the learned strategies themselves.


To overcome this limitation, we consider an already trained initial policy. This policy is then imitated using imitation learning [Ross et al., 2011]. In imitation learning, several surrogate NN are independently trained in a supervised manner to reproduce the behavior of the initial policy. This makes it possible to investigate whether equivalent policies also exhibit equivalent conceptual organizations.

It is important to emphasize that the use of surrogate policies does not make CAR a surrogate explainability method, as defined in Subsection ?. The surrogate policies remain deep NN. Their role is not to simplify the decision process but to provide multiple independent realizations of the same policy function, enabling the study of representation convergence.

1.2.2 Training surrogate policies

Throughout the remainder of this manuscript, the initial policy we aim to explain is denoted as π . This initial policy is obtained without imposing any constraints on its origin. It may result from classical or deep reinforcement learning [Sutton and Barto, 2018], alpha-beta algorithm [Knuth and Moore, 1975], or even from demonstrations provided by humans [Ross and Bagnell, 2010]. The crucial requirement is the availability of a dataset containing a sufficient number of projections from the observation space $\mathcal{O}$ to the action space $\mathcal{A}$ induced by π .

The replicated policies are referred to as surrogate policies, denoted by π' . For any observation O , the distribution over the action space produced by π' should approximate that of π as closely as possible.

Figure 1.1: Example of a surrogate policy architecture π' .

The surrogate policy π' is implemented as a NN parameterized by Θ . This network is composed of two components: an encoder and a decoder, respectively highlighted by the blue and black dashed rectangles in Figure 1.1. The encoder f_{enc} maps the observation O to the embedding. The decoder, which reconstructs the action distribution parameters from the embedding.

To train π' to accurately replicate the action distribution, we define a global loss function $\mathcal{L}$ optimized via gradient descent as introduced in Section ???. This loss, defined in Equation (1.1), is computed as the MSE between the parameters of the action distributions, obtained by applying the function p , which maps a distribution to its parameterization. In practice, this function does not need to be explicitly defined, as neural policies directly output the parameters of the action distribution. Once training is complete, this loss ensures that π' closely approximates the behavior of the original policy π . Figure 1.2 summarizes this overall procedure.

$$\min_{\Theta} \sum_{O \in \mathcal{O}} \mathcal{L}(O) = \min_{\Theta} \sum_{O \in \mathcal{O}} \text{MSE}(p(\pi(O)), p(\pi'_{\Theta}(O))) \quad (1.1)$$

We denote by $\mathcal{E}$ the set of embeddings obtained by applying f_{enc} to a subset of observations $\mathbb{O} \subseteq \mathcal{O}$.

On the set $\mathcal{E}$, we use a clustering function f_{clust} to partition the point cloud of embedded observations into distinct clusters. Each observation is assigned a cluster label based on its proximity in the LS, resulting in the clustering $\mathcal{C}$. Algorithm 1 summarizes the complete CAR procedure.

Figure 1.3 provides an illustration of the CAR method. The point cloud




Figure 1.2: Illustration of the Clustering Agent Representation method.

shows a two-dimensional projection of the LS of a surrogate policy trained on Space Invaders. The two-dimensional projection is used only for visualization, while the clusters are computed in the original high-dimensional LS. Colors indicate the resulting clusters and their corresponding observations. Each cluster corresponds to a distinct, human-interpretable stage of the game. The resulting organization suggests a correspondence between the position of observations in the LS and the semantic properties of the corresponding game states. To objectively assess this correspondence we introduce a set of evaluation metrics detailed in the following section.

1.3 Evaluation Metrics

A key contribution of this work is the development of three evaluation metrics, designed to assess the effectiveness of CAR:

- **Behavior Reconstruction** : it measures how well a surrogate policy faithfully reproduces the behavior of the original policy by comparing their cumulative rewards.
- **Clustering Universality** : it evaluates the stability of clusterings by measuring the similarity of clusters obtained from independently trained surrogate policies.
- **Abstraction Score** : it quantifies the extent to which LS clusters capture abstractions rather than directly reflecting the structure of observations or actions.

The remainder of this section details the motivation for each metric and describes how they are calculated. In the first subsection, we develop the *Behavior Reconstruction*. We then introduce the Adjusted Rand Index (ARI)

Algorithm 1: Clustering Agent Representation (CAR)

Input:
 Original policy π
 Observation dataset $\mathbb{O}$
 Number of surrogate policies n
 Surrogate policy architecture π' (encoder f_{enc} , decoder)
 Learning rate η
 Clustering function f_{clust}

Output:
 Trained surrogate policies $\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}$
 Clusterings $\mathcal{C}_1, \dots, \mathcal{C}_n$

Behavioral cloning dataset built from the target policy π
1 $\mathcal{D} \leftarrow \{(O, p(\pi(O))) \mid O \in \mathbb{O}\};$

MSE loss between action-distribution parameters, as defined in Equation (??)
2 $\mathcal{L} \leftarrow \text{Mean Squared Error (MSE)};$

Independently train n surrogate policies by behavioral cloning
3 for $i = 1, \dots, n$ **do**
4 Initialize Θ_i randomly;
 Behavioral cloning of π via SGD, see Algorithm ??
5 $\Theta_i \leftarrow \text{SGD}(\mathcal{D}, \pi', \mathcal{L}, \eta, \Theta_i);$

Cluster the latent representations of each surrogate policy
6 for $i = 1, \dots, n$ **do**
 Embeddings produced by the encoder of π'_{Θ_i}
7 $\mathcal{E}_i \leftarrow \{f_{enc_i}(O) \mid O \in \mathbb{O}\};$
 Partition of the embedded observations into clusters
8 $\mathcal{C}_i \leftarrow f_{clust}(\mathcal{E}_i);$

9 Return $(\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}), (\mathcal{C}_1, \dots, \mathcal{C}_n)$

164 measure, which serves as the foundation for the two other metrics introduced
 165 in the subsequent subsections, namely *Clustering Universality* and *Abstraction*
 166 *Score*.

167 The following notations will be adopted for the rest of the manuscript: we
 168 consider n surrogate policies, denoted as $\pi'_1, \dots, \pi'_n$, trained as described in
 169 Section 1.2. Each surrogate policy is associated with a projection function
 170 $f_{enc_1}, \dots, f_{enc_n}$, which is applied to the set of observations $\mathbb{O}$, yielding embed-
 171 dings $\mathcal{E}_1, \dots, \mathcal{E}_n$. A clustering function is then applied to each embedding set,
 172 producing partitions $\mathcal{C}_1, \dots, \mathcal{C}_n$.

173 1.3.1 Behavior Reconstruction

174 We argue that explaining a function b that closely approximates another function
 175 a is essentially the same as explaining a . If b accurately captures the key decision
 176 patterns of a , understanding b provides meaningful insights into how a works
 177 [Guidotti et al., 2018]. In our setting, this means that generating explanations




Figure 1.3: Two-dimensional projection of the LS embeddings of a surrogate policy trained on Space Invaders, together with the CAR clusters and their corresponding observations.

for π' that faithfully reproduce π amounts to explaining π directly. To verify that π' is indeed a faithful copy, we introduce the *Behavior Reconstruction* metric.

This metric is computed by comparing the cumulative rewards obtained by the original policy π and the surrogate policies π' over m evaluation episodes. The behavior reconstruction metric is then defined as:

$$\text{Behavior Reconstruction} = \frac{1}{n \sum_{i=1}^m g_{\pi}^i} \sum_{j=1}^n \sum_{k=1}^m g_{\pi'_j}^k \quad (1.2)$$

In imitation learning, it is unlikely to expect the surrogate policy π' to outperform the initial policy π it aims to approximate through supervised learning [Ross and Bagnell, 2010]. Consequently this metric ranges in $[0, 1]$. Values approaching 1 indicate minimal loss in reproducing π , while values close to 0 indicate poor reproduction.

We deliberately adopt a reward-based definition rather than alternatives based on actions or action distributions (e.g., training loss, accuracy, KL divergence [Shlens, 2014], or optimal transport [Peyré and Cuturi, 2020]). Since surrogate

191 policies are trained in a supervised manner, they cannot develop strategies
 192 beyond those demonstrated by the initial policy. This means that if the rewards
 193 achieved are comparable, the surrogate policy is effectively copying the original
 194 policy’s behavior.

195 1.3.2 Adjusted Rand Index

196 Within CAR, concepts are assumed to correspond to clusters in the LS. Combined
 197 with the *Representation Convergence Hypothesis* (Hyp. 1), this leads to a direct
 198 experimental prediction: if independently trained surrogate policies converge
 199 toward the same conceptual organization, then clustering their LR should produce
 200 similar partitions.

201 When a particular observation is processed by the network, it is expected to
 202 activate the same concept regardless of the specific neural implementation of the
 203 policy. If independently trained policies converge toward the same conceptual
 204 organization, a given observation should activate the same concept across all
 205 models. As a consequence, observations grouped together within a cluster for
 206 one policy should also be grouped together for another independently trained
 207 policy.

208 In the clustering literature, the most used measures for comparing two
 209 partitions is the ARI [Hubert and Arabie, 1985]. This measure quantifies the
 210 agreement between cluster assignments independently of the spatial organiza-
 211 tion of the clusters. Both metrics introduced later in this chapter, *Clustering*
 212 *Universality* and *Abstraction Score*, rely on the ARI. Since the ARI is a complex
 213 measure, we review its definition and discuss its main properties.

214 The ARI is derived from the Rand Index (RI) [Rand, 1971], which measures
 215 the agreement between two partitions by considering every possible pair of point
 216 cloud. For each pair, four situations are possible: the point may belong to the
 217 same cluster in both partitions, to different clusters in both partitions, or they
 218 may disagree by being grouped together in only one of the two partitions. The
 219 RI measures the proportion of pairs for which the two partitions agree:

$$\text{RI} = \frac{\text{number of agreeing pairs}}{\text{total number of pairs}}.$$

220 Although the RI is intuitive, it suffers from an important limitation. In most
 221 clustering problems, the majority of pairs of points belong to different clusters.
 222 Consequently, even two completely independent partitions will agree on a large
 223 number of pairs, simply because both assign them to different clusters. This
 224 effect becomes more pronounced as the number of points or clusters increases.

225 The ARI addresses this issue by accounting for the agreement that would
 226 be expected purely by chance. It compares the observed agreement RI to
 227 the expected agreement under random partitions $\mathbb{E}[\text{RI}]$. The score is then
 228 normalized by the maximum agreement achievable beyond this random baseline,
 229 $\max(\text{RI}) - \mathbb{E}[\text{RI}]$. As a result, an ARI of 0 corresponds to the level of agreement
 230 expected from random clusterings, while an ARI of 1 indicates perfect agreement
 231 between the two partitions:

$$\text{ARI} = \frac{\text{RI} - \mathbb{E}[\text{RI}]}{\max(\text{RI}) - \mathbb{E}[\text{RI}]}$$

In practice, the ARI is computed directly from the cluster assignments of the two partitions. Let clustering A be composed of partitions $\{A_i\}$ and clustering B of partitions $\{B_j\}$. Let $a_i = |A_i|$ denote the number of samples in partition i of clustering A , and $b_j = |B_j|$ the number of samples in partition j of clustering B . Furthermore, let $n_{ij} = |A_i \cap B_j|$ be the number of samples shared by partition i of clustering A and partition j of clustering B , and let n be the total number of samples. With these definitions, the ARI can be computed as:

$$\text{ARI} = \frac{\sum_{i,j} \binom{n_{ij}}{2} - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}{\frac{1}{2} \left[\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2} \right] - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}$$

Because it is bounded by 1, a value such as 0.5 can be misleadingly interpreted as “halfway to a perfect match.” This interpretation is incorrect because the ARI is normalized against the agreement expected from two random partitions. As a result, values close to 1 require the two clusterings to be nearly identical, while a value of 0.5 already indicates substantial agreement [Amezquita et al., 2020].

1.3.3 Clustering Universality

Having introduced the ARI, we can now define the *Clustering Universality* metric. While the ARI quantifies the similarity between two conceptual partitions, the objective of CAR is to assess the consistency of conceptual organizations across an entire set of independently trained surrogate policies.

To this end, the same set of observations is projected into the LS of each surrogate policy and clustered. Under the *Representation Convergence Hypothesis* (Hyp. 1), these independently obtained partitions are expected to be similar. *Clustering Universality* quantifies this by computing the mean pairwise ARI:

$$\text{Clustering Universality} = \frac{1}{\binom{n}{2}} \sum_{i=1}^n \sum_{j=0}^{i-1} \text{ARI}(C_i, C_j). \quad (1.3)$$

This metric takes values in $[-1, 1]$. A value close to 1 indicates that surrogate policies produce highly similar partitions, suggesting that the clusters reflect intrinsic properties of the policy. Conversely, a value near 0 implies that the similarity between partitions is no better than that of two randomly generated clusterings, meaning that the clustering fails to capture any meaningful structure in the policy. A negative value, on the other hand, signals that the agreement between partitions is worse than what would be expected by chance, indicating a systematic disagreement between the clusterings.

Figure 1.4 illustrates *Clustering Universality* in the ideal case. It shows three LS obtained by projecting the same seven observations, labeled 1 to 7, through three independently trained surrogate policies. Although the embeddings occupy different positions, the clusters group exactly the same elements in all three cases. The partitioning is therefore identical, resulting in an ARI of 1.

1.3.4 Abstraction Score

If *Clustering Universality* is high, it becomes important to determine what underlies this stability. Two possibilities arise:




Figure 1.4: Illustration of *Clustering Universality* in the ideal case where projections of the same observations are grouped within the same clusters across all surrogate policies.

- (i) the LS merely reproduces the structure of the observation space or action distribution, offering no abstraction, as defined in Definition ??;
- (ii) the LS captures similarities in decision-making processes, yielding behaviorally meaningful clusters.

To distinguish among these cases, we introduce the *Abstraction Score*. This metric contextualizes latent clusterings by examining how their patterns align with reference structures derived from the raw observation and action data. We consider : $\mathcal{C}_{obs}$, the clustering in the observation space $\mathbb{O}$, and $\mathcal{C}_{act}$, the clustering in the action-space distribution parameters. The clusterings $\mathcal{C}_{obs}$ and $\mathcal{C}_{act}$ are performed using the same algorithm, distance metric, and hyperparameters as those used for $\mathcal{C}_i$ in the LS, ensuring that any biases introduced by the clustering procedure itself are minimized.

For each latent clustering $\mathcal{C}_i$, we define its similarity to these baselines using the ARI:

$$A_i^{obs} = \text{ARI}(\mathcal{C}_i, \mathcal{C}_{obs})$$

$$A_i^{act} = \text{ARI}(\mathcal{C}_i, \mathcal{C}_{act})$$

The *Abstraction Score* is subsequently calculated as one minus the average of the maximum absolute similarity values across these two reference baselines:

$$\text{Abstraction Score} = 1 - \frac{1}{n} \sum_{i=1}^n \max(|A_i^{obs}|, |A_i^{act}|) \quad (1.4)$$

The *Abstraction Score* ranges from $[0, 1]$. A score around 0 means that the latent clusters are strongly aligned with either the observation or the action space, reflecting minimal abstraction, cases (i). A score close to 1 indicates that the latent clusters are independent of both the observation and action spaces, capturing abstractions that meaningfully reflect the agent’s behavior, case (ii).

The following section presents the results of applying these metrics to specific examples.

Environment	<i>Behavior Reconstruction</i>	<i>Abstraction Score</i>	<i>Clustering Universality</i>
Breakout	0.83 $\pm$ 0.02	0.99 $\pm$ 0.00	0.58 $\pm$ 0.09
Mario Bros	0.98 $\pm$ 0.02	0.99 $\pm$ 0.00	0.55 $\pm$ 0.10
Pong	0.95 $\pm$ 0.00	0.99 $\pm$ 0.00	0.48 $\pm$ 0.04
Space Invaders	0.97 $\pm$ 0.07	0.99 $\pm$ 0.00	0.51 $\pm$ 0.08
Surround	0.99 $\pm$ 0.00	0.99 $\pm$ 0.00	0.58 $\pm$ 0.07

Table 1.1: Evaluation of *Behavior Reconstruction*, *Abstraction Score* and *Clustering Universality* across the five Atari environments.

1.4 Results

This section quantitatively evaluates this behavior across several Atari benchmark environments [Bellemare et al., 2013], which are widely used in RL research [Mnih et al., 2013, Schulman et al., 2017, Zahavy et al., 2017]. The environments considered in this study are:

- **Breakout:** the agent controls a paddle to bounce a ball and break as many bricks as possible, the ball speeding up after each bounce.
- **Mario Bros:** the agent must defeat enemies, avoid fireballs, and collect mushrooms [AtariAge, 2024].
- **Pong:** a two-player game in which the agent tries to prevent the ball from passing its side while attempting to score against the opponent.
- **Space Invaders:** the agent controls a cannon and must destroy waves of invaders before they reach the ground.
- **Surround:** a two-player grid-based game in which each move blocks the traversed cell, the goal being to trap the opponent.

In all these environments, the agent receives visual observations and operates in discrete action spaces. The initial policies π are trained using Proximal Policy Optimization (PPO) [Schulman et al., 2017] with a deep actor-critic architecture. The network consists of four convolutional layers with configurations (16, 4, 2), (32, 4, 2), (64, 4, 2), (128, 4, 2), where each tuple represents (number of channels, kernel size, stride), followed by three fully connected layers with 512, 256, and 128 units, respectively. To ensure that these initial policies effectively learn to maximize rewards, we follow the recommendations of [Machado et al., 2017] and compare the performance of our initial policies against the benchmark scores reported in the same study to verify that they meet the required performance standards.

For each initial policy, three surrogate policies π' sharing the same architecture are trained as described in Section 1.2. These surrogate policies are trained on approximately one million observation-action pairs each.

We first evaluate the ability of the surrogate policies to replicate the behavior of the initial policies (*Behavior Reconstruction*, Table 1.1). Across all




Figure 1.5: Evolution of *Clustering Universality* with respect to the number of clusters, for each of the five Atari environments.

environments, the surrogate policies achieve an average *Behavior Reconstruction* of 0.94 ± 0.02 , indicating that they reliably reproduce the decisions of the initial policies over 100 evaluation episodes.

For the remainder of the analysis, a random subset of 50,000 observations is drawn from the training dataset and embedded into the LS of each surrogate policy, using the representation extracted directly after the last convolutional layer. This projection results in a point cloud $\mathcal{E}$. For clustering, we opted for the simplest approach and applied **K-means clustering** [MacQueen, 1967] to the point cloud, using the Euclidean distance between the projected points to form the clusters.

Across all environments, policies, and numbers of clusters, the average *Abstraction Score* is 0.99 ± 0.00 . This indicates that the clusters formed in the LS are not mere reflections of the raw pixel inputs or of the action distributions.

Figure 1.5 reports the evolution of *Clustering Universality* with respect to the number of clusters. From three clusters onward, all curves stabilize around 0.5 in every environment. Since *Clustering Universality* is computed from the ARI, this result indicates a strong similarity in the clustering of the LS across surrogate policies within each environment [Amezquita et al., 2020], regardless of the number of clusters chosen. This stability indicates that there is no single optimal number of clusters for explaining an agent’s policy. This raises the question of how many clusters should be selected to perform an analysis of LS. Our analysis suggests that this depends on the desired level of granularity.

With fewer clusters, CAR highlights broad strategic patterns. For instance, in the Mario Bros environment, the two main groups correspond to whether the POW block is present (highlighted in red in Figure 1.6A) or not. This is particularly interesting, as the use of the POW block is a key determinant of strategy in this environment [AtariAge, 2024].

With a larger number of clusters, the analysis captures more fine-grained tactical behaviors, such as ‘Attacking the left.’ and ‘Navigating through the level.’ (see Figure 1.6B).




Figure 1.6: Effect of the number of clusters on the groupings identified by CAR in the Mario Bros environment.

Both perspectives are valid, depending on whether the goal is to study strategic or tactical behaviors. Similar interpretable structures can be observed across all evaluated environments.

1.5 Conclusion

The three hypotheses investigated in this manuscript, *Mechanistic Hypothesis* (Hyp. ??), *Concept Hypothesis* (Hyp. ??), and *Representation Convergence Hypothesis* (Hyp. 1), formed the conceptual basis for the development of CAR. The experiments conducted with CAR provide empirical support for these hypotheses, with the high *Abstraction Score* and stable *Clustering Universality* obtained across all environments. These findings suggest that this hypothesis-driven approach provides a promising direction for understanding the structure of learned representations in RL agents.

Beyond its role as an experimental framework, CAR provides a general and automated method for analysing agent representations. Both the analysis procedure and the metrics are fully automated, requiring no manual intervention or environment-specific adjustments. This makes CAR a general method that can be applied to a wide range of RL settings.

However, several implementation choices were made for have not yet been evaluated. These include the NN architectures, the NN layer selected for analysis,

the distance used to compare embeddings, and the clustering algorithm. Their impact should therefore be studied jointly to determine how changes in the analysis procedure affect the discovered CS.

A more fundamental limitation concerns the semantic interpretation of the clusters. While CAR provides a quantitative description of the structure present in the agent’s representations, the semantics assigned to individual clusters currently rely on manual interpretation. The labels associated with the clusters in Figure 1.6 represent our interpretation of the discovered structure. This process is time-consuming and subjective. Developing methods to automatically extract the semantics of each cluster is therefore an important direction for future work.

More broadly, when considering CAR within the field of explainability, it is important to emphasize that substantial work remains to be done. The results presented in this chapter show that CAR can provide information about the concepts represented in the agent’s LS, which can therefore be considered an explanation in the sense defined in Section ???. How these explanations can be used to support validation, diagnosis, or certification remains to be established ??.

These limitations directly motivate several avenues for future research. Studying the impact of the selected LS layer, distance measures, and clustering algorithms would clarify how much the discovered structure depends on these methodological choices. Automating cluster interpretation and developing human-centered evaluation protocols would strengthen semantic validity of the resulting explanations. Finally, connecting CAR explanations more directly to the agent’s decisions and failures would help determine how they can support the broader objectives of explainability.

Acronyms

- ARI** Adjusted Rand Index 5, 8–10, 12
- CAR** Clustering Agent Representation 1–5, 7–9, 12–14
- CS** Concept Structure 1–3, 14
- LR** Latent Representation 1–3, 8
- LS** Latent Space 1–5, 7–10, 12, 14
- MSE** Mean Squared Error 4, 6
- NN** Neural Network 1–4, 13
- PPO** Proximal Policy Optimization 11
- RI** Rand Index 8
- RL** Reinforcement Learning 1, 3, 11, 13

Bibliography

- [Amezquita et al., 2020] Amezquita, R. A., Lun, A. T., Becht, E., Carey, V. J., Carpp, L. N., Geistlinger, L., Marini, F., Rue-Albrecht, K., Risso, D., Soneson, C., et al. (2020). Orchestrating single-cell analysis with Bioconductor, volume 17. Nature Publishing Group US New York. **Book**.
- [AtariAge, 2024] AtariAge (2024). Mario bros. atari 2600. **Web**.
- [Bellemare et al., 2013] Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. Journal of Artificial Intelligence Research, 47:253–279. **Article**.
- [Guidotti et al., 2018] Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Pedreschi, D., and Giannotti, F. (2018). A survey of methods for explaining black box models. **Article**.
- [Hubert and Arabie, 1985] Hubert, L. and Arabie, P. (1985). Comparing partitions. Journal of classification, 2(1):193–218. **Article**.
- [Knuth and Moore, 1975] Knuth, D. E. and Moore, R. W. (1975). An analysis of alpha-beta pruning. Artificial intelligence, 6(4):293–326. **Article**.
- [Machado et al., 2017] Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., and Bowling, M. (2017). Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. **Article. Web**.
- [MacQueen, 1967] MacQueen, J. (1967). Multivariate observations. In Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability, volume 1, pages 281–297. **Article**.
- [Mnih et al., 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.
- [Peyré and Cuturi, 2020] Peyré, G. and Cuturi, M. (2020). Computational optimal transport. **Article**.
- [Rand, 1971] Rand, W. M. (1971). Objective criteria for the evaluation of clustering methods. Journal of the American Statistical association, 66(336):846–850.

- [Ross and Bagnell, 2010] Ross, S. and Bagnell, D. (2010). Efficient reductions for imitation learning. In Teh, Y. W. and Titterton, M., editors, Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics, volume 9 of Proceedings of Machine Learning Research, pages 661–668, Chia Laguna Resort, Sardinia, Italy. PMLR. **Article**.
- [Ross et al., 2011] Ross, S., Gordon, G., and Bagnell, D. (2011). A reduction of imitation learning and structured prediction to no-regret online learning. In Gordon, G., Dunson, D., and Dudík, M., editors, Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, volume 15 of Proceedings of Machine Learning Research, pages 627–635, Fort Lauderdale, FL, USA. PMLR. **Article**.
- [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms.
- [Shlens, 2014] Shlens, J. (2014). Notes on kullback-leibler divergence and likelihood. **Article**.
- [Sutton and Barto, 2018] Sutton, R. S. and Barto, A. G. (2018). Reinforcement Learning: An Introduction. MIT Press. **Book**.
- [Templeton et al., 2024] Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., et al. (2024). Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet. Transformer Circuits Thread. **Web**.
- [Zahavy et al., 2017] Zahavy, T., Zrihem, N. B., and Mannor, S. (2017). Graying the black box: Understanding dqns. **Article**.