

1 Explainability in Reinforcement Learning

2 Maxime Alaarabiou

3 Friday 11th September, 2026

Chapter 1

Experiments

This chapter presents the experimental study of the Clustering Agent Representation (CAR) method introduced in the previous chapter. It focuses on the experimental validation of the method and its ability to consistently identify meaningful conceptual structures. The remainder of this chapter is organized as follows.

- Section 1.1 describes the software implementation of the CAR pipeline.
- Section 1.2 details the experimental configuration, including the environments, the reinforcement learning setup, and the training of the surrogate policies.
- Section 1.3 presents and analyses the results obtained by applying CAR and its metrics.
- Section ?? summarizes the findings and discusses their implications.

1.1 Software Implementation

Deep Learning (DL) and Reinforcement Learning (RL) are inherently empirical fields, in which the effectiveness of a method is assessed through experimentation. For this reason, particular attention is given to the implementation of CAR. It relies on established libraries that provide efficient and widely used implementations of its main components. The implementation is made publicly available as open-source to support the objective and reproducible replication of the reported results.¹

CAR is organized as a pipeline of five stages, shown in Figure 1.1. Each stage is an independent executable, and stages communicate only through artifacts written to disk. This decoupling allows any stage to be rerun or inspected independently. The parameters of each experiment are gathered in a single configuration object, described in Section 1.2. Every result can therefore be traced back to the settings that produced it.

Table 1.1 summarizes the main third-party libraries used at each stage of the pipeline.

¹https://github.com/Maxalaar/Clustering_Agent_Latent_Space




Figure 1.1: The CAR pipeline.

1.1.1 Initial policy training

The first stage trains the initial policy using RLlib [Moritz et al., 2018, Liang et al., 2018]. RLlib provides a highly scalable implementation of RL algorithms and supports a wide range of state-of-the-art methods. It supports distributed training, allowing the computational workload to be spread across multiple processes and machines.

The use of RLlib allows CAR to be applied to any environment implemented using the Gymnasium interface [Towers et al., 2024]. This flexibility allows CAR to be applied to a wide range of environments. By default, the RL algorithm used to train the initial policy is Proximal Policy Optimization (PPO) [Schulman et al., 2017]. Other algorithms supported by RLlib, such as Asynchronous Proximal Policy Optimization (APPO) [Espeholt et al., 2018], Deep Q-Networks (DQN) [Mnih et al., 2013], and DreamerV3, can also be selected through the same configuration.

The policy and critic networks are implemented using PyTorch [Paszke et al., 2019]. PyTorch provides efficient GPU computation and supports a wide range of neural network architectures. This provides a standardized implementation that is compatible with the broader DL ecosystem.

Pipeline Stage	Library	Description
Initial policy training	RLlib [Moritz et al., 2018, Liang et al., 2018]	Distributed RL library used to train the initial policy.
	PyTorch [Paszke et al., 2019]	Deep learning framework implementing the policy and critic networks.
	Gymnasium [Towers et al., 2024]	Common interface used to plug in any RL environment.
Trajectory dataset generation	RLlib [Moritz et al., 2018, Liang et al., 2018]	Distributed environment runners used to collect trajectories in parallel.
	Apache Arrow [Apache Arrow, 2016]	Columnar file format (Parquet) used to store the trajectory datasets.
Surrogate policy training	PyTorch Lightning [Falcon, 2019]	Structured training loop used for the behavior cloning of the surrogate policies.
Surrogate policy evaluation	RLlib [Moritz et al., 2018, Liang et al., 2018]	Wraps each surrogate policy as an RL policy to evaluate it in the environment.
Latent space analysis	cuML [RAPIDS, 2023]	GPU-accelerated clustering algorithms used to analyze the latent representations.

Table 1.1: Main software libraries used at each stage of the CAR pipeline.

1.1.2 Trajectory dataset generation

The second stage builds the dataset used to train the surrogate policies. The distributed execution capabilities of RLlib are used to collect trajectories in parallel. Multiple environment runners execute the initial policy simultaneously, distributing the computational workload across the available resources.

The initial policy is run in the environment, and the interactions between the agent and the environment are recorded. For each step, the dataset stores the observation together with the parameters of the action distribution produced by the policy. Trajectories are stored as Parquet files [Apache Arrow, 2016]. Parquet is a compressed columnar format widely used for large-scale data processing. Its columnar structure allows data to be read efficiently and supports concurrent access from multiple processes. This is particularly useful here because the datasets contain on the order of one million observation–action pairs and cannot be loaded entirely into memory.

1.1.3 Surrogate policy training

The third stage trains the surrogate policies using behavior cloning. This provides a supervised learning approach for replicating the initial policy. The training procedure is implemented using PyTorch Lightning [Falcon, 2019]. PyTorch Lightning provides a structured interface for supervised training while remaining compatible with the PyTorch ecosystem. It handles common training operations such as data loading, optimization, checkpointing, and GPU execution, while allowing the underlying neural network implementation to remain fully controlled.

73 1.1.4 Surrogate policy evaluation

74 The fourth stage measures how faithfully each surrogate policy reproduces the
 75 behavior of the initial policy. Each surrogate is wrapped so that it can be used
 76 as an RLlib policy and is then evaluated in the original environment. Its mean
 77 cumulative reward over multiple episodes is compared with that of initial policy.
 78 Evaluation episodes are collected using parallel environment runners.

79 This stage provides the value of the *Behavior Reconstruction*. It therefore
 80 verifies whether the surrogate policies adequately reproduce the behavior of the
 81 initial policy before their latent representations are analyzed. If a surrogate does
 82 not reproduce the original behavior sufficiently well, its training can be repeated
 83 before proceeding to the next stage.

84 1.1.5 Latent space analysis

85 The fifth stage performs the main analysis of CAR. The latent representations
 86 produced by the surrogate policies form the point clouds that are analyzed using
 87 clustering algorithms.

88 The clustering step relies on the cuML library [RAPIDS, 2023], which pro-
 89 vides GPU-accelerated implementations of several clustering algorithms. This
 90 significantly reduces the computational cost of clustering and makes it possible
 91 to analyze point clouds containing tens of thousands of representations.

92 The resulting partitions are used to compute the *Clustering Universality* and
 93 the *Abstraction Score*. The outputs of this stage are the resulting metric values
 94 and the material required for the qualitative analysis presented in the following
 95 sections.

96 1.2 Experimental Configuration

97 This section describes the configuration of the experiments reported in this
 98 chapter. It covers the environments, the neural architecture of the policies, and
 99 the main parameters of each pipeline stage. In general, for each environment, a
 100 single initial policy is trained and then imitated by three surrogate policies.

101 1.2.1 Environments

102 All experiments use environments from the Atari benchmark [Bellemare et al., 2013],
 103 accessed through the Arcade Learning Environment. The use of Atari environ-
 104 ments offers several advantages. First, Atari is a widely used benchmark in
 105 the RL community, with many studies using these environments as standard
 106 evaluation settings. This makes the experimental results easier to compare
 107 with existing work. Second, these environments were not specifically designed
 108 for RL, which reduces the risk of introducing biases related to the design of
 109 the environments themselves. Finally, the benchmark provides a diverse set of
 110 well-known environments with varying levels of complexity. This diversity makes
 111 it possible to evaluate CAR across different types of tasks and levels of difficulty.

112 The configuration of the Atari environments follows common practice in RL
 113 research [Machado et al., 2017].

114 In each environment, the agent receives visual observations in the form of
 115 raw pixels. Each image is resized to 64×64 pixels and converted to grayscale




Figure 1.2: The six Atari environments used to evaluate CAR.

116 to reduce memory requirements. An exception is made for Pac-Man, where
 117 color information is important for determining the behavior of other agents in
 118 the environment. The original color information is therefore preserved for this
 119 environment.

120 To provide information about the dynamics of the environment, four con-
 121 secutive observations are stacked along the channel dimension. The resulting
 122 observation has a shape of $64 \times 64 \times 4$. This allows the agent to infer the motion
 123 of objects from successive observations rather than relying on a single frame.

124 Frame skipping is also used to reduce the duration of each episode. The agent
 125 selects a new action every five frames, while the selected action is repeated for
 126 the intervening frames. This provides sufficient temporal resolution for effective
 127 control while reducing the number of decisions required during an episode.

128 The actions correspond to the different buttons available on the Atari con-
 129 troller. The resulting action space is therefore discrete. Since Atari games are
 130 fully deterministic, sticky actions are introduced to add stochasticity to the
 131 trajectories. With a probability of 0.05, the environment repeats the action
 132 selected at the previous step instead of applying the new action. This prevents
 133 the agent from simply memorizing a fixed sequence of actions during training.

134 With the environment configuration established, the next section describes
 135 the objective associated with each environment. A visualization of the six
 136 environments is provided in Figure 1.2 to give an overview of the experimental
 137 settings.

138 **Breakout.** The agent moves a paddle at the bottom of the screen (Figure 1.2a)
 139 to bounce a ball into a wall of bricks, each broken brick scoring a point, and the
 140 ball speeding up as the wall is cleared.

141 **Mario Bros.** A single-screen platform game (Figure 1.2b). The player clears
 142 each phase by hitting enemies from below to flip them, then kicking them off
 143 the platforms before they recover. A central POW block flips every enemy at
 144 once but can only be used a few times per phase.

145 **Ms. Pac-Man.** The player moves through a maze (Figure 1.2c), eating pellets
 146 while avoiding four ghosts, with power pellets briefly making the ghosts edible.
 147 The maze structure and the many possible local situations, such as a nearby
 148 ghost, an intersection, or an active power pellet, produce a large number of
 149 distinguishable states.

150 **Pong.** A minimal two-player game (Figure 1.2d) in which the agent controls
 151 one paddle and must send the ball past the opponent.

152 **Space Invaders.** The agent moves a cannon along the bottom of the screen
 153 (Figure 1.2e), firing at descending rows of aliens while dodging their shots and
 154 using destructible shields, with a bonus saucer occasionally crossing the top.

155 **Surround.** A two-player game on a grid (Figure 1.2f). Each player leaves a
 156 solid trail behind a moving cursor, and the goal is to force the opponent to crash
 157 into a trail or a wall.

158 1.2.2 Neural architecture

159 The surrogate policies use a convolutional neural network to process the visual
 160 observations. The input consists of a stack of four consecutive 64×64 frames.
 161 The convolutional part consists of four layers with 16, 32, 64, and 128 channels,
 162 using 4×4 kernels with a stride of 2. Each convolutional layer is followed by
 163 a LeakyReLU activation. The resulting representation is then flattened and
 164 passed through a multilayer perceptron that produces the parameters of the
 165 action distribution.

166 The architecture of the multilayer perceptron varies across environments. For
 167 Pong and Space Invaders, it contains a single hidden layer with 128 units. For
 168 the other four environments, it contains three hidden layers with 512, 256, and
 169 128 units. Figure 1.3 illustrates both configurations. The size of the output layer
 170 matches the number of actions available in each environment: 4 for Breakout,
 171 18 for Mario Bros, 9 for Ms. Pac-Man, 6 for Pong, 6 for Space Invaders, and 5
 172 for Surround.

173 The latent space analyzed by CAR is not restricted to a specific layer or
 174 architecture. Any latent space produced by the network can be analyzed using
 175 the same methodology. In the experiments presented in this chapter, the default
 176 analysis considers the representation produced by the last convolutional layer.
 177 The layer-wise analysis additionally considers the representations produced by
 178 the hidden layers of the multilayer perceptron. The final output layer is excluded
 179 because it directly produces the action-distribution parameters.




Figure 1.3: Architecture of the surrogate policies.

1.2.3 Parameters

Having defined the environments and neural architectures, the remaining parameters of the CAR pipeline must now be specified. There are too many parameters to discuss them all in detail, as each stage of the pipeline introduces its own configuration choices. The main parameters shared across the experiments are summarized in Table 1.2.

Given the large number of possible configurations, it is not feasible to evaluate all possible combinations. The choices made in this work are therefore based primarily on standard practices from the literature and on practical experience. Whenever possible, we used the simplest suitable configuration. For the clustering algorithm and distance metric in particular, the simplest suitable choices are used to establish a baseline.

These choices do not restrict the CAR methodology itself. Most components can be replaced or configured differently depending on the application. The evaluation metrics provide an objective basis for comparing these alternatives and determining which configurations produce the most consistent and meaningful results.

1.3 Results

This section reports the results of applying CAR to the six environments. Sections 1.3.1 and 1.3.2 first verify that the two learning stages of the pipeline behave as expected. The first section verifies that the initial policies reach a competent level of play, while the second checks that the surrogate policies faithfully reproduce their behavior, using the *Behavior Reconstruction*. Section 1.3.3 then presents the analysis of the latent representations using the *Abstraction Score* and *Clustering Universality*.

1.3.1 Initial policy training

Figure 1.4a shows how the initial policies improve during PPO training. For each environment, the mean episode return is normalized so that 0% and 100% correspond to the lowest and highest mean return observed during that run. In

Environment configuration	
Frame size	64×64
Frames stacked per observation	4
Frame skip	5
Sticky-action probability	0.05
Initial policy training	
Algorithm	PPO
Optimizer	Adam
Learning rate	1.5×10^{-4}
GAE λ	0.95
Policy clipping ε	0.1
Value-function clipping	10
Entropy coefficient	0.01
KL coefficient	0.5
Epochs per update	10
Gradient clipping (global norm)	100
Surrogate policy training	
Surrogate policies per environment	3
Observation–action pairs per dataset	$\geq 10^6$
Batch size	2048
Optimizer	Adam
Learning rate	1×10^{-4}
Latent space analysis	
Embedding per surrogate policy	50 000
Clustering algorithm	k -means
Distance	Euclidean distance
Number of clusters K	2 to 30


Table 1.2: Main experimental parameters.

every environment the return increases and then reaches a plateau, at which point training is stopped. The number of iterations required to reach this plateau varies from a few hundred, for Pong, to several thousand, for Surround and Ms. Pac-Man. For Ms. Pac-Man the return decreases after its peak. Since the best checkpoint is retained rather than the last one, this does not affect the policy that is subsequently analyzed.


To confirm that the trained policies play competently, their scores are compared with the reference scores reported by [Machado et al., 2017]. Each initial policy reaches a score of the same order as these references, so it constitutes a reasonable target for imitation.

1.3.2 Surrogate policy training

Figure 1.4b shows, for each environment, the training action loss of the three surrogate policies. In every case the loss decreases smoothly, and the three



(a) Training of the initial policies with PPO.



(b) Training action loss of the surrogate policies.




Figure 1.4: Training of the initial and surrogate policies.

independently trained policies follow almost identical trajectories, converging to the same value.

We further evaluate the ability of the surrogate policies to replicate the behavior of the initial policies using the *Behavior Reconstruction* (Table 1.3). The surrogate policies reach a *Behavior Reconstruction* of 0.96 on average over 100 evaluation episodes, ranging from 0.83 for Breakout to 1.00 for Ms. Pac-Man and Surround. They therefore reproduce the decisions of the initial policies closely enough for their latent representations to be analyzed in place of those of the initial policies.

Environment	<i>Behavior Reconstruction</i>	<i>Abstraction Score</i>	<i>Clustering Universality</i>
Breakout	0.83 $\pm$ 0.03	0.82 $\pm$ 0.07	0.48 $\pm$ 0.11
Mario Bros	0.99 $\pm$ 0.03	0.71 $\pm$ 0.14	0.55 $\pm$ 0.10
Ms. Pac-Man	1.00 $\pm$ 0.03	0.43 $\pm$ 0.09	0.65 $\pm$ 0.08
Pong	0.95 $\pm$ 0.01	0.99 $\pm$ 0.00	0.49 $\pm$ 0.05
Space Invaders	0.97 $\pm$ 0.07	0.69 $\pm$ 0.06	0.67 $\pm$ 0.07
Surround	1.00 $\pm$ 0.01	0.55 $\pm$ 0.07	0.61 $\pm$ 0.09

Table 1.3: Evaluation of *Behavior Reconstruction*, *Abstraction Score* and *Clustering Universality* across the six Atari environments.

1.3.3 Latent space analysis



With the behavioral fidelity of the surrogate policies established, this section analyzes their latent representations using the *Abstraction Score* and *Clustering Universality*, also reported in Table 1.3.

For the remainder of the analysis, a random subset of 50,000 observations is embedded into the Latent Space (LS) of each surrogate policy, using the representation extracted directly after the last convolutional layer, and the resulting point cloud $\mathcal{E}$ is partitioned with k -means clustering [MacQueen, 1967].

Figure 1.5a reports the evolution of *Clustering Universality* with respect to the number of clusters. After the first few values of K , every curve settles onto a plateau whose level depends on the environment, from about 0.45–0.50 for Breakout and Pong to about 0.65 for Ms. Pac-Man and Space Invaders. Since *Clustering Universality* is computed from the Adjusted Rand Index (ARI), and an ARI of 0.5 already corresponds to substantial agreement [Amezquita et al., 2020], these values indicate that independently trained surrogate policies recover a similar cluster structure in every environment, and that this structure is stable with respect to the number of clusters. The strength of this agreement varies across environments, but remains clearly above chance everywhere. The absence of a single optimal number of clusters suggests that K should be chosen according to the level of granularity desired for the analysis.

The *Abstraction Score*, in contrast, varies markedly across environments. It is close to 1 for Pong (0.99) and high for Breakout, Mario Bros and Space Invaders (0.69 to 0.82): in these environments the latent clusters are largely independent of a direct clustering of the observations or of the action distributions. It is lower for Surround (0.55) and Ms. Pac-Man (0.43). In these two environments the position of the agent on the screen largely determines both the raw observation and the decision to take, so a clustering of the LS naturally coincides with a clustering of the observations, and the *Abstraction Score* cannot separate the two; a low score does not by itself imply that the representation is not abstract.

With fewer clusters, CAR highlights broad strategic patterns. For instance, in the Mario Bros environment, the two main groups correspond to whether the POW block is present (highlighted in red in Figure 1.6A) or not. This is particularly interesting, as the use of the POW block is a key determinant of

(a) Evolution of *Clustering Universality* with respect to the number of clusters.(b) Evolution of *Abstraction Score* with respect to the number of clusters.

▲ Breakout ■ Mario Bros ● Ms. Pac-Man
 ✕ Pong ● Space Invaders ◆ Surround


Figure 1.5: Evolution of *Clustering Universality* and *Abstraction Score* with respect to the number of clusters.

strategy in this environment [AtariAge, 2024].


With a larger number of clusters, the analysis captures more fine-grained tactical behaviors, such as “Attacking the left” and “Navigating through the level” (see Figure 1.6B).

Both perspectives are valid, depending on whether the goal is to study strategic or tactical behaviors. Similar interpretable structures can be observed across the evaluated environments.

(A) The case with 2 clusters.



The POW block is present.



The POW block is absent.

(B) The case with 5 clusters.




The POW block is present.




Move on the left platform.



Navigating through the level.



Attacking the left.



Attacking the right.




Figure 1.6: Effect of the number of clusters on the groupings identified by CAR in the Mario Bros environment.

271 Acronyms

272 **APPO** Asynchronous Proximal Policy Optimization 2

273 **ARI** Adjusted Rand Index 10

274 **CAR** Clustering Agent Representation 1–7, 10, 12

275 **DL** Deep Learning 1, 2

276 **DQN** Deep Q-Networks 2

277 **GPU** Graphics Processing Unit 3

278 **LS** Latent Space 10

279 **PPO** Proximal Policy Optimization 2, 7–9

280 **RL** Reinforcement Learning 1–4

Bibliography

- [Amezquita et al., 2020] Amezquita, R. A., Lun, A. T., Becht, E., Carey, V. J., Carpp, L. N., Geistlinger, L., Marini, F., Rue-Albrecht, K., Risso, D., Soneson, C., et al. (2020). Orchestrating single-cell analysis with Bioconductor, volume 17. Nature Publishing Group US New York. **Book**.
- [Apache Arrow, 2016] Apache Arrow (2016). Apache Arrow and Parquet. <https://arrow.apache.org>.
- [AtariAge, 2024] AtariAge (2024). Mario bros. atari 2600. **Web**.
- [Bellemare et al., 2013] Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. Journal of Artificial Intelligence Research, 47:253–279. **Article**.
- [Espeholt et al., 2018] Espeholt, L., Soyer, H., Munos, R., Simonyan, K., Mnih, V., Ward, T., Doron, Y., Firoiu, V., Harley, T., Dunning, I., Legg, S., and Kavukcuoglu, K. (2018). IMPALA: Scalable distributed deep-RL with importance weighted actor-learner architectures. In Proceedings of the 35th International Conference on Machine Learning (ICML).
- [Falcon, 2019] Falcon, W. (2019). PyTorch Lightning. <https://github.com/Lightning-AI/pytorch-lightning>.
- [Liang et al., 2018] Liang, E., Liaw, R., Nishihara, R., Moritz, P., Fox, R., Goldberg, K., Gonzalez, J. E., Jordan, M. I., and Stoica, I. (2018). RLlib: Abstractions for distributed reinforcement learning. In Proceedings of the 35th International Conference on Machine Learning (ICML).
- [Machado et al., 2017] Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., and Bowling, M. (2017). Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. **Article. Web**.
- [MacQueen, 1967] MacQueen, J. (1967). Multivariate observations. In Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability, volume 1, pages 281–297. **Article**.
- [Mnih et al., 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.

- 314 [Moritz et al., 2018] Moritz, P., Nishihara, R., Wang, S., Tumanov, A., Liaw, R.,
315 Liang, E., Elibol, M., Yang, Z., Paul, W., Jordan, M. I., and Stoica, I. (2018).
316 Ray: A distributed framework for emerging AI applications. In 13th USENIX
317 Symposium on Operating Systems Design and Implementation (OSDI), pages
318 561–577.
- 319 [Paszke et al., 2019] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J.,
320 Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison,
321 A., Köpf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy,
322 S., Steiner, B., Fang, L., Bai, J., and Chintala, S. (2019). PyTorch: An
323 imperative style, high-performance deep learning library. In Advances in
324 Neural Information Processing Systems (NeurIPS).
- 325 [RAPIDS, 2023] RAPIDS (2023). RAPIDS cuML: GPU-accelerated machine
326 learning. <https://github.com/rapidsai/cuml>.
- 327 [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and
328 Klimov, O. (2017). Proximal policy optimization algorithms.
- 329 [Towers et al., 2024] Towers, M., Kwiatkowski, A., Terry, J., Balis, J. U.,
330 De Cola, G., Deleu, T., Goulão, M., Kallinteris, A., Krimmel, M., KG,
331 A., et al. (2024). Gymnasium: A standard interface for reinforcement learning
332 environments.