

# 1 Explainability in Reinforcement Learning

2 Maxime Alaarabiou

3 Monday 21<sup>st</sup> September, 2026

# Chapter 1

## Experiments

This chapter presents the experimental study of the Clustering Agent Representation (CAR) method introduced in the previous chapter. It focuses on evaluating the method’s ability to identify conceptual structures. The remainder of this chapter is organized as follows:

- Section 1.1 describes the software implementation of the CAR pipeline.
- Section 1.2 details the experimental configuration, covering the environments, the reinforcement learning setup, and the surrogate policy training.
- Section 1.3 presents the training of the initial and surrogate policies and evaluates how faithfully the surrogate policies reproduce the behavior of the initial policy.
- Section 1.4 provides a quantitative analysis of the latent representations using the *Clustering Universality* (CU), *Abstraction Score* (AS), and *Adjusted Clustering Universality* (ACU) metrics.
- Section 1.5 provides a qualitative analysis of the resulting clusters by manually examining their semantic content.
- Section 1.6 summarizes the findings and discusses their implications.

### 1.1 Software Implementation

This section details the experimental evaluation of the proposed method. Deep Learning (DL) and Reinforcement Learning (RL) are strongly empirical fields, where experimental evaluation plays an important role in assessing how methods perform in practice. Particular attention is therefore given to the implementation of CAR. It relies on established libraries that provide efficient implementations of its main components. The code is made publicly available as open-source to support reproducible replication of the reported results.<sup>1</sup>

The CAR implementation is organized as a pipeline of five stages. Each stage performs a specific operation and produces the artifacts required by the subsequent stages. The stages are implemented as independent executables

---

<sup>1</sup>[https://github.com/Maxalaar/Clustering\\_Agent\\_Latent\\_Space](https://github.com/Maxalaar/Clustering_Agent_Latent_Space)



Figure 1.1: The CAR pipeline.

33 allowing each stage to be rerun independently. The CAR pipeline is illustrated in  
 34 Figure 1.1, and Table 1.1 summarizes the libraries used throughout the pipeline.  
 35 The remainder of this section describes the purpose and implementation of each  
 36 stage.

### 37 1.1.1 Initial policy training

38 The first stage trains the initial policy using RLlib [Moritz et al., 2018, Liang  
 39 et al., 2018]. RLlib provides scalable implementations of RL algorithms and  
 40 supports distributed training. The use of RLlib also allows CAR to be applied to  
 41 environments implemented using the Gymnasium interface [Towers et al., 2024].  
 42 By default, the initial policy is trained using Proximal Policy Optimization  
 43 (PPO) [Schulman et al., 2017]. Other RLlib algorithms, such as Asynchronous  
 44 Proximal Policy Optimization (APPO) [Espeholt et al., 2018], Deep Q-Networks  
 45 (DQN) [Mnih et al., 2013], and DreamerV3 [Hafner et al., 2023], can also be  
 46 selected through the same configuration.

47 The policy and critic networks are implemented using PyTorch [Paszke et al.,  
 48 2019]. PyTorch is the standard deep learning frameworks and is widely used for  
 49 developing and training Neural Network (NN). It provides GPU acceleration  
 50 and supports a wide range of NN architectures. This makes it well suited for  
 51 implementing the different policy and critic architectures used in the experiments.

| Pipeline Stage                | Library                                         | Description                                                                       |
|-------------------------------|-------------------------------------------------|-----------------------------------------------------------------------------------|
| Initial policy training       | RLlib [Moritz et al., 2018, Liang et al., 2018] | Distributed RL library used to train the initial policy.                          |
|                               | PyTorch [Paszke et al., 2019]                   | Deep learning framework implementing the policy and critic networks.              |
|                               | Gymnasium [Towers et al., 2024]                 | Common interface used to plug in any RL environment.                              |
| Trajectory dataset generation | RLlib [Moritz et al., 2018, Liang et al., 2018] | Distributed environment runners used to collect trajectories in parallel.         |
|                               | Apache Arrow [Apache Arrow, 2016]               | Columnar file format (Parquet) used to store the trajectory datasets.             |
| Surrogate policy training     | PyTorch Lightning [Falcon, 2019]                | Structured training loop used for the behavior cloning of the surrogate policies. |
| Surrogate policy evaluation   | RLlib [Moritz et al., 2018, Liang et al., 2018] | Wraps each surrogate policy as an RL policy to evaluate it in the environment.    |
| Latent space analysis         | cuML [RAPIDS, 2023]                             | GPU-accelerated clustering algorithms used to analyze the latent representations. |

Table 1.1: Main software libraries used at each stage of the CAR pipeline.

### 1.1.2 Trajectory dataset generation

The second stage builds the dataset used to train the surrogate policies. RLlib is used to collect trajectories in parallel using multiple environment runners. For each step, the dataset stores the observation together with the parameters of the action distribution produced by the initial policy.

Trajectories are stored as Parquet files [Apache Arrow, 2016]. Its columnar structure allows efficient data access and supports concurrent processing. This is useful here because the datasets contain around one million observation–action pairs and cannot be loaded entirely into memory.

### 1.1.3 Surrogate policy training

The third stage trains the surrogate policies using behavior cloning on the trajectory dataset. This provides a supervised learning approach for replicating the initial policy. The training procedure is implemented using PyTorch Lightning [Falcon, 2019]. PyTorch Lightning provides a simplified interface for training NN with PyTorch. It handles common training operations such as data loading, optimization, checkpointing, and GPU execution.

### 1.1.4 Surrogate policy evaluation

The fourth stage evaluates how faithfully the surrogate policies reproduce the behavior of the initial policy. Each surrogate is wrapped as an RLlib policy and evaluated in the original environment. Its mean cumulative reward over multiple episodes is compared with that of the initial policy. Evaluation episodes are collected using parallel environment runners.

This stage computes the *Behavior Reconstruction*. It therefore verifies that the surrogate policies reproduce the behavior of the initial policy sufficiently well before their latent representations are analyzed. If a surrogate does not meet the required level of behavior reconstruction, its training can be repeated.

### 1.1.5 Latent space analysis

The fifth stage performs the main analysis of CAR. The trajectory dataset and trained surrogate policies are used to generate the latent representations. These representations form the point clouds that are analyzed using clustering algorithms.

The clustering step relies on the cuML library [RAPIDS, 2023], which provides GPU-accelerated clustering algorithms. This reduces the computational cost and allows point clouds containing tens of thousands of representations to be analyzed.

The resulting partitions are used to compute the *Clustering Universality* and the *Abstraction Score*. The stage also produces the material required for the qualitative analysis, including the clustered point clouds and the cluster assignment of each observation.

## 1.2 Experimental Configuration

This section describes the configuration of the experiments reported in this chapter. It covers the environments, the neural architectures used, and the main parameters of the pipeline.

### 1.2.1 Environments

All experiments presented in this manuscript use environments from the Atari benchmark [Bellemare et al., 2013], accessed through the Arcade Learning Environment. The use of Atari environments offers several advantages:

- **Standard benchmark:** Atari is widely used in the RL community, making the experimental results easier to compare with existing work.
- **Limited environment-specific bias:** The environments were not specifically designed for RL, reducing the risk of introducing biases related to their design.
- **Diversity:** The benchmark provides a diverse set of well-known environments with different types of tasks and levels of complexity, allowing CAR to be evaluated under different conditions.

The configuration of the Atari environments follows common practice in RL research [Machado et al., 2017]. In each environment, the agent receives visual observations in the form of raw pixels. Each image is resized to  $64 \times 64$  pixels and converted to grayscale to reduce memory requirements. To provide information about the dynamics of the environment, four consecutive observations are stacked along the channel dimension. The resulting observation has a shape of  $64 \times 64 \times 4$ . This allows the agent to infer the motion of objects from successive observations.

The actions correspond to the different buttons available on the Atari controller. The resulting action space is therefore discrete. Since Atari games are fully deterministic, sticky actions are introduced to add stochasticity to the trajectories. With a probability of 0.05, the environment repeats the action selected at the previous step instead of applying the new action. This prevents the agent from simply memorizing a fixed sequence of actions during training.




Figure 1.2: The six Atari environments used to evaluate CAR.

To reduce the duration of each episode, frame skipping is used. The agent selects a new action every five frames and repeats it for the following frames. This reduces the number of decisions required during an episode while maintaining sufficient temporal resolution for effective control.

The experiments use the six environments described below:

- **Breakout:** The agent moves a paddle at the bottom of the screen (Figure 1.2a) to bounce a ball into a wall of bricks. Each broken brick scores one point, and the ball speeds up as the wall is cleared.
- **Mario Bros.:** A single-screen platform game (Figure 1.2b) in which the player clears each phase by hitting enemies from below to flip them, then kicking them off the platforms before they recover. A central POW block can flip all enemies at once but can only be used a few times per phase.
- **Ms. Pac-Man:** The player moves through a maze (Figure 1.2c), eating pellets while avoiding four ghosts. Power pellets temporarily make the ghosts edible.
- **Pong:** A two-player game (Figure 1.2d) in which the agent controls one paddle and must send the ball past the opponent.
- **Space Invaders:** The agent moves a cannon along the bottom of the screen (Figure 1.2e), firing at descending rows of aliens while avoiding their



Figure 1.3: Architecture of the surrogate policies.

shots and using destructible shields. A bonus saucer occasionally crosses the top of the screen.

- **Surround:** A two-player game on a grid (Figure 1.2f) in which each player leaves a trail behind a moving cursor. The goal is to force the opponent to crash into a trail.

### 1.2.2 Neural architecture

The surrogate policies use a convolutional NN to process the visual observations. The convolutional part consists of four layers with 16, 32, 64, and 128 channels, using  $4 \times 4$  kernels with a stride of 2. Each convolutional layer is followed by a LeakyReLU activation. The resulting representation is then flattened and passed through a multilayer perceptron that produces the parameters of the action distribution.

The architecture of the multilayer perceptron varies across environments. For **Pong** and **Space Invaders**, it contains a single hidden layer with 128 units. For the other four environments, it contains three hidden layers with 512, 256, and 128 units. The size of the output layer matches the number of actions available in each environment: 4 for **Breakout**, 18 for **Mario Bros**, 9 for **Ms. Pac-Man**, 6 for **Pong**, 6 for **Space Invaders**, and 5 for **Surround**.

Figure 1.3 illustrates both configurations, together with the shape of the representation produced by each convolutional layer. Throughout the remainder of this chapter, the analyzed representations are numbered from 1 to 5. These representations correspond to the four convolutional layers and the flattened representation, which are common to all six environments.

The representations produced by the multilayer perceptron are numbered separately. For **Pong** and **Space Invaders**, which use a single hidden layer, the representation is numbered 6a. For the other four environments, which use three hidden layers, the representations are numbered 6b, 7b, and 8b.

### 1.2.3 Other CAR parameters

Several other parameters must be defined for the different stages of the CAR pipeline. These parameters are less critical to the overall methodology than those discussed above, but they still need to be specified to reproduce the experiments.

| Environment configuration            |                      |
|--------------------------------------|----------------------|
| Frame size                           | $64 \times 64$       |
| Frames stacked per observation       | 4                    |
| Frame skip                           | 5                    |
| Sticky-action probability            | 0.05                 |
| Initial policy training              |                      |
| Algorithm                            | PPO                  |
| Optimizer                            | Adam                 |
| Learning rate                        | $1.5 \times 10^{-4}$ |
| GAE $\lambda$                        | 0.95                 |
| Policy clipping $\varepsilon$        | 0.1                  |
| Value-function clipping              | 10                   |
| Entropy coefficient                  | 0.01                 |
| KL coefficient                       | 0.5                  |
| Epochs per update                    | 10                   |
| Gradient clipping (global norm)      | 100                  |
| Surrogate policy training            |                      |
| Surrogate policies per environment   | 3                    |
| Observation–action pairs per dataset | $\geq 10^6$          |
| Batch size                           | 2048                 |
| Optimizer                            | Adam                 |
| Learning rate                        | $1 \times 10^{-4}$   |
| Latent space analysis                |                      |
| Embedding per surrogate policy       | 20 000               |
| Clustering algorithm                 | $k$ -means           |
| Distance                             | Euclidean distance   |
| Number of clusters $K$               | 2 to 20              |


Table 1.2: Main experimental parameters.

Since each stage introduces its own configuration choices, there are too many parameters to discuss them all in detail. The main parameters used across the experiments are summarized in Table 1.2.


These settings are specific to the experiments presented in this chapter and do not restrict the CAR methodology itself. The different stages of the pipeline are modular and can be configured or replaced depending on the application. The evaluation metrics provide an objective basis for comparing these alternatives and assessing their effect on the consistency and meaningfulness of the resulting clusters.

### 1.3 Policy Performance and Surrogate Fidelity

This section reports the results of training the initial and surrogate policies for CAR, applied to the six environments. The first part verifies that the initial policies reach a competent level of play and evaluates how faithfully the surrogate policies reproduce their behavior using the *Behavior Reconstruction*.



(a) Training of the initial policies with PPO.



(b) Training action loss of the surrogate policies.



Figure 1.4: Training of the initial and surrogate policies.

### 1.3.1 Initial policy training

Figure 1.4a shows how the initial policies improve during PPO training. For each environment, the mean episode return is normalized so that 0% and 100% correspond to the lowest and highest mean return observed during that run. In every environment the return increases and then reaches a plateau, at which point training is stopped. The number of iterations required to reach this plateau varies from a few hundred, for **Pong**, to several thousand, for **Surround** and **Ms. Pac-Man**. For the subsequent analysis, the checkpoint with the best performance is retained rather than the final checkpoint.

To confirm that the trained policies play competently, their scores are compared with the reference scores reported by [Machado et al., 2017]. Each initial policy reaches a score of the same order as these references, so it constitutes a reasonable target for imitation.

### 1.3.2 Surrogate policy training

Figure 1.4b shows, for each environment, the training action loss of the three surrogate policies. In all cases, the loss decreases smoothly, and the three independently trained policies converge to the same value.

We further evaluate the ability of the surrogate policies to replicate the

| Environment           | <i>Behavior Reconstruction</i> |
|-----------------------|--------------------------------|
| <b>Breakout</b>       | 0.83 $\pm$ 0.03                |
| <b>Mario Bros</b>     | 0.99 $\pm$ 0.03                |
| <b>Ms. Pac-Man</b>    | 1.00 $\pm$ 0.03                |
| <b>Pong</b>           | 0.95 $\pm$ 0.01                |
| <b>Space Invaders</b> | 0.97 $\pm$ 0.07                |
| <b>Surround</b>       | 1.00 $\pm$ 0.01                |

Table 1.3: Evaluation of *Behavior Reconstruction* across the six Atari environments.

behavior of the initial policies using the *Behavior Reconstruction* (Table 1.3). The surrogate policies reach a *Behavior Reconstruction* of 0.96 on average over 100 evaluation episodes, ranging from 0.83 for **Breakout** to 1.00 for **Ms. Pac-Man** and **Surround**. They therefore reproduce the decisions of the initial policies closely enough for their latent representations to be analyzed.

## 1.4 Clustering Quality of the Latent Space


This section analyzes the Latent Space (LS) of the surrogate policies using CU, AS, and ACU. To do so, a random subset of 20 000 observations is selected for each environment. This subset is then embedded into the LS of each surrogate policy at each analyzed layer. This produces a point cloud for each combination of environment, surrogate policy, and layer. Each point cloud is partitioned using *k*-means clustering [MacQueen, 1967] for each number of clusters from 2 to 20.

The three metrics are then computed from these point clouds and their corresponding clusterings. The full results across all environments and analyzed layers are reported in Appendix ???. Two complementary analyses are then conducted on these results. First, the evolution of the three metrics across the different layers is examined to study the effect of representation depth. Second, for each environment, the layer achieving the highest average ACU is selected and analyzed in more detail across the different numbers of clusters.


### 1.4.1 Analysis across layers

The first analysis examines how the three metrics evolve with the depth of the analyzed layer. For each layer, CU, AS, and ACU are averaged over the tested cluster numbers. Figure 1.5 reports these averages as a function of layer depth.


Across environments, CU tends to be highest in the earliest layers and decreases with depth. AS follows the opposite trend and generally increases as the representation becomes less tied to the raw observation, the action distribution, and the network’s initialization. This increase is however not monotonic in the deepest layers. For **Mario Bros** and **Ms. Pac-Man**, AS decreases at the last layer, from 0.83 to 0.68 and from 0.64 to 0.60 respectively, as the latent clustering becomes closer to the clustering of the action distribution, this layer being directly followed by the output layer. As ACU combines these two effects, its evolution depends on the environment.



(a) CU with respect to the layer depth.



(b) AS with respect to the layer depth.



(c) ACU with respect to the layer depth.




Figure 1.5: CU, AS, and ACU averaged over cluster numbers as a function of layer depth.

235 The corresponding linear fits show that CU decreases from 0.72 at the first  
 236 layer to 0.50 at the last, with a slope of approximately  $-0.04$  per layer, while AS  
 237 increases from 0.28 to 0.69, with a slope of approximately  $+0.06$  per layer. Over  
 238 the same range, ACU increases more moderately from 0.00 to 0.19, with a slope  
 239 of approximately  $+0.03$  per layer. This suggests that, despite the environment-  
 240 specific variations, deeper representations tend on average to provide a better  
 241 trade-off between CU and AS.

#### 242 1.4.2 Analysis of the selected layers

243 The second analysis examines the effect of the number of clusters on the three  
 244 metrics. For each environment, the layer with the highest average ACU across  
 245 the tested cluster numbers is selected. Table 1.4 reports the selected layer,


246 together with the corresponding average CU, AS, and ACU.

247 Figure 1.6 reports the evolution of CU, AS, and ACU as the number of clusters  
248 varies. **Pong** achieves the highest average ACU among the six environments,  
249 with a value of  $0.45 \pm 0.06$ . This results from a moderate CU combined with  
250 an AS that saturates from the third layer onward. In contrast, **Ms. Pac-Man**  
251 achieves the lowest AS, with  $0.60 \pm 0.05$ . The resulting ACU,  $0.17 \pm 0.06$ , is the  
252 lowest across the six environments.


253 The quantitative analysis provides a measure of the consistency and abstrac-  
254 tion of the identified clusters, but does not directly reveal their semantic content.  
255 The following section therefore provides a qualitative analysis of the clusters to  
256 examine the concepts they represent.

| Environment    | Layer | CU              | AS              | ACU             |
|----------------|-------|-----------------|-----------------|-----------------|
| Breakout       | 6b    | $0.59 \pm 0.06$ | $0.82 \pm 0.05$ | $0.41 \pm 0.06$ |
| Mario Bros     | 3     | $0.62 \pm 0.12$ | $0.62 \pm 0.15$ | $0.24 \pm 0.10$ |
| Ms. Pac-Man    | 8b    | $0.58 \pm 0.08$ | $0.60 \pm 0.05$ | $0.17 \pm 0.06$ |
| Pong           | 4     | $0.50 \pm 0.04$ | $0.95 \pm 0.03$ | $0.45 \pm 0.06$ |
| Space Invaders | 4     | $0.63 \pm 0.11$ | $0.73 \pm 0.06$ | $0.35 \pm 0.09$ |
| Surround       | 8b    | $0.54 \pm 0.10$ | $0.71 \pm 0.05$ | $0.25 \pm 0.06$ |


Table 1.4: Layer that maximizes the average ACU across  $K \in [2, 20]$  for each environment, together with the corresponding average CU, AS, and ACU over  $K$ .



(a) CU with respect to the number of clusters.



(b) AS with respect to the number of clusters.



(c) ACU with respect to the number of clusters.




Figure 1.6: CU, AS, and ACU with respect to the number of clusters, for each of the six environments, for the layer that maximizes the average ACU in that environment.

## 1.5 Semantic Interpretation of the Clusters

The metrics reported in Section 1.4 evaluate the LS clusterings quantitatively but do not indicate what the resulting clusters actually represent. This section complements them with a qualitative, hand-labeled inspection of the clusters obtained for each environment.

For each environment, two clusterings are selected for the qualitative analysis. They are obtained using two different numbers of clusters and may correspond to different layers of the network. The choice of the layer and the number of clusters is based on the following criteria:

- **High ACU:** The selected clustering should achieve a sufficiently high ACU, indicating a good balance between consistency and abstraction.
- **Manual analysis:** The number of clusters must remain small enough to allow the resulting clusters to be examined manually.
- **Limited use of two-cluster partitions:** Clusterings with two clusters often achieve a high ACU, but provide limited information about the organization of the LS. Therefore, at most one of the two selected clusterings uses two clusters, while the other uses a larger number of clusters to provide a more detailed view of the latent space.

### 1.5.1 Breakout

Two clusterings are examined for **Breakout** (Table 1.5), obtained at layer 4 with  $K = 2$  and at layer 6b with  $K = 5$ .

At layer 4 (Figure 1.7), the partition captures the position of the ball relative to the brick wall. The ball remains below the wall in (cluster A). The ball has passed behind the wall in (cluster B), allowing it to continue clearing bricks without further interaction with the paddle.

At layer 6b (Figure 1.8), the partition again isolates situations in which the ball has passed behind the wall (cluster E). It also identifies the digging of a tunnel through the left side of the wall (cluster C). Finally, it distinguishes different paddle positions at the beginning of the episode (clusters D, F, and G).

The clustering at layer 4 with  $K = 2$  is almost perfectly reproducible across surrogate policies, with a CU of 0.99 and an ACU of 0.87. The clustering at layer 6b with  $K = 5$  is less consistent, with a CU of 0.79 and an ACU of 0.61, but reveals additional information about the organization of the LS.

| Layer 4  | $K = 2$                                             | CU 0.99 | AS 0.88 | ACU 0.87 |
|----------|-----------------------------------------------------|---------|---------|----------|
| A        | The ball is below the brick wall.                   |         |         |          |
| B        | The ball passes behind the brick wall.              |         |         |          |
| Layer 6b | $K = 5$                                             | CU 0.79 | AS 0.82 | ACU 0.61 |
| C        | Digging a tunnel through the left side of the wall. |         |         |          |
| D        | Start of the episode, paddle on the right.          |         |         |          |
| E        | The ball is behind the brick wall.                  |         |         |          |
| F        | Start of the episode, paddle in the middle.         |         |         |          |
| G        | Start of the episode, paddle on the left.           |         |         |          |

Table 1.5: CU, AS, and ACU of the two clusterings analyzed in the **Breakout** environment.




Figure 1.7: Clusters A–B identified by CAR in the **Breakout** environment (layer 4,  $K = 2$ ).




Figure 1.8: Clusters C–G identified by CAR in the **Breakout** environment (layer 6b,  $K = 5$ ).

### 1.5.2 Mario Bros

Two clusterings are examined for **Mario Bros** (Table 1.6), both obtained at layer 3, with  $K = 3$  and  $K = 5$ .

With  $K = 3$  (Figure 1.9), the partition is organized around the location of the action. Mario fights enemies on the left side of the screen in (cluster A). He fights enemies on the right side of the screen in (cluster C). The POW block is present in (cluster B), isolating a fixed and recognizable element of the level.

With  $K = 5$  (Figure 1.10), the partition provides a more detailed description of the situations occurring in the level. Mario is on the top-left platform in (cluster D). Mario attacks enemies at the top left in (cluster F). Mario hits enemies from below in (cluster H). Mario simply moves through the level in (cluster G). The POW block remains isolated in (cluster E).

Increasing  $K$  from 3 to 5 lowers CU from 0.84 to 0.75. The AS remains nearly unchanged, with values of 0.66 and 0.65. This results in a decrease in ACU from 0.50 to 0.40.

| Layer 3 | $K = 3$                                               | CU 0.84 | AS 0.66 | ACU 0.50 |
|---------|-------------------------------------------------------|---------|---------|----------|
| A       | Mario fights enemies on the left side of the screen.  |         |         |          |
| B       | The POW block is present.                             |         |         |          |
| C       | Mario fights enemies on the right side of the screen. |         |         |          |
| Layer 3 | $K = 5$                                               | CU 0.75 | AS 0.65 | ACU 0.40 |
| D       | Mario is on the top-left platform.                    |         |         |          |
| E       | The POW block is present.                             |         |         |          |
| F       | Mario attacks the enemies at the top right.           |         |         |          |
| G       | Mario moves through the level.                        |         |         |          |
| H       | Mario hits enemies from below on the left side.       |         |         |          |

Table 1.6: CU, AS, and ACU of the two clusterings analyzed in the **Mario Bros** environment.




Figure 1.9: Clusters A–C identified by CAR in the **Mario Bros** environment (layer 3,  $K = 3$ ).




Figure 1.10: Clusters D–H identified by CAR in the **Mario Bros** environment (layer 3,  $K = 5$ ).

### 1.5.3 Ms. Pac-Man

Two clusterings are examined for **Ms. Pac-Man** (Table 1.7), obtained at layer 6b with  $K = 3$  and at layer 4 with  $K = 5$ .

With  $K = 3$  (Figure 1.11), the partition captures the progress and state of the game rather than a specific spatial feature of the observation. The start of the episode is captured in (cluster B). The middle of the episode is captured in (cluster A). The end of the episode is captured in (cluster C).

With  $K = 5$  (Figure 1.12), the partition provides a more detailed description of the episode. The initial animation, during which the character cannot move, is isolated in (cluster H). The beginning of actual play is captured in (cluster F). A situation specific to the game mechanics, in which Pac-Man is in booster mode and waits in front of the ghost spawn to eat the ghosts, is isolated in (cluster D). The end of the episode, when at most one power pellet remains, is separated into two similar clusters (clusters E and G).

The two clusterings have the same CU of 0.88. The clustering at layer 4 has a lower AS of 0.47 compared with 0.52 at layer 6b. This results in a lower ACU of 0.35 compared with 0.40.

| Layer 6b | $K = 3$                                                                             | CU 0.88 | AS 0.52 | ACU 0.40 |
|----------|-------------------------------------------------------------------------------------|---------|---------|----------|
| A        | Middle of the episode.                                                              |         |         |          |
| B        | Start of the episode.                                                               |         |         |          |
| C        | End of the episode.                                                                 |         |         |          |
| Layer 4  | $K = 5$                                                                             | CU 0.88 | AS 0.47 | ACU 0.35 |
| D        | Pac-Man is in booster mode and waits in front of the ghost spawn to eat the ghosts. |         |         |          |
| E        | End of the episode, at most one power pellet remains (similar to cluster G).        |         |         |          |
| F        | Start of the episode.                                                               |         |         |          |
| G        | End of the episode, at most one power pellet remains (similar to cluster E).        |         |         |          |
| H        | Start-of-episode animation, the character cannot move.                              |         |         |          |

Table 1.7: CU, AS, and ACU of the two clusterings analyzed in the **Ms. Pac-Man** environment.




Figure 1.11: Clusters A–C identified by CAR in the **Ms. Pac-Man** environment (layer 6b,  $K = 3$ ).




Figure 1.12: Clusters D–H identified by CAR in the **Ms. Pac-Man** environment (layer 4,  $K = 5$ ).

#### 1.5.4 Pong

Two clusterings are examined for **Pong** (Table 1.8), obtained at layer 2 with  $K = 2$  and at layer 3 with  $K = 3$ .

At layer 2 (Figure 1.13), the partition separates the observations according to the number of digits in the player's score. A one-digit score corresponds to (cluster A). A two-digit score corresponds to (cluster B). This reflects the progress of the game, but only through its visual display at the top of the screen. The partition therefore captures a feature of the observation that provides no information about how the agent plays.

At layer 3 (Figure 1.14), the partition instead follows the vertical position of the ball. The ball is at the top of the field in (cluster C). It is in the middle of the field in (cluster D). It is at the bottom of the field in (cluster E). This feature directly determines where the agent needs to place its paddle.

The clustering at layer 2 is perfectly reproducible and perfectly abstract, with CU, AS, and ACU all equal to 1.00. The clustering at layer 3 is less consistent, with a CU of 0.67, and reaches a lower ACU of 0.65.

| Layer 2 | $K = 2$                            | CU 1.00 | AS 1.00 | ACU 1.00 |
|---------|------------------------------------|---------|---------|----------|
| A       | The player's score has two digits. |         |         |          |
| B       | The player's score has one digit.  |         |         |          |
| Layer 3 | $K = 3$                            | CU 0.67 | AS 0.98 | ACU 0.65 |
| C       | The ball is at the top.            |         |         |          |
| D       | The ball is in the middle.         |         |         |          |
| E       | The ball is at the bottom.         |         |         |          |

Table 1.8: CU, AS, and ACU of the two clusterings analyzed in the **Pong** environment.




Figure 1.13: Clusters A–B identified by CAR in the **Pong** environment (layer 2,  $K = 2$ ).




Figure 1.14: Clusters C–E identified by CAR in the **Pong** environment (layer 3,  $K = 3$ ).

### 1.5.5 Space Invaders

Two clusterings are examined for **Space Invaders** (Table 1.9), obtained at layer 4 with  $K = 4$  and at layer 2 with  $K = 5$ .

At layer 4 (Figure 1.15), the partition tracks the progression of the episode. The start of the episode is captured in (cluster D). The middle of the episode is captured in (cluster B). The end of the episode is captured in (cluster C). A specific behavior in which the agent has cleared the leftmost column and is positioned to catch the flying saucer that occasionally crosses the top of the screen is captured in (cluster A).

At layer 2 (Figure 1.16), the partition again tracks the progression of the episode. The start of the episode is captured in (cluster G). The middle of the episode is captured in (cluster F). The end of the episode is captured in (cluster I). Attacking the first two columns of the formation is captured in (cluster E). Trying to catch the flying saucer is captured in (cluster H).

The clustering at layer 4 has a CU of 0.92 and a AS of 0.69, resulting in a ACU of 0.61. The clustering at layer 2 has the same CU of 0.92, but a lower AS of 0.49, resulting in a lower ACU of 0.41.

| Layer 4 | $K = 4$                                                                   | CU 0.92 | AS 0.69 | ACU 0.61 |
|---------|---------------------------------------------------------------------------|---------|---------|----------|
| A       | The left column has been cleared; the flying saucer occasionally appears. |         |         |          |
| B       | Middle of the episode.                                                    |         |         |          |
| C       | End of the episode.                                                       |         |         |          |
| D       | Start of the episode.                                                     |         |         |          |
| Layer 2 | $K = 5$                                                                   | CU 0.92 | AS 0.49 | ACU 0.41 |
| E       | Attacks the first two columns of the formation.                           |         |         |          |
| F       | Middle of the episode.                                                    |         |         |          |
| G       | Start of the episode.                                                     |         |         |          |
| H       | Tries to catch the flying saucer.                                         |         |         |          |
| I       | End of the episode.                                                       |         |         |          |

Table 1.9: CU, AS, and ACU of the two clusterings analyzed in the **Space Invaders** environment.




Figure 1.15: Clusters A–D identified by CAR in the **Space Invaders** environment (layer 4,  $K = 4$ ).




Figure 1.16: Clusters E–I identified by CAR in the **Space Invaders** environment (layer 2,  $K = 5$ ).

### 1.5.6 Surround

Two clusterings are examined for **Surround** (Table 1.10), obtained at layer 8b with  $K = 4$  and at layer 7b with  $K = 6$ .

At layer 8b (Figure 1.17), the partition captures the start of the episode in (cluster B). The outcome of the episode, when the opponent loses, is captured in (cluster D). The player's trail is confined to the bottom-left area in (cluster A). The player reaches the top of the level from the left and then moves along the top in (cluster C).

At layer 7b (Figure 1.18), the partition again captures the start of the episode in (cluster J). The outcome of the episode, when the opponent loses, is captured in (cluster F). The player moves in the upper-left corner in (cluster E). The player moves in the lower-left corner in (cluster G). The end of the game is played in the lower-left corner in (cluster H). The player crosses the level from the top left to the bottom right in (cluster I).

The clustering at layer 8b has a CU of 0.67, a AS of 0.71, and a ACU of 0.38. The clustering at layer 7b reaches the same values, with a CU of 0.67, a AS of 0.71, and a ACU of 0.38.

| Layer 8b | $K = 4$                                                             | CU 0.67 | AS 0.71 | ACU 0.38 |
|----------|---------------------------------------------------------------------|---------|---------|----------|
| A        | The agent's trail remains confined to the bottom-left area.         |         |         |          |
| B        | Start of the episode.                                               |         |         |          |
| C        | The agent moves along the top of the level, arriving from the left. |         |         |          |
| D        | The opponent loses by running into a trail.                         |         |         |          |
| Layer 7b | $K = 6$                                                             | CU 0.67 | AS 0.71 | ACU 0.38 |
| E        | The player moves in the upper-left corner.                          |         |         |          |
| F        | The opponent loses the confrontation.                               |         |         |          |
| G        | The player moves in the lower-left corner.                          |         |         |          |
| H        | The end of the game is played in the lower-left corner.             |         |         |          |
| I        | The player starts at the top left and ends at the bottom right.     |         |         |          |
| J        | Start of the episode.                                               |         |         |          |

Table 1.10: CU, AS, and ACU of the two clusterings analyzed in the **Surround** environment.




Figure 1.17: Clusters A–D identified by CAR in the **Surround** environment (layer 8b,  $K = 4$ ).




Figure 1.18: Clusters E–J identified by CAR in the **Surround** environment (layer 7b,  $K = 6$ ).

## 1.6 Conclusion

The quantitative results show that CAR consistently recovers similar clustering structures across independently trained surrogate policies. Across the six environments, the CU frequently exceeds 0.5, indicating that independently trained policies can develop similarly structured latent representations [Amezquita et al., 2020]. The CU also tends to decrease as the network gets deeper, with the highest values generally observed in the earliest layers. Considering consistency alone would therefore favor the earliest layers of the network.

To complement this analysis, the level of abstraction of these representations must also be considered. For this purpose, AS is analyzed and shows that the earliest layers tend to exhibit lower levels of abstraction. Their representations remain strongly related to the observations and the initial structure of the network. In the deepest layers, the abstraction can decrease again in some environments, as the latent clustering becomes closer to the clustering of the action distribution.

This highlights the importance of considering both consistency and abstraction when identifying meaningful latent representations. The ACU addresses this limitation by combining these two properties into a single measure, providing a practical criterion for identifying layers that achieve a suitable balance between representational consistency and abstraction.

These quantitative analyses are complemented by the qualitative analysis, which shows that the clusters identified by CAR correspond to interpretable properties of the environment and the agent’s behavior. Across the different environments, the clusters capture recurring aspects such as episode progression, specific game situations, and behavioral patterns. These observations provide qualitative evidence that the structures recovered by CAR can reflect semantically aspects of the underlying task.

The qualitative analysis also reveals a trade-off associated with the number of clusters. Increasing the number of clusters allows more specific situations to be distinguished, but tends to reduce the stability of the resulting partitions across surrogate policies. Finer clusters have less clear semantic boundaries, reducing agreement across independently trained policies. The choice of the number of clusters therefore involves a trade-off between recovering stable, shared concepts and capturing finer distinctions between situations.

Overall, these results provide empirical support for the assumptions underlying CAR. Independently trained surrogate policies exhibit a degree of convergence in the organization of their Latent Representation (LR) that cannot be explained solely by the considered baselines. The qualitative analysis also shows that these clusters can correspond to meaningful properties of the environment and the agent’s behavior. Taken together, these findings provide empirical support for the *Representation Convergence Hypothesis* (Hyp. ??).

# Acronyms

|     |                                                                                    |
|-----|------------------------------------------------------------------------------------|
| 414 | <b>ACU</b> <i>Adjusted Clustering Universality</i> 1, 9–14, 16, 18, 20, 22, 24, 26 |
| 415 | <b>APPO</b> Asynchronous Proximal Policy Optimization 2                            |
| 416 | <b>AS</b> <i>Abstraction Score</i> 1, 9–12, 14, 16, 18, 20, 22, 24, 26             |
| 417 | <b>CAR</b> Clustering Agent Representation 1–7, 15, 17, 19, 21, 23, 25, 26         |
| 418 | <b>CU</b> <i>Clustering Universality</i> 1, 9–12, 14, 16, 18, 20, 22, 24, 26       |
| 419 | <b>DL</b> Deep Learning 1                                                          |
| 420 | <b>DQN</b> Deep Q-Networks 2                                                       |
| 421 | <b>GPU</b> Graphics Processing Unit 3                                              |
| 422 | <b>LR</b> Latent Representation 26                                                 |
| 423 | <b>LS</b> Latent Space 9, 13, 14                                                   |
| 424 | <b>NN</b> Neural Network 2, 3, 6                                                   |
| 425 | <b>PPO</b> Proximal Policy Optimization 2, 7, 8                                    |
| 426 | <b>RL</b> Reinforcement Learning 1–4                                               |

# Bibliography

- [Amezquita et al., 2020] Amezquita, R. A., Lun, A. T., Becht, E., Carey, V. J., Carpp, L. N., Geistlinger, L., Marini, F., Rue-Albrecht, K., Risso, D., Soneson, C., et al. (2020). Orchestrating single-cell analysis with Bioconductor, volume 17. Nature Publishing Group US New York. **Book**.
- [Apache Arrow, 2016] Apache Arrow (2016). Apache Arrow and Parquet. <https://arrow.apache.org>.
- [Bellemare et al., 2013] Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. Journal of Artificial Intelligence Research, 47:253–279. **Article**.
- [Espeholt et al., 2018] Espeholt, L., Soyer, H., Munos, R., Simonyan, K., Mnih, V., Ward, T., Doron, Y., Firoiu, V., Harley, T., Dunning, I., Legg, S., and Kavukcuoglu, K. (2018). IMPALA: Scalable distributed deep-RL with importance weighted actor-learner architectures. In Proceedings of the 35th International Conference on Machine Learning (ICML).
- [Falcon, 2019] Falcon, W. (2019). PyTorch Lightning. <https://github.com/Lightning-AI/pytorch-lightning>.
- [Hafner et al., 2023] Hafner, D., Pasukonis, J., Ba, J., and Lillicrap, T. (2023). Mastering diverse domains through world models.
- [Liang et al., 2018] Liang, E., Liaw, R., Nishihara, R., Moritz, P., Fox, R., Goldberg, K., Gonzalez, J. E., Jordan, M. I., and Stoica, I. (2018). RLlib: Abstractions for distributed reinforcement learning. In Proceedings of the 35th International Conference on Machine Learning (ICML).
- [Machado et al., 2017] Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., and Bowling, M. (2017). Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. **Article. Web**.
- [MacQueen, 1967] MacQueen, J. (1967). Multivariate observations. In Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability, volume 1, pages 281–297. **Article**.
- [Mnih et al., 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.

- 461 [Moritz et al., 2018] Moritz, P., Nishihara, R., Wang, S., Tumanov, A., Liaw, R.,  
462 Liang, E., Elibol, M., Yang, Z., Paul, W., Jordan, M. I., and Stoica, I. (2018).  
463 Ray: A distributed framework for emerging AI applications. In 13th USENIX  
464 Symposium on Operating Systems Design and Implementation (OSDI), pages  
465 561–577.
- 466 [Paszke et al., 2019] Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J.,  
467 Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison,  
468 A., Köpf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy,  
469 S., Steiner, B., Fang, L., Bai, J., and Chintala, S. (2019). PyTorch: An  
470 imperative style, high-performance deep learning library. In Advances in  
471 Neural Information Processing Systems (NeurIPS).
- 472 [RAPIDS, 2023] RAPIDS (2023). RAPIDS cuML: GPU-accelerated machine  
473 learning. <https://github.com/rapidsai/cuml>.
- 474 [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and  
475 Klimov, O. (2017). Proximal policy optimization algorithms.
- 476 [Towers et al., 2024] Towers, M., Kwiatkowski, A., Terry, J., Balis, J. U.,  
477 De Cola, G., Deleu, T., Goulão, M., Kallinteris, A., Krimmel, M., KG,  
478 A., et al. (2024). Gymnasium: A standard interface for reinforcement learning  
479 environments.