

1 Explainability in Reinforcement Learning

2 Maxime Alaarabiou

3 Sunday 30th August, 2026

Chapter 1

Deep Learning

This chapter introduces the fundamental concepts of Deep Learning (DL), the training paradigm used to optimize deep Neural Network (NN). The objective is to provide the necessary foundation in DL for the ideas developed throughout the rest of this manuscript. Its structure is as follows:

Section 1.1: An overview of the historical developments that led to the emergence of the field.

Section 1.2: A formalization of the objective underlying DL algorithms.

Section 1.3: An examination of the structure of artificial NN.

Section 1.4: An explanation of how the learning process shapes NN.

Section 1.5: A discussion of how such models are evaluated.

Section 1.6: A summary of the concepts introduced in the chapter and a discussion of the theoretical limitations of DL.

1.1 Historical Developments

NN and DL emerged from an ambition to study the expressive capabilities of organic brain. The first major milestone came in 1943 with the proposal of a mathematical model of an artificial neuron, with a formulation of networked interactions [McCulloch and Pitts, 1943]. This work established the key idea that networks of simple units can give rise to complex computations. The focus was on representational power, not on the ability of such systems to learn. As a result, this work remained largely theoretical and had limited practical applications.

A major step toward operationalizing the idea of NNs was achieved with the perceptron algorithm in 1958 [Rosenblatt, 1958]. It was designed to adjust its parameters from input-output pairs in order to approximate a target function. The perceptron fundamental limitations were exposed in 1969, showing that it could not learn functions as simple as the exclusive-or (XOR) [Minsky and Papert, 1969]. At the same time, it was demonstrated that stacking multiple perceptrons could, in principle, represent such functions. However, no

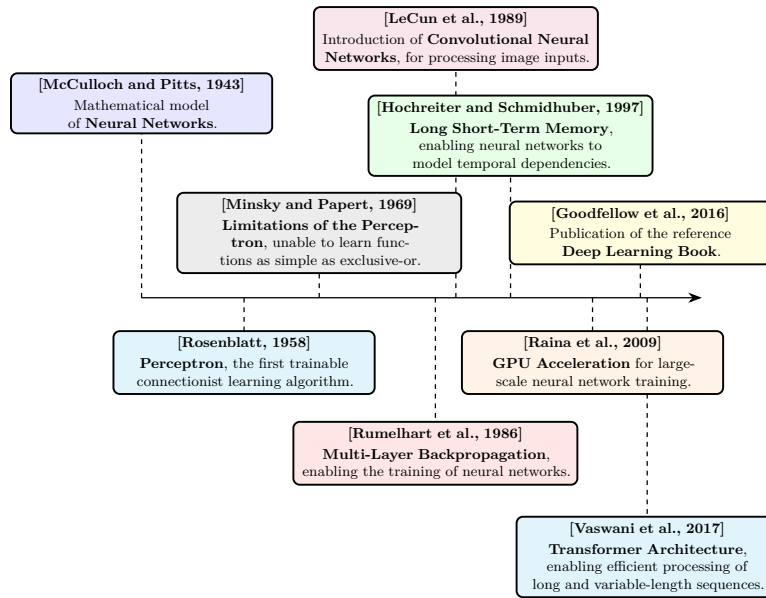


Figure 1.1: Evolution of deep learning.

effective training procedure for these multi-layer architectures was available at the time. It was only in 1986 that researchers introduced gradient descent as an effective method for training deep NN [Rumelhart et al., 1986]. From this point onward, the emergence of DL is retrospectively situated, even though the term itself appeared later. DL is defined as the training of deep NN using gradient-based optimization methods.

It took several decades before DL reached the level of prominence it has today. One major obstacle was the absence of theoretical guarantees regarding the convergence of learning algorithms. Although the universal approximation theorem shows that NN can represent a wide class of functions [Hornik et al., 1989], it did not ensure that such representations could be effectively learned through optimization methods like gradient descent. This lack of rigorous guarantees initially reduced interest within the academic community. As a result, early progress relied largely on empirical breakthroughs.

These empirical advances were made possible by two key developments: the availability of large-scale datasets and significant improvements in computational infrastructure. DL methods require vast amounts of data to perform well, and the rise of the internet alongside the digitization of information enabled the creation of such datasets [Hilbert and López, 2011]. The use of graphics processing units for large-scale training, demonstrated in 2009, greatly increased computational capacity and made it feasible to train large-scale models [Raina et al., 2009, Krizhevsky et al., 2012].

At the same time, researchers developed architectures specifically designed for different types of data, allowing DL to address an increasingly wide range of problems. Convolutional Neural Network (CNN) were introduced in 1989 to process images, using local, weight-shared filters that exploit the spatial structure of pixel grids [LeCun et al., 1989]. Long Short-Term Memory (LSTM)

networks were later proposed in 1997 to process temporal sequences, introducing gating mechanisms able to retain relevant information over long time horizons [Hochreiter and Schmidhuber, 1997]. More recently, the Transformer architecture, relying entirely on the attention mechanism, was introduced in 2017 to process long, variable-length sequences of strongly interconnected elements, allowing every element to attend directly to every other element regardless of their distance within the sequence [Vaswani et al., 2017]. This growing body of knowledge was consolidated in 2016 in a reference textbook that formalized DL as a coherent field of study [Goodfellow et al., 2016].

These changes paved the way for the rise of DL as it is known today. DL has achieved remarkable success across a wide range of domains, including computer vision [Krizhevsky et al., 2012], complex games [Silver et al., 2017a], content recommendation [Covington et al., 2016], bio-chemistry [Jumper et al., 2021], as well as the now essential digital assistants based on Large Language Model (LLM) [Brown et al., 2020]. The diversity of these applications explains the enthusiasm for this technology, which continues to redefine the boundaries of what was once considered automatable.

Figure 1.1 provides a timeline of the key milestones that shaped deep learning.

1.2 Objective

Now that the main developments that led to the rise of DL have been outlined, it is necessary to formalize what is actually being achieved when performing DL. To ensure clarity, the supervised learning setting will be considered, specifically applied to a regression task. This framework serves as a standard reference in the field, as all building blocks of DL applications are built upon this fundamental formalism.

In this setting, learning consists of progressively shaping a NN that serves as an approximation of a target function. At initialization, the network does not yet provide a meaningful representation of this function. Since a NN is a parametric function, its parameters are progressively updated during training so as to make it a better approximation of the target function. It gradually becomes capable of capturing the relationship between inputs and outputs through successive parameter updates during training. For this reason, DL is naturally situated within the framework of machine learning. A standard definition of machine learning is:

“A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”
[Mitchell, 1997]

100 When this formulation is adapted to DL, the following correspondences hold:

- 101 • The computer program corresponds to the learning algorithm, which is
102 responsible for adjusting the network's parameters.
- 103 • The task T consists of learning a NN that provides a good approximation
104 of the target function.
- 105 • The performance measure P corresponds to a function used to evaluate
106 the quality of this approximation.
- 107 • The experience E takes the form of a dataset composed of input–output
108 pairs derived from the function to be approximated.

109 The function to be approximated is denoted by f . As with any function, it
110 defines a mapping from an input to an output: $x \in \mathcal{X}$ denotes an input and
111 $y \in \mathcal{Y}$ the corresponding output. The function f is therefore characterized by an
112 input space $\mathcal{X}$ and an output space $\mathcal{Y}$, and can be formally written as $f : \mathcal{X} \rightarrow \mathcal{Y}$.
113 For simplicity of exposition, both $\mathcal{X}$ and $\mathcal{Y}$ are assumed to be vector spaces
114 throughout this manuscript. This assumption is adopted solely to facilitate the
115 presentation, as the framework can be readily extended to more general settings.
116 The notation $\hat{\cdot}$ indicates that a given object is an approximation of another.
117 Since the goal of the trained NN is to approximate f , this approximation is
118 denoted by $\hat{f} : \mathcal{X} \rightarrow \mathcal{Y}$. Given an input $x \in \mathcal{X}$, the output of the approximating
119 function $\hat{f}(x)$ is denoted by $\hat{y}$.

120 It is important to emphasize that the target function f is generally not known
121 explicitly. If it were, there would be no need to learn an approximation, since
122 the true function would already provide the exact mapping. In practice, only a
123 finite set of observations generated by f is available, organized in what is called
124 a dataset. To distinguish between the different samples, an index i is introduced,
125 allowing input–output pairs of the form $(x^{(i)}, y^{(i)})$ to be defined. This leads to
126 the following formulation for a dataset $\mathcal{D}$ of size N , where N denotes the number
127 of available examples:

$$\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N. \quad (1.1)$$

128 Now that the way the target function f is represented has been made explicit,
129 one can ask what it means for a function $\hat{f}$ to be a good approximation of this
130 function. It means that, for a given input $x^{(i)}$, the NN output $\hat{y}^{(i)}$ is close to the
131 true output $y^{(i)}$. To quantify the quality of this approximation, a function $\mathcal{L}$ is
132 introduced that measures the discrepancy between predictions and target values.
133 In DL, this function called the loss function plays a central role as it guides the
134 learning process. In regression settings, a common choice is the Mean Squared
135 Error (MSE), defined as:

$$\mathcal{L}(\hat{y}, y) = (\hat{y} - y)^2. \quad (1.2)$$

136 Thanks to this loss, it is now possible to rigorously define the performance
137 measure P . This performance measure corresponds to the average loss over the
138 dataset and can therefore be written as:

$$P = \frac{1}{N} \sum_{i=1}^N \mathcal{L}(\hat{f}(x^{(i)}), y^{(i)}). \quad (1.3)$$

139 This formulation captures the mathematical core of the learning objective,
140 namely finding a NN $\hat{f}$ that minimizes the prediction error over the dataset $\mathcal{D}$.

Since the target function f may take many different forms, $\hat{f}$ must be expressive enough to approximate a broad range of such functions. NN architectures are well suited to this requirement, as they can represent a wide range of functional mappings, whose composition is detailed in the following section.

1.3 Neural Network Structure

NN are called connectionist models. Rather than modeling information processing through a set of explicit rules, they rely on a flow of signals distributed within a computational structure. Their architecture is composed of a large number of simple computational units, which operate on partial representations and exchange signals across the network. A complex pattern of weighted connections between these units enables the transformation and propagation of information.

To simplify the presentation, the focus is placed on the MultiLayer Perceptron (MLP) architecture. Although more specialized architectures exist, the MLP forms a fundamental building block of DL. All the essential mechanisms of DL are present in it, making it a particularly suitable framework for introducing key concepts.

This section therefore begins by examining the artificial neuron, the fundamental computational unit underlying the MLP. It then explains how these neurons are organized into layers. Finally, it concludes with a general description of the neural network as a whole.

1.3.1 Neurons

The artificial neuron is the basic building block of NN. It is through these units that all computations within the network are carried out. A neuron receives a set of input signals from other neurons, where these inputs are represented as real-valued numbers. It produces an output, also a real-valued number, which is transmitted to subsequent neural units. Its internal state, called the activation, reflects its level of response and encodes the information processed for a given input. This activation, or lack thereof, is then used as input information by other neurons to determine their own activations, thereby propagating information through the network.

A neuron is characterized by:

- a set of incoming connections from other neurons,
- a weight associated with each of these connections,
- a bias term,
- a differentiable nonlinear activation function,
- a scalar activation state.

The activation of a neuron is determined by a weighted combination of its input signals, to which a bias is added, followed by the application of a nonlinear activation function. Mathematically, for a neuron receiving inputs $(a_1, \dots, a_n)$, its activation a is given by:

$$a = \sigma \left(\sum_{i=1}^n w_i a_i + b \right), \quad (1.4)$$




Figure 1.2: Artificiale neurone.

181 where:

- 182 • w_i is the weight associated with the i -th incoming connection,
- 183 • b is the bias,
- 184 • σ is the activation function.

185 Figure 1.2 shows a schematic view of this basic unit. The set of weights
 186 $w_1, \dots, w_i$ with the bias b constitute the parameters of the neuron, and they are
 187 the elements that evolve during training. The magnitude of a weight w (relative
 188 to the others) indicates how strongly a neuron depends on the activation of
 189 the neurons to which it is connected. A large positive weight means that a
 190 high activation in the connected neuron tends to increase the activation of the
 191 receiving neuron, whereas a negative weight implies the opposite effect. In this
 192 way, the weights modulate the strength and nature of the connections between
 193 neurons, thereby defining the overall pattern of interaction within the network.

194 The activation function σ models how a neuron transforms incoming signals
 195 into its output activation. Originally, this function was inspired by biological
 196 neurons, with the goal of mimicking the firing behavior of organic cells. However,
 197 modern DL no longer relies on this biological interpretation. In practice, any
 198 function that is nonlinear, differentiable, and computationally efficient can be
 199 used. Nonlinearity is a crucial property, since without it a NN would reduce
 200 to a composition of linear transformations, and would therefore be unable to
 201 approximate complex functions.

202 1.3.2 Layers

203 In MLP architectures, neurons are organized into successive layers, as illustrated
 204 in Figure 1.3. This layered structure determines how neurons are connected. In
 205 this setting, each neuron receives inputs from all neurons in the preceding layer




Figure 1.3: Multilayer Perceptron Architecture.

and sends its output to all neurons in the following layer. As a result, neurons within the same layer do not interact with each other.

Three types of layers are distinguished:

- an input layer, which directly receives input,
- one or more hidden layers (also referred to as intermediate layers), which perform intermediate transformations,
- an output layer, which produces the final prediction.

The input and output layers play a specific role. The input layer has no preceding layer: its neurons do not compute activations but instead directly encode the components of the input vector x . Conversely, the output layer has no subsequent layer, and its neurons produce the model's prediction $\hat{y}$. This structure imposes a direct correspondence between the dimensionality of the input space and the number of neurons in the input layer, and similarly between the output space and the number of neurons in the output layer.

When describing NN, especially in MLP, it is more common to reason at the level of layers rather than individual neurons. This perspective allows interactions between layers to be modeled in a matrix form. Such a representation enables efficient parallel computation and forms the basis of modern DL implementations [Krizhevsky et al., 2012]. Within this framework, it is useful to introduce the notion of layer activations. The activation of a layer corresponds to the activations of all its neurons, organized into a vector. For a layer l composed of n_l neurons, the activations are given by:

$$\mathbf{h}^{(l)} = [a_1^{(l)}, \dots, a_{n_l}^{(l)}]. \quad (1.5)$$

Using this formulation, the interaction between two consecutive layers can be written compactly as:

$$\mathbf{h}^{(l)} = \sigma^{(l)} \left(\mathbf{W}^{(l)} \mathbf{h}^{(l-1)} + \mathbf{b}^{(l)} \right), \quad (1.6)$$




Figure 1.4: Interaction between layers.

where:

- $\mathbf{h}^{(l-1)}$ represents the activations vector of the previous layer,
- $\mathbf{W}^{(l)} \in \mathbb{R}^{n_l \times n_{l-1}}$ is the weight matrix connecting layer $l - 1$ to layer l ,
- $\mathbf{b}^{(l)} \in \mathbb{R}^{n_l}$ is the bias vector associated with layer l ,
- $\sigma^{(l)}$ is the activation function applied element-wise to the neurons of layer l .

This matrix-vector representation is visualised in Figure 1.4, which highlights how each neuron in layer l receives weighted signals from all neurons in the previous layer.

It is important to examine the role played by the intermediate layers of a NN. At least one hidden layer is required to represent sufficiently complex functions, and in practice, modern architectures often include many more. Theoretical results have shown that, for a fixed number of parameters, increasing the number of layers enhances the expressive capacity of NN [Telgarsky, 2016]. A common interpretation is that each layer acts as a processing step, so that more complex target functions require more such steps to be well approximated, although this interpretation remains grounded in empirical observation rather than formal theoretical justification [Goodfellow et al., 2016].

1.3.3 Network

With a layer-wise representation, the propagation of information can be described as a sequence of transformations applied from the input layer (indexed 1) to the output layer (indexed L):

$$\hat{f}(x) = \mathbf{h}^{(L)} = \sigma^{(L)} \left(\mathbf{W}^{(L)} \sigma^{(L-1)} \left(\dots \sigma^{(1)} \left(\mathbf{W}^{(1)} x + \mathbf{b}^{(1)} \right) \dots \right) + \mathbf{b}^{(L)} \right) = \hat{y}. \quad (1.7)$$

For a more compact description of information propagation, the layer activations are defined recursively as:

$$\mathbf{h}^{(l)} = \begin{cases} x & \text{if } l = 1, \\ \sigma^{(l)} \left(\mathbf{W}^{(l)} \mathbf{h}^{(l-1)} + \mathbf{b}^{(l)} \right) & \text{if } 2 \leq l \leq L. \end{cases} \quad (1.8)$$

To simplify notation, it is convenient to group all parameters of the network into a single set. The model parameters are thus denoted by:

$$\theta = \{\mathbf{W}^{(l)}, \mathbf{b}^{(l)}\}_{l=1}^L. \quad (1.9)$$

In this manuscript, θ is a vector in $\mathbb{R}^P$, where P denotes the total number of parameters, such that $\theta = [w_1, \dots, w_P]$. The notation $\hat{f}_\theta$ expresses the dependence of the NN on its parameters. Increasing the number of parameters, either by adding more neurons per layer or by increasing the number of layers, can enhance the expressive power of the model, allowing it to approximate more complex functions. However, this also increases computational cost and training time.

It is important to distinguish between the architecture and the parameters of a NN. When referring to the architecture of $\hat{f}$, the number of layers, the number of neurons per layer, and the choice of activation functions are considered. The parameters correspond to θ , which controls the strength of the connections between neurons. The following sections examine how these parameters can be adjusted to obtain a better approximation of the target function.

1.4 Learning

Now that the structure of a NN has been detailed, the next step is to examine how its parameters are adjusted to accomplish the assigned task. This brings the discussion to the core of DL, namely the learning process itself. As a reminder, the objective of training a NN is to adjust its parameters so that the network provides a good approximation of the target function. Mathematically, this objective is expressed as:

$$\min_{\theta} \frac{1}{N} \sum_{i=1}^N \mathcal{L} \left(\hat{f}_\theta(x^{(i)}), y^{(i)} \right). \quad (1.10)$$

For conciseness, the per-example loss is denoted $\ell^{(i)}(\theta) \stackrel{\text{def}}{=} \mathcal{L} \left(\hat{f}_\theta(x^{(i)}), y^{(i)} \right)$ throughout the remainder of this section, so that the training objective can equivalently be written as $\min_{\theta} \frac{1}{N} \sum_{i=1}^N \ell^{(i)}(\theta)$.

The NN $\hat{f}$ used to approximate the target function f is highly complex. It involves a large number of parameters as well as many nonlinear components. As a result, the optimum cannot be determined analytically by directly studying the properties of the function. Instead, approximate optimization methods are required to find suitable parameter values for NN. The class of algorithms used for this optimization is known as gradient descent methods. Numerous variants of gradient descent have been developed for NN training, each designed to improve specific aspects of the optimization process. For simplicity, the focus here is on Stochastic Gradient Descent (SGD) with a batch size of 1. As in previous cases, this choice is made because this setting captures the essential intuition behind the training method.

This section is organized into three steps. The first derives the gradient of the loss using the chain rule. The second shows how this gradient is used to update the network parameters. The third discusses the convergence properties of the resulting algorithm.

1.4.1 Computing the Gradient

The first step of SGD is to randomly initialize the parameter vector $\theta \in \mathbb{R}^P$ for a given network architecture $\hat{f}$. This initial state is denoted by $\theta^{(1)}$. The goal is to iteratively update this vector so that the network gets better at approximating the target function. The effect of each parameter w_k on the loss $\ell^{(i)}(\theta)$ is quantified by the corresponding partial derivative $\frac{\partial \ell^{(i)}(\theta)}{\partial w_k}$.

However, w_k can belong to any layer of the network. Computing this derivative therefore requires propagating the effect of a change in w_k through all the layers that follow it, up to the loss. The key idea is to view each layer as a function, denoted $\phi^{(l)}$ for layer l , that takes as input the activation vector $\mathbf{h}^{(l-1)} \in \mathbb{R}^{n_{l-1}}$ produced by the previous layer and returns as output the activation vector $\mathbf{h}^{(l)} \in \mathbb{R}^{n_l}$, where n_{l-1} and n_l denote the number of neurons in layers $l-1$ and l , respectively. Layer l is therefore the function:

$$\phi^{(l)}: \mathbb{R}^{n_{l-1}} \rightarrow \mathbb{R}^{n_l}, \quad \mathbf{h}^{(l)} = \phi^{(l)}(\mathbf{h}^{(l-1)}). \quad (1.11)$$

With this view, the whole network is the composition of its L layer functions, $\hat{f}_\theta = \phi^{(L)} \circ \dots \circ \phi^{(1)}$. Each layer function $\phi^{(l)}$ is itself parameterized by the weights belonging to layer l , so it can equally be viewed as a function of these parameters. This is what makes it possible, later on, to differentiate $\phi^{(l)}$ with respect to w_k .

Because each layer function has several outputs that jointly depend on several inputs, the ordinary derivative is no longer sufficient to describe how it varies, and the Jacobian matrix is used instead as its generalization. The Jacobian matrix quantifies how a small change in each input component affects each output component. Consider layer l 's function $\phi^{(l)}: \mathbb{R}^{n_{l-1}} \rightarrow \mathbb{R}^{n_l}$: its Jacobian $J_{\phi^{(l)}}(\mathbf{h}^{(l-1)})$, evaluated at $\mathbf{h}^{(l-1)}$, collects the partial derivatives of every output component with respect to every input component:

$$J_{\phi^{(l)}}(\mathbf{h}^{(l-1)}) = \begin{bmatrix} \frac{\partial h_1^{(l)}}{\partial h_1^{(l-1)}} & \dots & \frac{\partial h_1^{(l)}}{\partial h_{n_{l-1}}^{(l-1)}} \\ \vdots & \ddots & \vdots \\ \frac{\partial h_{n_l}^{(l)}}{\partial h_1^{(l-1)}} & \dots & \frac{\partial h_{n_l}^{(l)}}{\partial h_{n_{l-1}}^{(l-1)}} \end{bmatrix}. \quad (1.12)$$

Consider now two consecutive layers, $\phi^{(l)}$ followed by $\phi^{(l+1)}$, so that $\mathbf{h}^{(l+1)} = \phi^{(l+1)}(\phi^{(l)}(\mathbf{h}^{(l-1)}))$. The chain rule states that the Jacobian of this composition, evaluated at $\mathbf{h}^{(l-1)}$, is the product of the Jacobians of each layer, each evaluated at the appropriate point:

$$J_{\phi^{(l+1)} \circ \phi^{(l)}}(\mathbf{h}^{(l-1)}) = J_{\phi^{(l+1)}}(\mathbf{h}^{(l)}) J_{\phi^{(l)}}(\mathbf{h}^{(l-1)}). \quad (1.13)$$

Since the network is the composition of its layer functions $\phi^{(1)}, \dots, \phi^{(L)}$, with each layer depending only on the one before it, a parameter w_k belonging to layer l affects the loss $\ell^{(i)}(\theta)$ through a chain of dependencies: it first changes the output $\mathbf{h}^{(l)}$ of its own layer, which changes $\mathbf{h}^{(l+1)}$, and so on until the output layer $\mathbf{h}^{(L)}$, from which the loss is computed. This layer-by-layer propagation is referred to as backpropagation:

$$\frac{\partial \ell^{(i)}(\theta)}{\partial w_k} = J_{\mathcal{L}}(\mathbf{h}^{(L)}, y^{(i)}) J_{\phi^{(L)}}(\mathbf{h}^{(L-1)}) \dots J_{\phi^{(l+1)}}(\mathbf{h}^{(l)}) J_{\phi^{(l)}}(w_k), \quad (1.14)$$

where:

- $J_{\mathcal{L}}(\mathbf{h}^{(L)}, y^{(i)})$ denotes the Jacobian of $\mathcal{L}$ with respect to its first argument, evaluated at $(\mathbf{h}^{(L)}, y^{(i)})$,
- each subsequent $J_{\phi^{(j+1)}}(\mathbf{h}^{(j)})$ denotes the Jacobian of the function $\phi^{(j+1)}$ with respect to its input $\mathbf{h}^{(j)}$,
- $J_{\phi^{(l)}}(w_k)$ denotes the Jacobian of the function $\phi^{(l)}$ with respect to the parameter w_k .

Multiplying these local Jacobians together therefore propagates the effect of w_k , layer by layer, all the way to the loss.

Once $\frac{\partial \ell^{(i)}(\theta)}{\partial w_k}$ can be computed for every parameter, all of these values are stacked into one vector, called the gradient:

$$\nabla_{\theta} \ell^{(i)}(\theta) = \left[\frac{\partial \ell^{(i)}(\theta)}{\partial w_1}, \dots, \frac{\partial \ell^{(i)}(\theta)}{\partial w_P} \right]. \quad (1.15)$$

Each component answers the same question for a different parameter, namely, if this parameter is slightly increased, does the loss go up or down, and by how much. A positive value means increasing the parameter increases the loss, whereas a negative value means the opposite. The gradient thus indicates, for every parameter, which direction to move in order to reduce the loss.

1.4.2 Update rule

Now that information is available on how changes in the parameters affect the loss, the objective is to minimize it. To do so, the parameters are updated in the direction opposite to the gradient. Indeed, the gradient is followed in the opposite direction because it indicates the direction in which the loss increases, whereas the objective is to reduce it. This leads to the gradient descent update rule:

$$\theta^{(n+1)} = \theta^{(n)} - \eta \nabla_{\theta^{(n)}} \ell^{(i)}(\theta^{(n)}), \quad (1.16)$$

where $\eta > 0$ is the learning rate, which controls the step size.

Algorithm 1: Stochastic Gradient Descent**Input:**Dataset $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$ Architecture $\hat{f}$ Loss function $\mathcal{L}$ Learning rate $\eta > 0$ Initial parameters θ (optional)**Output:**Trained parameters θ

```

1 if  $\theta$  is given then
2    $\theta^{(1)} \leftarrow \theta$ ;
3 else
4   Initialize  $\theta^{(1)}$  randomly;
5  $n \leftarrow 1$ ;
6 while stopping criterion not satisfied do
7   Select a random example  $(x^{(i)}, y^{(i)}) \in \mathcal{D}$ ;
8   Model prediction for the selected example
    $\hat{y}^{(i)} \leftarrow \hat{f}_{\theta^{(n)}}(x^{(i)})$ ;
9   Compute the error between the prediction and the ground truth
    $\ell \leftarrow \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$ ;
10  Compute the gradient of the error since the loss function is differentiable
    $g \leftarrow \nabla_{\theta^{(n)}} \ell$ ;
11  Error minimization via a gradient descent step
    $\theta^{(n+1)} \leftarrow \theta^{(n)} - \eta g$ ;
12   $n \leftarrow n + 1$ ;
13 Return  $\theta^{(n)}$ 

```

353 This parameter η must remain small, as the gradient only provides reliable
354 information in a local neighborhood of the current parameters. As a result,
355 a single step is not sufficient to significantly reduce the loss, so the process is
356 iterated many times. Producing a sequence of parameter vectors $\theta_1, \theta_2, \dots$ that
357 progressively improve the model. In practice, training is stopped when successive
358 updates no longer produce a significant decrease in the loss. Algorithm 1
359 summarizes this whole training procedure.

360 To develop an intuition for this update rule, Figure 1.5 illustrates it applied to
361 the simple convex loss function $\mathcal{L}(w_1, w_2) = w_1^2 + w_2^2$, where the red path shows
362 the parameter updates converging to the minimum. The point θ_1 represents the
363 initial parameter values, with w_1 and w_2 randomly initialized, and the subsequent
364 points, from θ_2 to θ_6 , represent the successive updates produced by applying
365 Equation 1.16, progressively moving the trajectory toward the minimum of the
366 loss function.




Figure 1.5: Gradient descent in a case $\mathcal{L}(w_1, w_2) = w_1^2 + w_2^2$.

1.4.3 Convergence Properties

Having established how SGD updates the parameters, a natural question is whether this procedure is guaranteed to converge. It has been mathematically shown that gradient descent applied to a convex loss with respect to the parameters converges to a global optimum, independently of the initialization. However, NN training involves highly non-convex loss surfaces, so the optimization trajectory becomes strongly dependent on the initial conditions, and no such guarantee holds. Figure 1.6 illustrates the effect of non-convex functions, different starting points lead to distinct local minima.

Despite the lack of theoretical guarantees of convergence to a global optimum, the simple procedure of random initialization followed by iterative gradient-based updates is sufficient in practice to train complex NN.

1.5 Evaluation

Now that the process of training a model has been described, the question of how to evaluate it must be addressed. In the supervised learning setting, this evaluation relies on two aspects:

- its performance on the data it was trained on,
- its performance on data it was not trained on.

The first aspect measures how well the loss minimization process has succeeded. The second measures how well the model generalizes. Indeed, the dataset $\mathcal{D}$ represents a small subset of all possible inputs. To estimate how the model will perform in general, it is important to evaluate how it performs on data it has not seen during training.




Figure 1.6: Comparison of loss surfaces and gradient descent trajectories.

In order to quantify these two aspects, the dataset is split into two subsets:

- $\mathcal{D}_{\text{train}}$ is the set on which the model minimizes its loss during the training.
- $\mathcal{D}_{\text{test}}$ is kept aside during training in order to evaluate the model's ability to generalize.

When training is completed, the first step is to verify that the model performs well on the data it was trained on. To do so, the average loss on the training set, denoted $\bar{\ell}_{\text{train}}$, is studied. This is an instance of the performance measure P defined in Equation (1.3), restricted to $\mathcal{D}_{\text{train}}$:

$$\bar{\ell}_{\text{train}} = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x^{(i)}, y^{(i)}) \in \mathcal{D}_{\text{train}}} \mathcal{L}(\hat{f}_{\theta}(x^{(i)}), y^{(i)}). \quad (1.17)$$

If $\bar{\ell}_{\text{train}}$ is not sufficiently low, this means that the model has not been able to learn the task. It is not necessary to go further in the analysis of the model's quality. Before restarting the training process, various design choices are then adjusted. Based on what has been discussed so far in this chapter, these include modifications to the number of layers, the number of neurons per layer, activation functions, the learning rate. However, in practice, there are many other possible

variations. There is no universal rule for choosing the appropriate value for each of them. They are therefore chosen empirically, based on models that perform well on similar tasks, combined with intuition. Their selection is often costly and remains a challenging problem in DL.

Assuming now that the model achieves satisfactory performance on the training dataset $\mathcal{D}_{\text{train}}$. Its ability to generalize must now be studied. To do so, the corresponding average loss on $\mathcal{D}_{\text{test}}$ is computed:

$$\bar{\ell}_{\text{test}} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sum_{(x^{(i)}, y^{(i)}) \in \mathcal{D}_{\text{test}}} \mathcal{L}(\hat{f}_{\theta}(x^{(i)}), y^{(i)}). \quad (1.18)$$

If $\bar{\ell}_{\text{test}} \gg \bar{\ell}_{\text{train}}$, the model fails to generalize. This situation most often arises due to a phenomenon known as overfitting. In such cases, the model does not learn general patterns but instead memorizes the training examples. This typically occurs when the number of parameters is too large relative to the amount of training data. The simplest remedy is therefore to reduce the number of parameters in the model. However, it is still necessary to ensure that the model remains sufficiently expressive to perform well on the data.

If $\bar{\ell}_{\text{test}} \approx \bar{\ell}_{\text{train}}$, this indicates that the training procedure has successfully produced a NN that generalizes well beyond the training data [Kawaguchi et al., 2022]. The resulting model can therefore be considered suitable for the task for which it was trained, marking the end of the NN training process.

1.6 Conclusion

This chapter has introduced the fundamental concepts underlying DL, including the general objective of NN models, their structure, the training procedure, and their evaluation. For clarity of presentation, several simplifying assumptions have been made throughout this chapter. The problem has been considered in a regression setting, using a MLP trained with SGD and a batch size of one. These choices were made deliberately to expose the core mechanisms of DL as clearly as possible.

Despite these simplifications, the framework introduced here captures mechanisms that are shared across a broad range of DL methods: A sequence of parameter updates, each based on information about the loss around the current parameters. The apparent simplicity of this learning procedure therefore contrasts with the complexity of the functions that NN can ultimately represent. Understanding how such complex behavior emerges from a relatively simple optimization process remains an important challenge in the study of DL.

Mathematical tools allow us to describe the functions represented by NN and the procedures used to train them. However, they do not yet fully explain their remarkable empirical performance. Understanding the origin of this success therefore requires looking beyond the optimization procedure itself.

One common feature across NN architectures is the presence of intermediate representations. Whether considering a MLP, a CNN, an LSTM, or an architecture based on attention, the input is progressively transformed into intermediate representations before producing the final output. Their presence provides a common perspective for understanding how NNs transform information. The collection of these intermediate representations is referred to as the Latent Space

447 (LS). Studying this LS can provide insights into how NNs learn and generalize,
448 which motivates the following chapter.

Chapter 2

Latent Space Analysis

In classical machine learning approaches, practitioners had to manually design data representations in order to make them usable by simple models. Deep Learning (DL) introduces a paradigm shift: rather than relying on hand-crafted features to meaningfully represent data, it automatically learns these representations. Through the successive activation of hidden layers, Neural Network (NN) gradually construct representations that become increasingly relevant to the task under consideration [Rumelhart et al., 1986].

These representations are not meaningful at initialization, they emerge progressively during training as the model optimizes its objective. It is through these learned representations that a trained network is able to map an input to an output, including for examples never encountered during training. Since these representations are learned internally by the network and are not directly observable from the input or output spaces, they are commonly referred to as Latent Representation (LR). The collection of these LR learned throughout the network is referred to as a Latent Space (LS).

These LS are now at the core of a growing number of practical applications. Neural video compression exploits the compactness of LS to encode information efficiently [Lu et al., 2019]. In Reinforcement Learning (RL), agent internal representations are used for efficient exploration in extremely open environments [Burda et al., 2018]. Recommendation systems model users and content within a shared LS to compute meaningful similarities [He et al., 2017]. Retrieval-augmented generation systems index and query knowledge bases through latent embeddings [Lewis et al., 2021]. Finally, Large Language Model (LLM) and multimodal architectures rely on the quality of learned representations to align heterogeneous modalities such as text, images, and audio [Radford et al., 2021]. These examples represent only some of the most widespread applications of LR.

Given their extensive use, the study of LR has become a central topic in DL research. A major objective is to understand how information is organized within these spaces, which concepts are encoded, and how these representations

influence model behavior.

- Section 2.1 introduces the working hypothesis that underlies much of the research on LS analysis, which we refer to as the *Concept Hypothesis* (Hyp. 1).
- Sections 2.2, 2.3, 2.4, and 2.5 examine different approaches to LS analysis, respectively covering probing methods, semantic directions, sparse autoencoders, and clustering-based methods.
- Section 2.6 concludes with a comparative synthesis of these methods and a discussion of the open questions that remain.

2.1 Concept Hypothesis

In this section, we introduce the main working hypothesis underlying much of the literature on LS analysis. Although rarely stated explicitly, this hypothesis is implicitly present in the recurring use of notions such as abstraction, representation, LS, concept, and semantics. Throughout this manuscript, we refer to this hypothesis as *Concept Hypothesis* (Hyp. 1). The goal of this section is to make this conceptual framework explicit by defining its main components.

The *Concept Hypothesis* (Hyp. 1) originates from the need to explain the generalization ability of deep NN observed after training, as discussed in Section 1.5. A model is said to generalize when, for an input $x \in \mathcal{X}$ not seen during training, it still produces an output $\hat{y}$ close to the true output y . Such behavior cannot be explained by simple memorization of training samples. Instead, the model must extract and exploit underlying regularities in the target function. This requires a process of abstraction, as defined in Definition 2.1.1, whereby raw inputs are transformed into internal structures that are meaningful for the task [Bengio et al., 2014].

Definition 2.1.1: Abstraction

Abstraction is the process by which a model transforms an input into a task-relevant internal form. It captures and restructures information in order to represent relationships that are useful for prediction.

The ability of NN to perform abstraction arises from their intermediate activations. A NN does not directly map x to y , it computes a sequence of intermediate activations across hidden layers. These activations are numerical encodings of the input referred to as embeddings. Once training is complete, these embeddings are structured to support the performance of a specific task. To mark this distinction, we will refer to these trained embeddings as LR, as defined in Definition 2.1.2.

Definition 2.1.2: Latent Representation

A representation is the internal encoding of an input given by a intermediate layer of a trained NN.

Each hidden layer refines the representation of the previous layer. Each transformation progressively reshapes the representation into a form more directly




Figure 2.1: The same latent space activations seen without (a) and under (b) the *Concept Hypothesis* (Hyp. 1).

aligned with the output space [Rumelhart et al., 1986]. This process can be described as a sequence of transformations where each $h^{(l)}$ is a representation:

$$x \rightarrow h^{(1)} \rightarrow h^{(2)} \rightarrow \dots \rightarrow h^{(L-1)} \rightarrow \hat{y}$$

The set of all such representations produced by a layer over a dataset defines its LS, as defined in Definition 2.1.3.

Definition 2.1.3: Latent Space

The LS is the set of all representations produced by a layer over a dataset. It is shaped by learning and organizes representations according to task-relevant properties.

Under the *Concept Hypothesis* (Hyp. 1), the analysis of the LS is assumed to provide access to abstract structures learned by the model, as illustrated in Figure 2.1. The same set of latent activations can be viewed either as an unstructured point cloud or, under this hypothesis, as a collection of points organized by concept membership, in which case the structure of the LS becomes meaningful and amenable to analysis. In Figure 2.1, the colors represent hypothetical concept memberships assigned by the model rather than observed labels. These structures are referred to as concepts. In this manuscript, a concept is defined as an abstract representation that organizes information and supports reasoning (see Definition 2.1.4).

Definition 2.1.4: Concept

A concept is a task-dependent abstraction that captures regularities in the input space relevant for predicting the target output.

Our definition of concepts is consistent with both the pragmatic tradition in philosophy and the information-processing view of cognition. Within these frameworks, concepts arise from the need to simplify an otherwise intractable reality. Because the environment contains more information than any cognitive system can process, intelligent behavior relies on abstractions, called concepts, that discard irrelevant details while preserving the information required to achieve

a given objective [Peirce, 1878, Misak, 2013]. These concepts are realized as internal representations that reduce complexity and enable reasoning under limited computational resources [Newell, 1980, Simon, 1996]. Although these theories were originally developed to explain human cognition, the *Concept Hypothesis* (Hyp. 1) assumes that the same principles can be extended to deep NN.

In order to influence the model’s computations, concepts must necessarily admit a concrete realization within the LS. We refer to this realization as the Concept Structure (CS). The definition of the CS is provided in Definition 2.1.5.

Definition 2.1.5: Concept Structure

A CS is the specific form for how a concept is instantiated within the LS.

The methods reviewed in this chapter all build on the *Concept Hypothesis* (Hyp. 1) formalized in Hypothesis 1, but differ in how they characterize the CS of a concept. This choice directly determines how concepts can be identified and extracted from the LS. As a result, several families of LS analysis methods have emerged in the literature, each adopting a different view of how concepts are structured within the LS.

Hypothesis 1: Concept Hypothesis

NN generalize by performing abstraction. Through training, their LS becomes organized around concepts. Concepts are task-dependent abstractions that capture regularities relevant to the task. Analyzing the LS therefore provides access to the CS through which these concepts are instantiated in the model.

Now that the definition of a *Concept Hypothesis* (Hyp. 1) has been established, we can consider its relationship to human. According to our definition, we can imagine concepts that are useful for the machine but do not make sense to humans. This relationship is captured by the notion of semantics [Marconato et al., 2023], defined in Definition 2.1.6. Some concepts developed in LS can be readily associated with human concepts and are said to have high semantic value. Others may contribute to task performance while remaining difficult for humans to understand, and are said to have low semantic value.

Definition 2.1.6: Semantics

Semantics refers to the alignment between concepts in the model LS and concepts used by a human observer. A concept is semantic if it can be associated with a human interpretable concept.

It is important to note that this conceptual framework is not supported by a complete theoretical understanding of deep NN. Current mathematical tools do not provide a rigorous characterization of their internal representations and do not formally prove that concepts emerge as identifiable CS within LS. Nevertheless, this perspective constitutes a widely adopted working hypothesis in the field of representation and LS analysis.

The remainder of this chapter is devoted to reviewing methods for analyzing LS and extracting the concepts they encode. We begin with probing, one of the

most direct approaches for assessing whether concepts are represented in the LS.

2.2 Probing

Probing methods aim to determine whether a given semantic concept is encoded within the LS of a NN [Li et al., 2024, Karvonen, 2024]. More specifically, they investigate whether the presence or absence of a selected concept can be inferred from embeddings. The assumption is that if a concept can be reliably decoded from these representations, then the network likely encodes information related to it in order to perform its task.

A probing analysis begins by identifying a set of semantic concepts that the network could potentially use to accomplish its task. Each example in the dataset is then annotated to indicate the presence or absence of the studied concepts. We denote by c_j the presence or absence of concept j in example x . The dataset can therefore be written as follows:

$$\mathcal{D} = \{(x^{(i)}, y^{(i)}, c_1^{(i)}, \dots, c_n^{(i)})\}_{i=1}^N.$$

The data are subsequently projected into the LS of the pretrained model. We denote by h^l the representation vector obtained at layer l . From these LR, a classifier is trained in a supervised manner to predict the presence of the considered concept. This classifier is referred to as a probe and is denoted by p . Its purpose is to determine whether a given concept is present in the representation. Its purpose is therefore to assess whether information associated with the concept is accessible from the network’s internal representations. This can be formalized as follows:

$$p_{\theta}^{l,j} : h^l \mapsto c_j, \quad \theta^* = \arg \min_{\theta} \mathcal{L}(p_{\theta}^{l,j}(h^l), c_j)$$

The classifier used to detect the presence or absence of a concept must remain intentionally simple. Indeed, if the probe has too much capacity, it may reconstruct the target information from complex correlations that do not necessarily correspond to information explicitly encoded in the LS. In such a case, it becomes difficult to determine whether the concept is genuinely represented by the network or merely reconstructed through the expressive power of the classifier. For this reason, probing methods often rely on linear classifiers.

A linear separation suggests that the information is easily accessible to the NN and potentially usable during decision-making. If the classifier achieves performance significantly above chance level, this suggests that the LS contains exploitable information related to the considered concept. Figure 2.2 illustrates the decision boundaries of such linear probes for three concepts, where each colored region indicates the area classified positively by the corresponding probe.

An application of this technique can be found in studies investigating whether LLM develop internal representations of their environment [Li et al., 2024, Karvonen, 2024]. In these works, researchers train LLM architecture to play games such as chess or Othello in a purely textual setting. The model is provided only with sequences of moves and information about the actions it can take, all expressed in textual form. The central question is whether the LLM develops an internal representation of




Figure 2.2: Decision boundaries of the linear probes for concepts A, B, and C.

the game state. To address this question, the probed concepts correspond to the presence or absence of different game pieces on specific board positions.

As illustrated in Figure 2.3, probes are able to reconstruct the full board state from the model’s hidden representations, even though the only available information to the model is a sequence of text. In this figure, each row corresponds to a piece type, comparing the ground truth with the probe’s predictions and confidence. The close agreement between the ground truth and predicted positions suggests that the model develops a structured representation of the underlying environment from which the data are generated.

However, the fact that a concept can be decoded from LR does not necessarily imply that the model actually uses this information to perform its task. To assess whether representations play a causal role in the model’s output, it is possible to perform interventions on the LR. An intervention consists of modifying a LR in a controlled manner by activating or deactivating a given concept, and then studying its impact on the model’s output. If the resulting change in the model’s output is consistent with the applied perturbation, this suggests that the representation had a causal role.

To this end, the LR h^l is minimally modified such that the probe prediction for the target concept c_j changes state, while the predictions associated with other concepts remain unchanged. This goal can be formulated as an optimization problem over the latent activation itself. The optimization is performed directly on the representation $h^{l'}$, using Stochastic Gradient Descent (SGD). The solution to this problem is the optimal modified representation h^{l*} , defined as:

$$h^{l*} = \arg \min_{h^{l'}} \underbrace{\gamma \mathcal{L}_{\text{flip}}(p_{\theta}^{l,j}(h^{l'}), c_j)}_{\text{flip target concept } c_j} + \underbrace{\lambda \sum_{k \neq j} \mathcal{L}_{\text{preserve}}(p_{\theta}^{l,k}(h^{l'}), c_k)}_{\text{preserve other concepts } c_{k \neq j}} + \underbrace{\beta \|h^{l'} - h^l\|_2^2}_{\text{minimality}}$$




Figure 2.3: Probing analysis of a large language model trained to play chess [Karvonen, 2024].

where:

- h^l is the original LR of the input at layer l ,
- $h^{l'}$ is the modified LR being optimized,
- h^{l*} is the optimal modified representation,
- $\mathcal{L}_{\text{flip}}$: loss encouraging the prediction of the target concept c_j to change state (i.e., flipping the predicted label),
- $\mathcal{L}_{\text{preserve}}$: loss enforcing invariance of all non-target concepts c_k for $k \neq j$,
- $\|h^{l'} - h^l\|_2^2$: regularization term ensuring the modified representation remains close to the original LR,
- γ, λ, β : weighting coefficients controlling the trade-off between the intervention strength, the preservation of other concepts, and the minimality of the perturbation.

In the context of LLM studied previously, such interventions have been used to remove specific pieces from the model’s internal representation of the game state [Li et al., 2024, Karvonen, 2024]. The main observation is that, in the vast majority of cases, the model behaves as if the piece had effectively disappeared. The LLM continues to generate legal moves while correctly accounting for the absence of the removed piece. This provides evidence for the causal role of these internal representations in the model’s behavior.

A main limitation of the probing approach is that the concepts under investigation must be predefined before any analysis can be performed. This introduces a strong bias in what can be studied. Moreover, annotating a dataset for each concept is extremely time-consuming and requires significant human effort. Interventions in LS are also difficult to interpret, as it is not always clear

what constitutes a meaningful change in the network’s output under a given perturbation.

Moreover, probing methods provide a binary view of concepts, focusing on whether a given concept is present or absent in a LR. This binary perspective can be overly restrictive, as concepts may be represented in a more continuous or graded manner, with varying degrees of presence. To obtain a more detailed characterization of these representations, other approaches have been proposed, such as the analysis of interpretable directions in LS.

2.3 Interpretable Directions

To overcome the limitations of binary concept detection in probing methods, several approaches have been proposed to analyze the structure of LR. Among them, the notion of interpretable directions has emerged as a natural extension, aiming to capture the continuous nature of how concepts are encoded in representation spaces. Empirically, it has been observed that shifting a LR along a direction can induce coherent and semantically meaningful changes in the model’s output [Haas et al., 2024, Voynov and Babenko, 2020]. Such transformations can be mathematically express as:

$$h' = h + \alpha \mathbf{v}_i,$$

where:

- h denotes the original embedding,
- $\mathbf{v}_i$ is a unit vector associated with concept i ,
- α controls the strength of the intervention,
- h' is the resulting perturbed representation.

In many cases, the resulting change in the model’s output varies smoothly with α , suggesting an approximately linear CS with respect to certain semantic factors. This observation has led to the hypothesis that neural representations may organize concepts in a partially linear manner within the embedding space.

Vectors of this form are referred to as interpretable directions, as they correspond to transformations that can be associated with specific semantic attributes. The objective of this line of work is therefore to identify directions in LS that align with meaningful variations in the encoded concepts. This is illustrated schematically in Figure 2.4, where arrows represent semantic shifts from one concept toward others, and moving along these directions induces a coherent change in the model’s behavior, suggesting that the transition between concepts is continuously encoded along these axes.

This phenomenon has been extensively studied in generative image models. As illustrated in Figure 2.5, where each row corresponds to a distinct direction $\mathbf{v}_i$ and each column to an increasing value of α , controlled perturbations of LR along interpretable directions can induce specific semantic transformations, such as stroke thickness modification in digit images, head rotation in facial images, or texture and background changes in natural images [Voynov and Babenko, 2020]. These observations have led to the development of the field of LS editing, whose objective is to manipulate generated content directly through modifications in LS rather than through changes in the input prompt.




Figure 2.4: Interpretable directions in the latent space.

A central challenge remains identifying such interpretable directions. Most existing approaches rely on auxiliary models that encode human defined semantics, such as CLIP [Haas et al., 2024], to relate latent variations to textual concepts. Consequently, these methods inherit a limitation similar to that of probing techniques, since concept discovery depends on a pre-existing semantic framework defined by humans.

To overcome this limitation, a research direction known as unsupervised semantic discovery has emerged. Its goal is to uncover interpretable directions in LS without explicit semantic supervision. In these approaches, one model applies perturbations along random directions in LS, while a second model attempts to recognize the applied transformation. If a stable correspondence emerges between the perturbation and its prediction, this suggests that the considered direction corresponds to a coherent CS [Voynov and Babenko, 2020]. In this framework, human intervention occurs only a posteriori, to verify whether the discovered directions indeed correspond to interpretable semantic variations.

Despite these advances, several challenges remain. It appears that not all concepts encoded by NN are linearly identifiable in LS. In practice, interpretable direction methods typically discover only a few dozen directions per dataset, which intuitively seems insufficient to account for the full diversity of semantic variations that a generative model is capable of producing.

In addition, LR often exhibit a phenomenon known as entanglement, illustrated in Figure 2.6, where multiple semantic attributes are mixed within the same directions of the representation space. For instance, in a generative face model, moving along the “glasse” direction may simultaneously modify age and gender, whereas a disentangled representation would let each axis control a single semantic attribute while leaving the others unchanged. Ideally, one would like each semantic factor to vary independently along a specific direction, without affecting other attributes.

In practice, this is rarely the case, as a single direction in LS often encodes multiple semantic concepts simultaneously, such that modifying one concept



Figure 2.5: Interpretable directions in the latent space of a generative model [Voynov and Babenko, 2020].

can unintentionally affect several others. This lack of separability indicates that semantic information is not cleanly factored in the representation space, which motivates the search for more structured decompositions. This limitation has motivated a stronger hypothesis about how NN organize information internally, namely the *Superposition Hypothesis* (Hyp. 2), introduced in the next section. Under this hypothesis, Sparse Autoencoder (SAE) offer a way to analyze the structure of LR by recovering the underlying features encoded in them.

2.4 Sparse Autoencoder

Approaches based on latent directions face limitations. Only a restricted subset of concepts admit a linear representation, and even these directions are often entangled with multiple semantic concepts, as illustrated in Figure 2.6. To address this limitation, an extension of the *Concept Hypothesis* (Hyp. 1) has been proposed under the name *Superposition Hypothesis* (Hyp. 2), formalized in Hypothesis 2 [Elhage et al., 2022].




Figure 2.6: Illustration of entanglement in the latent space of a generative face model [Haas et al., 2024].

Hypothesis 2: Superposition Hypothesis

Concepts are preferentially represented by distinct linear directions in LS. When the number of concepts exceeds the available dimensions, multiple concepts are encoded in superposition within the same directions.

According to this hypothesis, in an unconstrained LS, each concept is represented in its own independent linear direction. However, NN operate in finite-dimensional spaces, so when the number of concepts exceeds the available dimensions, such a one-to-one correspondence becomes impractical. The network must therefore reuse its representational dimensions, allowing multiple concepts to be encoded using the same directions. This shared encoding is referred to as superposition [Bricken et al., 2023]. As a result, in a superposed representation, linear analysis of the LS may not isolate individual concepts, the resulting directions reflect a mixture of several semantic attributes.

The idealized non-superposed representation and can be mathematically expressed as follows:

$$\mathbf{h} = \sum_{i=1}^n \alpha_i \mathbf{v}_i, \quad \forall i \neq j, \quad \mathbf{v}_i^\top \mathbf{v}_j \approx 0, \quad \|\boldsymbol{\alpha}\|_0 \ll n,$$

where:

- $\mathbf{h}$ denotes a LR,
- n is the total number of concepts,
- $\mathbf{v}_i$ represents the direction associated with concept c_i ,
- α_i corresponds to the activation strength of concept c_i in the representation $\mathbf{h}$,
- $\mathbf{v}_i^\top \mathbf{v}_j \approx 0$ indicates that concept directions are approximately orthogonal, meaning each concept independently controls a single factor of variation,
- $\|\boldsymbol{\alpha}\|_0 \ll n$ indicates that only a small fraction of concepts are active for a given representation.

Under the *Superposition Hypothesis* (Hyp. 2), the phenomenon of superposition becomes an obstacle for any method that attempts to analyze LR directly




Figure 2.7: Transformation from a dense latent space (a) to a sparse SAE representation (b).

in the original activation space. This has motivated the development of methods that aim to recover a less superposed LR. A common approach is to map LS into a higher-dimensional space while imposing additional constraints that encourage individual concepts to be represented separately. This transformation is illustrated in Figure 2.7, while embeddings are entangled in the initial LS, the sparse representation aligns each concept with a dedicated feature axis.

SAE provide a practical implementation of this approach and can be defined as follows:

$$z = f_{\text{enc}}(h), \quad \hat{h} = f_{\text{dec}}(z)$$

$$\mathcal{L}_{\text{auto}} = \gamma \|h - \hat{h}\|_2^2 + \lambda \|z\|_1$$

where:

- $f_{\text{enc}}(h)$ is the encoding function that projects the original representation h into a higher-dimensional sparse space,
- $f_{\text{dec}}(z)$ is the decoding function that reconstructs the original representation from the sparse code,
- $z \in \mathcal{Z}$ is the sparse representation of h , with $\dim(z) \gg \dim(h)$,
- $\hat{h}$ is the reconstructed representation,
- $\gamma \|h - \hat{h}\|_2^2$ is the reconstruction loss ensuring that the original representation can be faithfully recovered,
- $\lambda \|z\|_1$ is the sparsity regularization term that encourages most components of z to be zero,
- γ, λ are weighting coefficients controlling the trade-off between reconstruction fidelity and sparsity.

In practice, both f_{enc} and f_{dec} are implemented as linear projections between the original representation space and the high-dimensional space $\mathcal{Z}$. This design follows the same principle as in probing methods, the transformation used to analyze the representation should remain as simple as possible. More expressive



Figure 2.8: A concept direction corresponding to the Golden Gate Bridge, discovered by a sparse autoencoder applied to a large language model [Templeton et al., 2024].

projection functions could themselves introduce complex structures into the transformed representation. In that case, it would become difficult to determine whether the observed concept reflects a CS already present in the model’s or structure introduced by the projection.

The parameters of these projection functions f_{enc} and f_{dec} are learned by minimizing $\mathcal{L}_{\text{auto}}$ using SGD. Under this formulation, each dimension of $\mathcal{Z}$ is interpreted as a dedicated concept direction.

The method proves to be powerful when successfully applied to language models [Templeton et al., 2024]. In these applications, SAE are used to desuperpose the LR of LLM. Once this representation is obtained, it has been shown that many concept vectors exhibit semantic properties. A example is the Golden Gate Bridge concept direction, illustrated in Figure 2.8, which activates systematically whenever the input text or image refers to the Golden Gate Bridge, regardless of the language or modality. This demonstrates that the learned concept is not tied to a specific surface form but captures a deeper semantic abstraction shared across modalities.

SAE also provide a way to intervene directly on the LR learned by the model. A concept can be artificially activated by modifying its corresponding feature in the sparse representation z , even when the input contains no reference to it. The modified representation is then projected back into the original activation space, after which the model proceeds with standard next-token generation from h^* :

$$h^* = f_{\text{dec}}(z^*), \quad \text{where} \quad z^* = z + \delta e_i,$$

where:

- z is the original sparse representation of the input,
- z^* is the modified sparse representation,
- e_i is the basis vector of $\mathcal{Z}$ associated with the target concept direction,
- δ controls the strength of the artificial activation,
- h^* is the resulting edited representation in the original activation space.

With this method, a reference to the Golden Gate Bridge can be injected into the LLM’s internal representation. This causes the model’s output to shift

toward discussing the Golden Gate Bridge, even when the prompt contains no reference to it. This provides evidence that the learned feature plays a causal role in the model’s behavior. Such interventions enable control over specific model behaviors through LS editing.

Despite their empirical success, *Superposition Hypothesis* (Hyp. 2) exhibit several important limitations that question their assumptions. These limitations can be summarized as follows:

- **Instability of the decomposition.** Repeating the SAE training process multiple times on the same model does not yield the same decomposition of the LS. Different runs produce different concept directions in $\mathcal{Z}$, even when achieving comparable reconstruction performance, which raises questions about the reliability of the discovered CS.
- **Lack of semantic coverage.** In the high-dimensional space $\mathcal{Z}$, a large fraction of dimensions do not correspond to any interpretable semantic concept. This abundance of non-semantic directions raises questions about whether the decomposition truly reflects the model’s CS.
- **Sparsity–semantics trade-off.** Increasing the sparsity coefficient λ beyond a certain threshold degrades the semantic quality of the learned features [Jermyn et al., 2024]. Suggesting that enforcing excessive sparsity comes at the cost of interpretability.
- **Failure of scaling to remove superposition.** Under the *Superposition Hypothesis* (Hyp. 2), increasing the dimensionality of the representation space should reduce superposition. However, this is not observed in practice. Even in very high-dimensional LS, representations remain superposed, indicating that scaling alone does not resolve the entanglement.

These limitations suggest that *Superposition Hypothesis* (Hyp. 2), while powerful, do not fully resolve the problem of CS. Convincing empirical results do not necessarily validate the underlying theoretical assumptions.

Probing, interpretable directions, and *Superposition Hypothesis* (Hyp. 2) all rest on a shared assumption that the CS can be recovered through linear decompositions of LR. This assumption has driven significant empirical progress, yet it remains theoretically fragile. The following section therefore introduces a approach that relaxes this constraint.

2.5 Clustering

All the LS analysis methods presented so far rely on assumptions about a linear CS. For probing methods, concepts must be at least partially linearly separable for simple probes. Latent directions assumes that variations associated with a concept can be captured by moving along a linear axis of the LS. The *Superposition Hypothesis* (Hyp. 2) assumes that concepts are represented through linear directions that become entangled because of representational capacity constraints. Although these assumptions have led to empirical successes, several observations suggest that the linearity hypothesis may not fully capture the CS.

For instance, some concepts can only be recovered using non-linear probes. Likewise, methods based on interpretable directions typically identify only a




Figure 2.9: Visualization of the latent space structure.

869 limited number of semantic factors. Finally, in sparse autoencoder approaches,
 870 enforcing excessive sparsity often degrades reconstruction quality and inter-
 871 pretability, suggesting that the underlying representation may not be perfectly
 872 described by a sparse linear decomposition. Taken together, these observations
 873 indicate that semantic concepts may not always exhibit a linear CS. This moti-
 874 vates the exploration of alternative approaches that rely on weaker assumptions
 875 regarding the geometry of LS.

876 To the best of our knowledge, clustering-based analysis is the only family
 877 of LS analysis methods that does not rely on any assumption of a linear CS
 878 [Ren et al., , Wei et al.,]. Rather than assuming that concepts correspond to
 879 directions, clustering methods only assume that conceptual similar inputs produce
 880 nearby LR. Conversely, inputs that differ significantly from a concept perspective
 881 are expected to be located farther apart in the LS. This is a considerably weaker
 882 assumption, in the sense that it is empirically verified in a wide range of settings.

883 Under this assumption, the discovery of concepts can be reformulated as the
 884 identification of groups of embeddings that occupy the same region of the LS.
 885 This is shown schematically in Figure 2.9, where the points represent embeddings
 886 associated with concepts A, B, and C, and the ellipses indicate the approximate
 887 extent of each group and highlight their geometric separation.

888 This idea forms the basis of the research field known as deep clustering. In
 889 these approaches, a collection of inputs is projected into a LS using a NN. The
 890 resulting embeddings form a point cloud on which standard clustering algorithms
 891 can be applied. The obtained clusters are then interpreted as representing
 892 distinct concepts.

893 Deep clustering has been particularly successful in unsupervised image classi-
 894 fication. For this type of task, a specific NN architecture called an autoencoder is
 895 commonly used. An autoencoder is composed of an encoder f_{enc} and a decoder
 896 f_{dec} . The encoder maps an input $x^{(i)} \in \mathcal{X}$ to a low-dimensional LR $z^{(i)} \in \mathbb{R}^d$.
 897 The decoder then maps this representation back to the input space to produce
 898 a reconstruction $\hat{x}^{(i)}$. The model is trained by SGD to minimize the recon-

struction error between the original input and its reconstruction using a Mean Squared Error (MSE) loss. This objective encourages the LS to preserve the most informative features of the input while enforcing a compact representation.

This can be expressed mathematically as follows:

$$x^{(i)} \xrightarrow{f_{\text{enc}}} z^{(i)} \xrightarrow{f_{\text{dec}}} \hat{x}^{(i)}, \quad d \ll \dim(\mathcal{X}), \quad \mathcal{L}_{\text{rec}} = \frac{1}{N} \sum_{i=1}^N \|x^{(i)} - \hat{x}^{(i)}\|_2^2.$$

Once the autoencoder has been trained sufficiently well, the reconstruction loss $\mathcal{L}_{\text{rec}}$ should be low enough for the learned representation to be useful.

A clustering algorithm can then be applied to the compressed representations of the images. Let $\mathcal{D}$ denote a dataset of inputs, through the encoder f_{enc} , this dataset is mapped to a point cloud in the LS:

$$\mathcal{Z} = \{z^{(i)} \in \mathbb{R}^d \mid z^{(i)} = f_{\text{enc}}(x^{(i)}), x^{(i)} \in \mathcal{D}\}.$$

A classical clustering algorithm is then applied to the set $\mathcal{Z}$ of embeddings to identify groups of similar representations. Each resulting cluster can therefore be interpreted as a group of images that the model considers similar.

An example of this methodology can be observed in Figure 2.10, where the approach is applied to the STL-10 dataset [Coates et al., 2011] using k -means clustering [MacQueen, 1967] with a number of clusters arbitrarily fixed to $K = 10$ [Xie et al.,]. For each cluster, the corresponding line shows the 10 samples closest to the associated centroid in the LS.

We observe that the method captures coherent semantic categories such as animals, vehicles, and flying entities. This result is particularly striking given that the only available information to the system consists of raw images. It suggests that the compression process induces a LS in which semantically meaningful CS emerge.

More advanced approaches combine representation learning and clustering within a single optimization process. In this setting, additional loss terms encourage embeddings to organize themselves around cluster centroids during training. This often produces LS that are easier to analyze and improves the quality of the discovered semantic CS [Guo et al., , Miklautz et al.,].

Despite its effectiveness, clustering-based analysis presents several important limitations. First, as with most unsupervised semantic discovery techniques, interpreting the resulting clusters remains partly subjective. Benchmark datasets often provide class labels for evaluation, but the discovered clusters may not match the categories that humans consider relevant.

Alternative semantic groupings may also be equally valid. Second, clustering methods are highly sensitive to design choices, as the number of clusters, the distance metric, and the clustering algorithm itself can all significantly influence the results, making parameter selection challenging and often requiring substantial empirical tuning.

Finally, to the best of our knowledge, no intervention technique has yet been proposed in the context of clustering-based analysis. Unlike probing or SAE analysis, which allow targeted modifications of LR to study their causal role, clustering methods currently offer no mechanism to alter the model’s behavior by acting directly on the discovered CS.




Figure 2.10: Clusters discovered by a deep clustering model applied to natural images [Xie et al.,].

2.6 Conclusion

The four families of methods reviewed in this chapter are summarized in Figure 2.11. These include probing, interpretable directions, sparse autoencoders, and clustering, all of which build upon the *Concept Hypothesis* (Hyp. 1) introduced in Section 2.1. Each of them approaches the question from a different angle, but they share a common objective: uncover how LR are structured and what information can be extracted from them.

However, a fundamental difficulty remains. There is currently no consensus on how to rigorously evaluate whether these methods truly capture the internal concepts of a model. Most existing metrics assess reconstruction quality or linear separability, but these do not directly measure whether the discovered CS correspond to meaningful abstractions from the model’s perspective. This leaves open the question of what it means, formally, for a concept to be represented in a NN.

Addressing this question may require contributions from disciplines beyond computer science and mathematics. As the notion of semantics is inherently tied to human interpretation, a purely formal or computational account may be insufficient. Insights from philosophy, cognitive science, or psychology could prove valuable in grounding the notion of concept in a way that is both theoretically rigorous and empirically tractable. Despite these open questions, the methods




Figure 2.11: Overview of the four latent space analysis methods studied in this chapter, applied to the same point cloud.

958 studied in this chapter provide a practical foundation for analyzing the internal
 959 representations of NN.

960 Building on the *Concept Hypothesis* (Hyp. 1) and the insights developed
 961 throughout this chapter, this manuscript investigates a representation-based
 962 approach to explainability in RL. To establish the necessary foundations, the
 963 next chapter introduces the principles of RL. The subsequent chapter then
 964 presents the proposed method for analyzing these representations and studying
 965 their potential for explainability in RL (Chapter 5).

Chapter 3

Reinforcement Learning

This chapter introduces the fundamental concepts of Reinforcement Learning (RL), providing the necessary foundation for the ideas developed throughout the rest of this manuscript. Its structure is as follows:

Section 3.1: An overview of the milestones that led to the emergence of RL as a unified field.

Section 3.2: A clarification of what is meant by an agent and how it relates to the RL algorithm.

Section 3.3: A formalization of the problem that RL algorithms aim to solve.

Section 3.4: An explanation of how policies are progressively refined through interaction with the environment to approach the optimal policy.

Section 3.5: A summary of the key ideas and a discussion of the broader significance of RL.

3.1 Historical Developments

RL is today presented as a major field within artificial intelligence. However, it originates from several research traditions that progressively converged toward a common framework. In this section, we outline the main stages in the emergence of this field, highlighting the key ideas and milestones that shaped its evolution.

3.1.1 Behavioral foundations (1900–1950)

The origins of RL can be traced back to 20th-century animal psychology studies [Thorndike, 1911]. These studies aimed to understand how behaviors are formed and modified. The central idea emerging from this work is that learning is driven by the consequences of actions. Behaviors followed by positive outcomes are more likely to be repeated, whereas those associated with negative outcomes tend to gradually decrease in frequency.

This perspective was later developed within behaviorism. In this school of psychology, rewards and punishments are used to gradually shape behavior

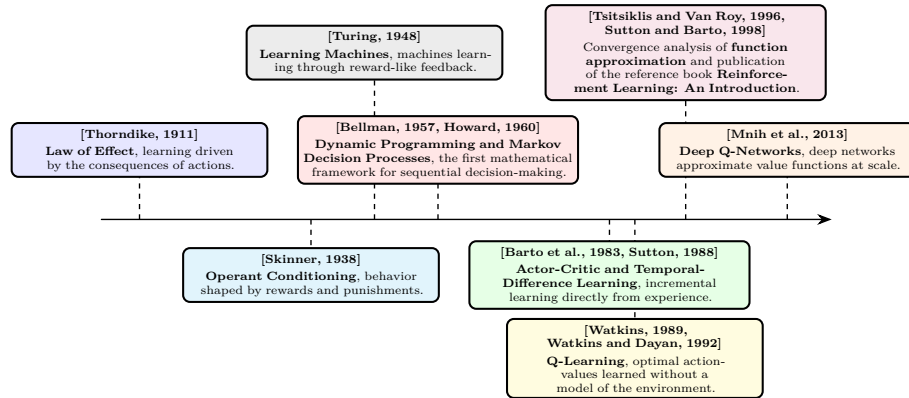


Figure 3.1: Evolution of reinforcement learning prior to its fusion with deep learning.

through repeated interaction with the environment [Skinner, 1938]. The term “reinforcement” in the context of learning comes from this school of thought, where it refers precisely to the process of strengthening a behavior through its consequences.

This idea led to the suggestion that learning principles observed in living organisms could be transferred to artificial systems. Machines were therefore envisioned as potentially learning through reward-like mechanisms, guided by evaluative feedback signals resembling a “pleasure–pain system” [Turing, 1948]. Although these studies did not involve formal mathematical models, they introduced the central idea of RL, trial-and-error learning guided by environmental feedback.

3.1.2 Sequential decision-making (1950–1970)

A second foundation of RL comes from optimal control theory. This field studies how to determine sequences of decisions that optimize a cumulative performance criterion over time. In this context, the decision-making entity is called an agent. The agent interacts with the environment by selecting actions. The function that maps an observed state to the chosen action is called a policy. The objective is to find an optimal policy, in the sense that it maximizes the expected cumulative reward over time.

This framework led to the formalization of Markov Decision Processes (MDP), which provide a mathematical model for sequential decision-making [Howard, 1960]. MDP describe a agent interacting with an environment through a set of states, actions, rewards, and transition rules. The earliest method proposed to solve MDP is dynamic programming [Bellman, 1957]. It proceeds by breaking down a complex decision problem into smaller, more manageable subproblems, solving each one, and combining the results to obtain an overall solution. However, dynamic programming requires a full and exact description of the environment, which is only feasible for small problems.

1023 3.1.3 Emergence of reinforcement learning (1970–2000)

1024 In order to address this issue, temporal-difference learning methods were in-
 1025 troduced [Sutton, 1988]. This mechanism enables incremental learning directly
 1026 from experience without requiring a complete model of the environment. It can
 1027 be interpreted as a sample-based approximation of the Bellman equations.

1028 Building on these ideas, the notion of value functions emerged as a way
 1029 to estimate the long-term benefit of being in a given state or taking a given
 1030 action. These functions assign a numerical score to states or state-action pairs,
 1031 reflecting the expected cumulative reward an agent can obtain from that point
 1032 onward. They thus provide a compact way to evaluate and compare different
 1033 courses of action. This line of work culminated in Q-learning, which showed
 1034 that optimal action-value functions can be learned directly from interaction
 1035 with the environment, without requiring a model of its dynamics [Watkins, 1989,
 1036 Watkins and Dayan, 1992].

1037 The publication of the foundational reference work [Sutton and Barto, 1998]
 1038 contributed to the formalization of RL as a coherent field. From that point
 1039 onward, the field can be considered well established and widely recognized as a
 1040 distinct area of research.

1041 3.1.4 Consolidation of reinforcement learning (2000–2013)

1042 A limitation of early algorithms such as Q-learning is their reliance on explicit
 1043 state-action tables. Although these methods remove the need to know the
 1044 full dynamics of the MDP, they still require storing a value for each possible
 1045 state-action pair. As a result, the agent must effectively visit and maintain
 1046 estimates for a potentially enormous number of states in the environment. This
 1047 requirement becomes infeasible in large-scale MDP, where the state space is too
 1048 large to be exhaustively explored or stored in memory.

1049 To address this issue, research shifted toward more general representations
 1050 of value functions and policies. Instead of maintaining explicit tables, function
 1051 approximation methods were introduced. These allowed value functions and
 1052 policies to be represented using parametric models such as linear approximators or
 1053 perceptron-like architectures [Barto et al., 1983, Tsitsiklis and Van Roy, 1996].
 1054 This change enabled generalization across states, reducing the dependence on
 1055 exhaustive exploration of the state space. The introduction of Deep Q-Networks
 1056 (DQN), which demonstrated that deep Neural Network (NN) could be used to
 1057 approximate value functions at scale [Mnih et al., 2013].

1058 Figure 3.1 provides a timeline of the key milestones that shaped the develop-
 1059 ment of reinforcement learning prior to its fusion with deep learning.

1060 3.1.5 Deep reinforcement learning (2013–2026)

1061 The introduction of Deep Learning (DL) into RL marked a major turning point
 1062 in the field. A series of algorithms significantly improved scalability. Deep Deter-
 1063 ministic Policy Gradient (DDPG) extended actor-critic methods to continuous
 1064 action spaces [Lillicrap et al., 2015]. Proximal Policy Optimization (PPO) later
 1065 became a standard algorithm in RL due to its simplicity, stability, and ability to
 1066 perform in both continuous and discrete action spaces [Schulman et al., 2017].




Figure 3.2: Evolution of deep reinforcement learning.

These advances established a practical foundation for large-scale policy optimization. Notably, AlphaZero demonstrated that the combination of deep NN, RL, and tree search could reach superhuman performance in the game of Go [Silver et al., 2017b]. This success was later extended to even more complex real-time strategy environments such as StarCraft II, where agents must handle long-term planning, partial observability, and large action spaces [Samvelyan et al., 2019]. Similarly, in Dota 2, large-scale multi-agent RL systems such as OpenAI Five showed that self-play combined with deep policy optimization can achieve professional-level performance in highly dynamic and stochastic environments [OpenAI et al., 2019].

More recently, RL has expanded beyond its traditional application domains. With the emergence of increasingly agentic views of Large Language Model (LLM), it has become possible to train these models to perform specific tasks using RL-based techniques. In particular, RL from human feedback has become a method for aligning LLM with human preferences. It does so by formalizing the problem as an MDP and using RL to solve it [Ouyang et al., 2022]. Similarly, DeepSeekMath shows how RL can enhance the mathematical reasoning abilities of LLM [Shao et al., 2024]. It demonstrates that an artificial agent can learn to reason in natural language in order to accomplish a given task. The skills acquired during this learning process also transfer to other similar tasks, leading to improved overall performance. For instance, training a LLM with RL on mathematical problem-solving can also enhance its programming capabilities.

Figure 3.2 provides a timeline of the key milestones that shaped the emergence of deep reinforcement learning.

3.2 Agent Paradigm

The RL framework is built around the agent paradigm, in which the decision-making entity is referred to as the agent. A standard definition is:




Figure 3.3: The agent paradigm.

“An agent is anything that can be viewed as perceiving its environment through sensors and acting upon that environment through actuators.”
[Russell and Norvig, 2021]

In the RL setting, this definition takes the following form: The agent perceives its environment through observations. These observations are then processed by a decision-making mechanism referred to as the policy. The policy determines how the agent modifies its environment through the actions it selects. Figure 3.3 illustrates this interaction.

The objective of an RL agent is to maximize the cumulative reward it receives through its interactions with the environment. To achieve this objective, the agent’s policy must progressively improve through experience. The procedure responsible for this learning process is the RL algorithm, which progressively modifies the policy based on the experience collected through interaction.

In this manuscript, the RL algorithm is considered as a process external to the agent. Other presentations of RL describe the agent as both the learner and the decision-making entity [Sutton and Barto, 1998]. We chose to distinguish the agent from the RL algorithm because this provides a more general definition of an agent.

An agent can therefore exist without learning, for example when its behavior is entirely specified by a set of programmed rules. Such an agent still perceives, processes information, and acts upon its environment, but its behavior is not the result of RL.

This distinction also makes the definition consistent with modern implementations of RL [Espeholt et al., 2018, Liang et al., 2018], in which the acting policy and the learning algorithm are separate software components. Several copies of the policy, each interacting with its own instance of the environment, collect experience in parallel, while a single learning algorithm uses this experience to update the shared policy parameters.

1122 3.3 Problem Formulation

1123 Now that the main developments shaping RL have been outlined, the RL
 1124 problem can be stated formally. In essence, RL aims to build a system that
 1125 makes sequential decisions in order to maximize a cumulative reward signal over
 1126 the long term. The purpose of this section is to translate this intuitive goal into
 1127 a rigorous mathematical framework. To this end, the presentation proceeds step
 1128 by step:

- 1129 1. Modeling the RL problem as a Markov Decision Process.
- 1130 2. Defining the notion of a policy, which formalizes the agent’s behavior.
- 1131 3. Describing the interaction between the agent and the environment through
 1132 trajectories.
- 1133 4. Introducing value functions to evaluate policies.
- 1134 5. Defining the notion of an optimal policy.

1135 3.3.1 Markov Decision Processes

1136 To specify the task and the agent’s capacity to perceive and act, the problem is
 1137 modeled as a MDP.

1138 Before providing a more detailed definition, it is useful to clarify the scope
 1139 of the MDP problems considered in this manuscript. We therefore restrict the
 1140 analysis to the following configuration:

- 1141 • The single-agent learning setting is considered. If any other agents are
 1142 present in the environment, they are treated as part of the environment
 1143 dynamics.
- 1144 • Finite-horizon MDP are considered. Each interaction sequence between the
 1145 MDP and the policy has a defined beginning and a defined end. Between
 1146 these two points, the number of interaction steps may vary depending on
 1147 the situation, but it always remains finite.
- 1148 • The partially observable framework is adopted, where the agent does
 1149 not have direct access to the true state of the system, but instead relies
 1150 on limited observations that provide only partial information about the
 1151 environment.

1152 The resulting framework can be formally described as a single-agent finite-
 1153 horizon partially observable Markov decision process. For simplicity, we will
 1154 refer to this setting as an “MDP” throughout this manuscript.

The following section aims to provide a mathematical description of the RL problem. The formulation adopted in this manuscript has been chosen to emphasize the main elements that will be used throughout the remainder of the manuscript. A MDP is defined by the tuple $(\mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{R}, p)$:

- $\mathcal{S}$: The set of states, representing all possible configurations of the environment. A state belonging to this set is denoted by $s \in \mathcal{S}$. Among these, two special subsets are distinguished: initial states, from which an interaction can begin, and terminal states, in which the interaction ends. The set of initial states is denoted by $\mathcal{S}^{\text{init}}$, with an initial state written as $s^{\text{init}} \in \mathcal{S}^{\text{init}}$. The set of terminal states is denoted by $\mathcal{S}^{\text{ter}}$, with a terminal state written as $s^{\text{ter}} \in \mathcal{S}^{\text{ter}}$.
- $\mathcal{O}$: The set of observations that the agent can perceive from the environment. An observation belonging to this set is denoted by $o \in \mathcal{O}$. Unlike the true state $s \in \mathcal{S}$, which is hidden from the agent, the observation $o \in \mathcal{O}$ provides partial information about the current state.
- $\mathcal{A}$: The set of actions that the agent may execute. An action belonging to this set is denoted by $a \in \mathcal{A}$. Actions influence the dynamics of the MDP, thereby shaping its future course. They represent the agent's capacity to act upon its environment.
- $\mathcal{R}$: The set of rewards that the agent can receive from the environment. A reward belonging to this set is denoted by $r \in \mathcal{R}$. $\mathcal{R}$ is a subset of $\mathbb{R}$. Rewards can be either positive or negative, depending on whether the corresponding outcome is considered beneficial or detrimental for the agent. The agent's goal is to maximize the cumulative sum of these rewards over time.
- p : The transition function, which defines the dynamics of the environment. Given a state s and an action a , this function defines a probability distribution over the joint outcome consisting of the next state s' , the observation o , and the reward r . This function is denoted by p , and is defined as: $p : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{P}(\mathcal{S} \times \mathcal{O} \times \mathcal{R})$. The notation $(s', o, r) \sim p(s, a)$ means that the tuple (s', o, r) is sampled from the distribution induced by $p(\cdot \mid s, a)$. The notation $p(s', o, r \mid s, a)$ denotes the probability of observing the tuple (s', o, r) under the distribution $p(\cdot \mid s, a)$.

The two quantities, observation o and state s , are not independent. They are jointly generated via $(s', o, r) \sim p(s, a)$, which implicitly defines the relationship between them. In our MDP formalization, this relationship is left implicit because it is never manipulated directly during learning.

However, to express mathematical equations, it is convenient to introduce the conditional distribution over states given an observation. We denote this distribution by $b : \mathcal{O} \rightarrow \mathcal{P}(\mathcal{S})$. The notation $b(s \mid o)$ denotes the probability of state s given the observation o , while $s \sim b(\cdot \mid o)$ indicates that the state s is sampled from this distribution.

3.3.2 Policy

The agent's decision-making capability is modeled by a policy function. The policy defines the stochastic mapping between observations $o \in \mathcal{O}$ and actions

1200 $a \in \mathcal{A}$. A policy function is denoted by π , and is defined as: $\pi : \mathcal{O} \rightarrow \mathcal{P}(\mathcal{A})$. The
 1201 notation $a \sim \pi(\cdot \mid o)$ means that the action a is sampled from the distribution
 1202 induced by $\pi(\cdot \mid o)$. The notation $\pi(a \mid o)$ denotes the probability of selecting
 1203 action a under the distribution $\pi(\cdot \mid o)$.

1204 It is important to note that this policy does not include any explicit mechanism
 1205 for estimating the underlying state from the observations. Instead, the policy
 1206 must learn from experience how to infer the information about the underlying
 1207 state that is not directly available in the observations. This reflects the design
 1208 of many modern RL implementations, where policies are often parameterized
 1209 by architectures capable of summarizing past observations, such as recurrent or
 1210 attention-based networks, without explicitly modeling the uncertainty over the
 1211 underlying state.

1212 3.3.3 Trajectory

1213 Now that both the MDP $(\mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{R}, p)$ and the policy π have been defined,
 1214 the interaction between them can be characterized. A sequence of consecutive
 1215 interactions between an environment and a policy is called a trajectory. A
 1216 trajectory is generated through the repeated interaction between the policy and
 1217 the environment, following the sequence below:

- 1218 1. The environment, being in state s , provides an observation o to the policy
 1219 π . Based on this observation, the policy outputs a probability distribution
 1220 over actions $\pi(\cdot \mid o)$, from which an action is sampled: $a \sim \pi(\cdot \mid o)$.
- 1221 2. The sampled action is executed in the environment. The environment then
 1222 generates a transition: $(s', o', r) \sim p(\cdot \mid s, a)$. It moves to a new state s' ,
 1223 and returns a new observation o' and a reward r .

1224 A trajectory t of length T can be written as:

$$t = (s_1, o_1, a_1, r_1, s_2, o_2, a_2, r_2, \dots, s_T, o_T, a_T, r_T).$$

1225 It is important to clearly distinguish between a trajectory distribution and its
 1226 realizations. The distribution over trajectories is denoted by τ , and a particular
 1227 trajectory is denoted by t . The distribution τ is determined by the initial state
 1228 s , the initial observation o , the environment dynamics p , and the policy π , and is
 1229 denoted by $\tau(s, o, p, \pi)$. The notation $t \sim \tau(s, o, p, \pi)$ means that the trajectory
 1230 t is sampled from the distribution induced by s , o , p , and π . The notation
 1231 $\tau(t \mid s, o, p, \pi)$ denotes the probability of observing the trajectory t under the
 1232 distribution $\tau(\cdot \mid s, o, p, \pi)$.

1233 An important property of the trajectory distribution is its recursive structure.
 1234 Since both the policy π and the transition function p are known, the distribution
 1235 over trajectories starting from s_1 with observation o_1 can be expressed in terms of
 1236 the distribution over trajectories starting from the next state s_2 with observation
 1237 o_2 . Specifically, the probability of a trajectory t factorizes as:

$$\tau(t \mid s_1, o_1, p, \pi) = \pi(a_1 \mid o_1) p(s_2, o_2, r_1 \mid s_1, a_1) \tau(t_2 \mid s_2, o_2, p, \pi),$$

1238 where $t_2 = (s_2, o_2, a_2, r_2, \dots, s_T, o_T, a_T, r_T)$ denotes the trajectory starting from
 1239 the second step.

There exist two special types of trajectories. First, a trajectory that terminates in a terminal state is called a terminal trajectory, and is denoted by $t_t \sim \tau_t(s, o, p, \pi)$. Second, a terminal trajectory that starts from an initial state and an initial observation is called an episode, and is denoted by $t_e \sim \tau_t(s^{\text{init}}, o, p, \pi)$.

3.3.4 Value Functions

To evaluate the ability of a policy π to collect rewards, the notion of return is introduced. The return of a trajectory $g(t, \gamma)$, is defined as the discounted sum of rewards collected along the trajectory:

$$g(t, \gamma) = r_1 + \gamma r_2 + \gamma^2 r_3 + \cdots + \gamma^{T-1} r_T = \sum_{i=1}^T \gamma^{i-1} r_i,$$

where $\gamma \in [0, 1]$ is the discount factor. This parameter controls the time horizon of the agent's objective. When γ is close to 0, the agent focuses almost exclusively on immediate rewards, leading to short-sighted behavior. As γ increases toward 1, future rewards gain more weight, and the agent adopts a more far-sighted strategy. In theory, since the goal is to maximize long-term cumulative reward, the discount factor would ideally be set to $\gamma = 1$. In practice, however, γ is often chosen slightly below 1 (e.g., 0.99) to stabilize learning, and to reflect the fact that distant future rewards are inherently more uncertain.

The return also admits a recursive form, which is central to the dynamic programming and RL algorithms presented later:

$$g(t, \gamma) = r_1 + \gamma g(t_2, \gamma),$$

From the notion of return, functions that estimate the expected return of a policy from a given starting point can be defined. These functions are called value functions. The first one considered is the state-value function V . It tells us how desirable a particular state is under a given policy π , given an observation o , the environment dynamics p , and the discount factor γ . Strictly speaking, this function should be written as $V(s, o, p, \gamma, \pi)$ to reflect all its dependencies. However, since both the environment dynamics p and the discount factor γ are fixed throughout learning, they are omitted from the notation. To emphasize that the value is tied to a specific policy, π is placed as a subscript. This yields the compact notation $V_\pi : \mathcal{S} \times \mathcal{O} \rightarrow \mathbb{R}$.

If $V_\pi(s, o) > V_\pi(s', o')$, then, under policy π , the agent prefers being in state s with observation o over state s' with observation o' . Formally, the state-value function is defined as:

$$V_\pi(s, o) = \mathbb{E}[g(t, \gamma)], \quad \text{with } t \sim \tau_t(s, o, p, \pi).$$

It is worth noting that the state-value function V_π is recursive, a property that follows directly from the recursive form of the return:

$$\begin{aligned} V_\pi(s, o) &= \mathbb{E}[g(t, \gamma)], \quad \text{with } t \sim \tau_t(s, o, p, \pi) \\ &= \mathbb{E}[r_1 + \gamma g(t_2, \gamma)] \\ &= \mathbb{E}[r_1 + \gamma V_\pi(s_2)] \end{aligned}$$

1274 If the expectation is expressed explicitly in terms of the distributions p and
 1275 π , the Bellman equation for V_π is obtained:

$$\begin{aligned} V_\pi(s, o) &= \mathbb{E}\left[r_1 + \gamma V_\pi(s_2)\right] \\ &= \sum_{a \in \mathcal{A}} \pi(a | o) \sum_{\substack{s' \in \mathcal{S} \\ o' \in \mathcal{O} \\ r \in \mathcal{R}}} p(s', o', r | s, a) \left[r + \gamma V_\pi(s', o')\right]. \end{aligned}$$

1276 Now that the value of a state has been expressed, the next step is to consider
 1277 the value of taking a specific action in a given state. Indeed, it is useful to
 1278 estimate the expected future gain of an action a in state s , in order to compare
 1279 different actions and determine which one is the most promising under a given
 1280 policy π . This function is called the state-action value function $Q_\pi : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$.
 1281 If $Q_\pi(s, a) > Q_\pi(s, a')$, then, under policy π , choosing action a in state s yields
 1282 a higher expected return than choosing action a' . It is defined as:

$$Q_\pi(s, a) = \mathbb{E}\left[r + \gamma g(t, \gamma)\right], \quad \text{with } t \sim \tau_t(s', o, p, \pi), \quad (s', o, r) \sim p(\cdot | s, a).$$

1283 The function Q_π can be related to V_π by recognizing that, after taking action
 1284 a in state s , the remaining return from the next state s' is precisely $V_\pi(s')$. This
 1285 yields:

$$\begin{aligned} Q_\pi(s, a) &= \mathbb{E}\left[r + \gamma g(t, \gamma)\right], \quad \text{with } t \sim \tau_t(s', o, p, \pi), \quad (s', o, r) \sim p(\cdot | s, a) \\ &= \mathbb{E}\left[r + \gamma V_\pi(s', o)\right] \\ &= \sum_{\substack{s' \in \mathcal{S} \\ o' \in \mathcal{O} \\ r \in \mathcal{R}}} p(s', o', r | s, a) \left[r + \gamma V_\pi(s', o)\right] \end{aligned}$$

1286 3.3.5 Optimal Policy

1287 The goal of RL is to find a policy that maximizes the expected cumulative
 1288 reward from any observation until the end of the trajectory. This policy is called
 1289 the optimal policy and is denoted by π^* . In a partially observable setting, the
 1290 optimal action is the one that maximizes the expected state-action value over
 1291 the distribution of possible states given the observation:

$$\pi^*(a | o) = \begin{cases} 1 & \text{if } a = \arg \max_{a' \in \mathcal{A}} \sum_{s \in \mathcal{S}} b(s | o) Q_{\pi^*}(s, a'), \\ 0 & \text{otherwise} \end{cases}$$

1292 As shown, both V_π and Q_π can be expressed explicitly in terms of the
 1293 environment dynamics p and the policy π . If p , b and π are fully known, one can
 1294 theoretically compute the optimal policy by unfolding the graph of all possible
 1295 trajectories starting from a given state s and observation o . This is possible
 1296 precisely because V_π and Q_π are recursive functions. Solving an MDP in this
 1297 way is known as dynamic programming.

1298 In practice, most environments do have too many states to allow an explicit
 1299 representation of the dynamics p . Storing and computing over the full transition

function becomes intractable. To overcome this limitation, instead of directly constructing the optimal policy, an iterative approach is adopted: starting from an initial policy, it is progressively improved so that it converges toward π^* . This iterative improvement through interaction with the environment is precisely what constitutes learning in RL, and it is the subject of the next section.

3.4 Learning

Now that the optimal policy π^* has been defined, the question is how to obtain it. The recursive Bellman equations for V_π and Q_π suggest dynamic programming approaches, but these scale poorly to large state spaces. RL offers an alternative by solving the MDP iteratively and generating a sequence of policies that progressively approach π^* through direct interaction with the environment. This section introduces two major families of RL algorithms, critic-only methods and actor-critic methods.

At a high level, every algorithm presented below repeats the same loop: the current policy interacts with the environment, the resulting trajectories are sent to the RL algorithm, and the algorithm uses them to produce a new policy.

3.4.1 Critic-Only Methods

The central idea behind all critic-only approaches is to iteratively construct an approximation of the state-action value function Q_π , denoted by $\hat{Q}_\pi$. Although $\hat{Q}_\pi$ is only an approximation of the true Q_π , the policy always acts greedily with respect to this approximation.

Since $\hat{Q}_\pi$ serves as the basis for the agent's decisions, it cannot rely on more information than the agent itself has access to. It must therefore be defined solely over observations, as a function $\hat{Q}_\pi : \mathcal{O} \times \mathcal{A} \rightarrow \mathbb{R}$. Its objective is to approximate the true Q_π , which depends on the underlying states s , using only the partial information available through o . In practice, this does not necessarily pose a problem, as observations often contain sufficient information to make accurate predictions about the underlying state. Moreover, the strong representational capacity of NN allows the critic to capture these relationships without explicitly representing the mapping between observations and states.

This leads to the following mathematical formulation:

$$\pi(a \mid o) = \begin{cases} 1 & \text{if } a = \arg \max_{a' \in \mathcal{A}} \hat{Q}_\pi(o, a'), \\ 0 & \text{otherwise.} \end{cases} \quad (3.1)$$

It is precisely this way of modeling the policy that gives this family of methods its name. The policy is not represented explicitly as a standalone function. Rather, it simply emerges from searching over the set of actions to maximize $\hat{Q}_\pi$.

To give a clear example of critic-only methods, they are presented in their simplest form: model-free, without temporal-difference learning, without replay buffer, without bootstrapping, and using only ε -greedy exploration. Naturally, many variants and extensions of this basic scheme exist, but the core idea remains the same.

When implementing a RL algorithm, one of the first questions is how to represent the value function. In critic-only methods, this means choosing a model for $\hat{Q}_\pi$. Any machine learning model capable of regression task can be used for this purpose. The set of parameters of the model is denoted by θ . In deep RL, a NN is used to model this function, in which case θ represents the weights and biases of the network.

In our simple setting, exploration is handled via ε -greedy selection. With probability $1 - \varepsilon$, the policy selects the greedy action that maximizes $\hat{Q}_\pi$, and with probability ε it selects an action uniformly at random. Formally, this ε -greedy policy is obtained by modifying the greedy policy defined above as follows:

$$\pi(a \mid o) = \begin{cases} 1 - \varepsilon + \frac{\varepsilon}{|\mathcal{A}|} & \text{if } a = \arg \max_{a' \in \mathcal{A}} \hat{Q}_\pi(o, a'), \\ \frac{\varepsilon}{|\mathcal{A}|} & \text{otherwise.} \end{cases} \quad (3.2)$$

The trajectory distribution induced by the resulting policy now depends on both π and ε , and is denoted by $t \sim \tau(s, o, p, \pi, \varepsilon)$. The parameter ε controls the degree of exploration. If the policy always chooses the action it currently believes to be optimal, it may miss better alternatives that $\hat{Q}_\pi$ has not yet learned to value correctly. Exploring suboptimal actions from time to time is therefore essential to verify that $\hat{Q}_\pi$ accurately reflects the true values, while still focusing most of the time on the most promising actions. In more advanced implementations, ε can be adapted over the course of learning to control exploration dynamically, for instance by decreasing it gradually over time.

The learning process consists of two steps that repeat until convergence:

1. **Data collection.** The current policy π interacts with the MDP. For a given starting state s_1^{init} and observation o_1 , the interaction produces a episode t_e sampled from $\tau_t(s_1^{\text{init}}, o_1, p, \pi, \varepsilon)$. For each step i along the episode, the return from that step onward is $g(t_{i:}) = \sum_{k=i}^T \gamma^{k-i} r_k$, which serves as the target value for the state-action pair (s_i, a_i) . This procedure is detailed in Algorithm 2.
2. **Model update.** The tuples $(s_i, a_i, g(t_{i:}))$ obtained in the previous step serve as supervised training data: the state-action pair (s_i, a_i) is the input, and the return $g(t_{i:})$ is the associated target. These examples are used to train $\hat{Q}_\theta$. The improved approximation is denoted by $\hat{Q}'_\theta$, which now predicts returns more accurately under π . Since $\hat{Q}'_\theta$ has changed, the policy derived from it also changes. This new policy is denoted by π' .

Because the policy is now π' , the approximation $\hat{Q}'_\theta$ must be trained to predict returns under π' . This requires returning to step 1 to collect fresh interaction data generated by π' . The two steps thus alternate until the policy and the value approximation stabilize. Algorithm 3 summarizes this whole critic-only learning loop.

Even this simple algorithm already illustrates a difficulty inherent to all forms of RL: the instability of the learning process. Each update from $\hat{Q}_\theta$ to $\hat{Q}'_\theta$ causes the policy to shift from π to π' . For the learning to remain stable, the change from $\hat{Q}_\theta$ to $\hat{Q}'_\theta$ should be as small as possible. This way, $\hat{Q}'_\theta$, trained to predict returns under π , remains reasonably accurate under π' . However, the smaller the change, the slower the learning progresses. A trade-off must therefore be

Algorithm 2: Data Collection**Input:**MDP $(\mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{R}, p)$ Policy π Discount factor $\gamma \in [0, 1]$ **Output:**Dataset $\mathcal{D}$ *Sample an initial state.***1** $i \leftarrow 1$;**2** Sample $s_i^{\text{init}} \in \mathcal{S}^{\text{init}}$;**3** Observe o_i ;*Episode generation by interaction with the MDP.***4 repeat***Action selection.***5** Sample $a_i \sim \pi(\cdot \mid o_i)$;*Environment transition.***6** Sample $(s_{i+1}, o_{i+1}, r_i) \sim p(\cdot \mid s_i, a_i)$;**7** $i \leftarrow i + 1$;**8 until** $s_i \in \mathcal{S}^{\text{ter}}$;*Return computation for each step of the trajectory.***9** $T \leftarrow i - 1$;**10 for** $i = 1, \dots, T$ **do****11** $g_i \leftarrow \sum_{k=i}^T \gamma^{k-i} r_k$;*Dataset construction***12** $\mathcal{D} \leftarrow \{(s_i, o_i, a_i, g_i)\}_{i=1}^T$;**13 Return** $\mathcal{D}$

1384 struck between stability and learning speed. Many improvements introduced
 1385 later in the literature aim at stabilizing this learning loop [Hessel et al., 2017].

1386 Convergence guarantees for such algorithms have been proved under a different
 1387 set of assumptions than the setting presented here. For example, in the classical
 1388 convergence results, the MDP is fully observable, the value function is stored
 1389 in a table, and the learning rate is decreased at a suitable rate over time
 1390 [Watkins and Dayan, 1992]. In modern implementations, the value function is
 1391 approximated by a NN. Theoretical convergence proofs no longer hold in this
 1392 setting. Nevertheless, deep methods have demonstrated remarkable empirical
 1393 success, making them the tool of choice for RL tasks.

1394 Critic-only approaches suffer from several limitations. The main limitation
 1395 is scalability to large action spaces. Since the policy must evaluate all possible
 1396 actions at each decision step to select the maximum, the computational cost
 1397 grows with the size of the action set. To address this issue, another family of
 1398 RL algorithms exists: actor-critic methods.

Algorithm 3: Deep Critic-Only Reinforcement Learning**Input:**

MDP $(\mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{R}, p)$
 Architecture $\hat{Q}$
 Exploration rate $\varepsilon \in [0, 1]$
 Discount factor $\gamma \in [0, 1]$
 Learning rate $\eta > 0$
 Initial parameters θ (optional)

Output:

Trained policy π_θ

```

1 if  $\theta$  is not given then
2   | Initialize  $\theta$  randomly;
   Exploration policy as defined in Equation (3.2)
3  $\pi_\theta \leftarrow \varepsilon$ -greedy policy with respect to  $\hat{Q}_\theta$ 
   MSE loss function as defined in Equation (1.2)
4  $\mathcal{L} \leftarrow$  Mean Squared Error (MSE);
5 while stopping criterion not satisfied do
   Data collection see Algorithm 2
6    $\mathcal{D} \leftarrow \text{DataCollection}(\text{MDP}, \pi_\theta, \gamma)$ 
   observation-action pairs and their returns, from  $\mathcal{D}$ 
7    $\mathcal{D}' \leftarrow \{((o, a), g) \mid (s, o, a, g) \in \mathcal{D}\}$ 
   Model update with SGD see Algorithm 1
8    $\theta \leftarrow \text{SGD}(\mathcal{D}', \hat{Q}, \mathcal{L}, \eta, \theta)$ 
   Final policy without exploration as defined in Equation (3.1)
9  $\pi_\theta \leftarrow$  policy with respect to  $\hat{Q}_\theta$ ;
10 Return  $\pi_\theta$ 

```

3.4.2 Actor-Critic Methods

Two principal modifications distinguish actor-critic methods from critic-only approaches. First, the objective is to approximate the state-value function V , instead of the state-action value function Q . Second, the policy π is given its own dedicated set of parameters. Two separate components are thus obtained: a value function approximation $\hat{V}_{\theta_1}$ and a parameterized policy π_{θ_2} .

Analogously to critic-only methods, $\hat{V}_{\theta_1}$ is trained using supervised learning on pairs obtained from interaction with the MDP. For each step i along a trajectory, the input is the state-observation pair (s_i, o_i) and the target is the observed return $g(t_{i:})$. The value function is thus trained to correct its predictions based on data collected by the current policy interacting with the environment.

The key difference lies in how the policy π itself is optimized. Since the policy now has its own set of parameters θ_2 , it can be updated directly. The update uses the notion of advantage $A : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{R}$, defined as:

$$A(s, a) = g(t) - \hat{V}_{\theta_1}(s, o),$$

where t is the trajectory that starts by taking action a in state s and then follows π for all subsequent steps.

When $A(s, a) > 0$, the action a performed better than expected, and the policy parameters θ_2 are adjusted to increase the probability $\pi_{\theta_2}(a | o)$. When $A(s, a) < 0$, the action a performed worse than expected, and the probability is decreased.

Both $\hat{V}_{\theta_1}$ and π_{θ_2} are improved simultaneously within the same loop: the critic provides a learned baseline for the actor, and the actor generates the data used to train the critic. This interplay reduces the instability inherent to RL. Algorithm 4 summarizes this whole actor-critic learning loop.

PPO (Proximal Policy Optimization), which is the current standard in RL, belongs to this family of actor-critic methods. It will be used throughout this manuscript for all RL experiments.

3.5 Conclusion

This chapter has provided an introduction to the fundamental principles of RL. As in Chapter 1, only the essential concepts required to understand the deep RL framework have been presented. The broader field of RL encompasses many additional methods and extensions.

Despite this limited scope, the defining characteristics of RL can already be observed in its simple learning procedure. A policy is iteratively improved using estimates obtained through the interaction between an agent and its environment. Learning therefore emerges from the repeated alternation of two elementary operations, which are performed until the policy and its associated value estimates converge.

This simple iterative structure brings RL conceptually close to DL. In both paradigms, simple operations repeated at scale can give rise to highly complex capabilities with limited human intervention. Their combination is particularly powerful. RL provides a general framework for learning through interaction, where a wide range of decision-making problems can be formulated as MDPs. DL, in turn, provides flexible function approximators capable of representing the value functions and policies required by this framework. Together, these two paradigms form the foundation of deep RL, whose empirical performance has led to remarkable results.

Originally inspired by biological learning processes, RL is now closer than ever to this initial intuition. Agents learn through interaction, adapt via trial and error, and progressively improve their behavior based on experience. This behavior is governed by deep neural architectures, which are the closest computational analog to organic brains. More than ever, RL stands as one of the core paradigms for achieving artificial general intelligence [Silver et al., 2021].

As the capabilities of RL continue to improve, a new challenge becomes increasingly important. The ability to learn complex behaviors with limited human intervention can make the resulting decisions difficult to understand. To address this challenge, a parallel research direction has consequently emerged. This is the subject of the next chapter, which focuses on explainability in the context of RL.

Algorithm 4: Deep Actor-Critic Reinforcement Learning

Input:

MDP $(\mathcal{S}, \mathcal{O}, \mathcal{A}, \mathcal{R}, p)$
 Architecture $\hat{V}$ (critic)
 Architecture π (actor)
 Discount factor $\gamma \in [0, 1]$
 Learning rate $\eta > 0$
 Initial parameters θ_1 of the critic (optional)
 Initial parameters θ_2 of the actor (optional)

Output:

Trained policy π_{θ_2}

```

1 if  $\theta_1$  is not given then
2   | Initialize  $\theta_1$  randomly;
3 if  $\theta_2$  is not given then
4   | Initialize  $\theta_2$  randomly;

   MSE loss function as defined in Equation (1.2)
5  $\mathcal{L} \leftarrow \text{MSE};$ 

6 while stopping criterion not satisfied do
   Data collection see Algorithm 2
7    $\mathcal{D} \leftarrow \text{DataCollection}(\text{MDP}, \pi_{\theta_2}, \gamma)$ 

   Advantage estimation: how much better/worse each action performed compared
   to the critic's expectation
8   for  $(s_i, o_i, a_i, g_i) \in \mathcal{D}$  do
9     |  $A(s_i, a_i) \leftarrow g_i - \hat{V}_{\theta_1}(s_i, o_i);$ 

   state and their returns, from  $\mathcal{D}$ 
10   $\mathcal{D}' \leftarrow \{(s, g) \mid (s, o, a, g) \in \mathcal{D}\}$ 

   Critic update with SGD see Algorithm 1
11   $\theta_1 \leftarrow \text{SGD}(\mathcal{D}', \hat{V}, \mathcal{L}, \eta, \theta_1)$ 

   Actor update: increase the probability of actions with positive advantage, decrease
   it for those with negative advantage
12   $\theta_2 \leftarrow \theta_2 + \eta \sum_{(s_i, o_i, a_i, g_i) \in \mathcal{D}} A(s_i, a_i) \nabla_{\theta_2} \log \pi_{\theta_2}(a_i \mid o_i);$ 

13 Return  $\pi_{\theta_2}$ 

```

Chapter 4

Explainability in Reinforcement Learning

Today, the majority of Artificial Intelligence (AI) agents with which humans interact are deployed in controlled or low-risk environments. For example, agents capable of playing games such as Go operate in settings where potential errors remain limited in terms of consequences [Silver et al., 2017a]. Even more autonomous systems, such as agents based on Large Language Model (LLM), remain constrained by human supervision mechanisms and are often restricted to tasks such as code generation or information retrieval [Chen et al., 2021, Lewis et al., 2021]. Safety-critical applications integrating machine learning components are still rare. Introducing autonomous agents in such environments, requires the establishment of strong guarantees [Seshia et al., 2022]. These requirements include in particular:

- validating the system’s functioning before deployment,
- identifying the origin of a malfunction in case of failure,
- and certifying that an observed failure will not occur again.

The main challenge lies in the fact that the performing agents currently available rely on both Reinforcement Learning (RL) and Deep Learning (DL) [Mnih et al., 2013]. As a result, the mechanisms underlying their decisions cannot be understood directly from the learned model. This lack of explainability makes their use difficult in safety-critical contexts. To address this limitation, the field of eXplainable Reinforcement Learning (XRL) has emerged as a subfield of eXplainable Artificial Intelligence (XAI). The objective of XRL is to develop explainability methods that make it possible to obtain explanations of RL agents. Due to the relevance of the topic and to foster broader real-world adoption of RL, a growing number of research teams are working on developing new explainability methods. This enthusiasm is reflected in a rapid increase in the number of publications, as illustrated in Figure 4.1.¹

¹The complete search protocol, including the databases consulted and the full set of query terms, has been archived on Zenodo <https://zenodo.org/records/19387785>.




Figure 4.1: Yearly growth of Explainable Reinforcement Learning research with at least one citation on arXiv.

The growing number and diversity of these methods make it increasingly difficult to obtain a unified view of the field. This motivated our work on developing a unified taxonomy of explainability methods for RL, initially presented in [Alaarabiou et al., 2024]. The taxonomy provides a common structure for organizing existing approaches according to the type of information they use and the form of explanation they provide. Building on this work, this chapter provides a detailed presentation of the proposed taxonomy and uses it to organize a broader state-of-the-art review of XRL.

This chapter is organized as follows:

- Section 4.1 begins by clarifying what we mean by an explanation in the context of XAI.
- Section 4.1 then introduces a taxonomy for categorizing existing XRL methods.
- Section 4.3 provides a comprehensive state-of-the-art review structured around this taxonomy.
- Section 4.4 concludes the chapter with a summary and discussion.

4.1 Explanation in Artificial Intelligence

In this section, we define what an explanation is in the context of XAI and how it can be obtained. We first introduce the concept of explanation in general. We then describe how this concept applies to XAI. Finally, we further detail the properties that characterize an explainability method in the context of XAI.

4.1.1 Explanation

Defining what constitutes an explanation is challenging because the concept is inherently multidisciplinary. When we refer to explanation, we are inherently referring to an explanation provided to a human. As soon as humans are involved, the problem is no longer purely computational or mathematical. It involves several fields of research, such as philosophy, psychology, and sociology. So far, these disciplines have not converged on a unified definition of explanation that can be directly applied to XAI [Miller, 2019, Bellucci et al., 2021].

For this reason, we propose the following definition of explanation. An explanation involves two entities: the object (Definition 4.1.1) being explained and the receiver of the explanation. The receiver is a human who aims to understand how the object operates and will be referred to as the observer throughout this manuscript (Definition 4.1.2). The observer seeks to understand the object and, in doing so, builds a mental model of its functioning. An explanation is therefore what allows the observer to refine and correct their mental model of how the object operates (Definition 4.1.4). This refinement is obtained through a transformation process that produces an explanation from the object being explained. We refer to this process as an explainability method (Definition 4.1.5), detailed further in Section 4.1.3. For an observer to understand an explanation, the information it contains must be interpretable (Definition 4.1.3). Information is considered interpretable when it can be directly understood by a human without requiring any additional transformation process.

Definition 4.1.1: Object

The object of an explanation is the system the observer wants to understand.

Definition 4.1.2: Observer

The observer is the human who aims to understand an object. The observer possesses a mental model of the object being explained. This mental model is an internal representation of how the observer believes the system operates. This representation may be incomplete or inaccurate.

Definition 4.1.3: Interpretable

Information is interpretable when it can be directly understood by an observer without requiring any additional transformation process.

Definition 4.1.4: Explanation

An explanation is any interpretable information that enables an observer to refine or correct their mental model of the object of the explanation.

Definition 4.1.5: Explainability Method

An explainability method is a transformation process that produces an explanation from the object being explained.

This definition of explanation provides the foundation for the approach to explainability adopted throughout this manuscript. This work adopts the perspective of *Mechanistic Hypothesis* (Hyp. 3) [Siegel and Craver, 2024, Zednik, 2019]. Within this framework, explainability is not regarded as a binary property but as a continuum that evolves with the development of scientific knowledge.

Accordingly, the opacity of a system should be interpreted as an epistemic limitation reflecting the current state of knowledge rather than as an intrinsic property of the system itself. This view is consistent with the methodological commitments of scientific inquiry. Scientific progress is driven by the continuous refinement of theories and analytical methods that provide progressively deeper

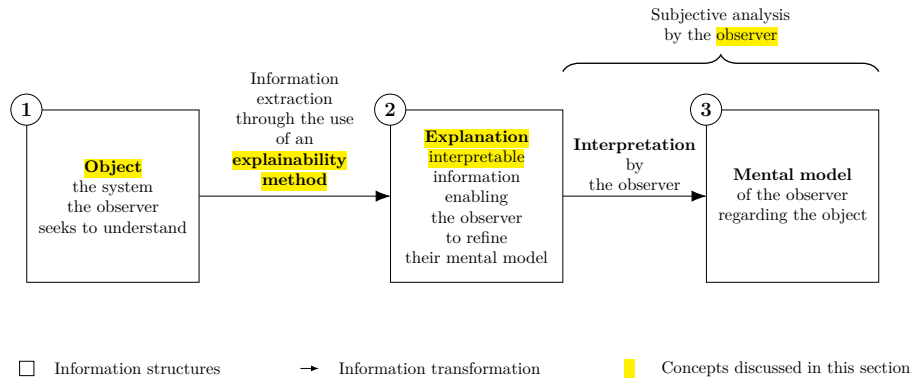


Figure 4.2: Explanation process.

1545 mechanistic accounts of observed phenomena. Consequently, explainability
 1546 should be regarded as a matter of degree, increasing as these mechanisms become
 1547 better understood.

Hypothesis 3: Mechanistic Hypothesis

Every system results from a sequence of causal computations. Because the rules governing these interactions are mechanical, the behavior of the system can be explained. If a system cannot be explained, this reflects a limitation of knowledge rather than an intrinsic property of the system.

1548

1549 This quantitative perspective allows us to express an intuitive idea: not all
 1550 explanations provide the same level of value [Doshi-Velez and Kim, 2017]. The
 1551 effectiveness of an explanation depends on its ability to improve the observer's
 1552 understanding of the system. For example, describing a marginal aspect of
 1553 a system is different from correcting an observer's misunderstanding of a key
 1554 mechanism.

1555 Figure 4.2 summarizes this whole explanation process, from the object of
 1556 the explanation to the mental model refined by the observer, as a chain of three
 1557 information structures, depicted as boxes and connected by arrows denoting a
 1558 transformation of information:

- 1559 (1) The object of the explanation, that is, the system the observer seeks to
 1560 understand.
- 1561 (2) The explanation, the interpretable information produced from the object.
 1562 This transformation is performed through the use of an explainability
 1563 method.
- 1564 (3) The mental model held by the observer regarding the object.

1565 A brace spanning boxes (2) and (3) indicates that this second half of the process,
 1566 from the explanation to the resulting mental model, constitutes a subjective
 1567 analysis carried out by the observer.

4.1.2 Explainable Artificial Intelligence

Now that the general framework of explanation has been established, we can examine how this concept applies in the context of AI. In the context of XAI, the object of explanation is an AI system. More precisely, XAI emerged as a response to the success of DL models [Gunning and Aha, 2019, Barredo Arrieta et al., 2020]. In these architectures, information is represented in a distributed manner across the network through patterns of activation and connections between computational units. Unlike symbolic approaches, individual components of Neural Network (NN) cannot be directly associated with human concepts. From this starting point, we identify two types of approaches:

- **Substitutive explainability:** In this approach, the underlying assumption is that there is nothing to explain within a NN. There is no semantic meaning associated with each computational component, and therefore nothing to interpret. This corresponds to the “black box” view of NN. Under this assumption, XAI consists of replacing as much of the DL system as possible with alternative AI methods [Rudin, 2019].
- **Intrinsic explainability:** In this approach, it is argued that NN inherently contain mechanisms that can be explained. The limitation is considered to lie not in the networks themselves, but in the absence of adequate tools to analyze their functioning [Olah et al., 2020].

This manuscript considers methods from both approaches in order to provide a comprehensive overview of the XAI field. According to the *Mechanistic Hypothesis* (Hyp. 3) adopted in this manuscript, we consider that any system can be explained. As a consequence, the perspective adopted throughout this manuscript is that of intrinsic explainability.

4.1.3 Explainability Method in XAI

Now that the general framework of XAI has been established, we need to further detail the properties of an explanation in this context. In our framework, an explanation in the field of XAI can be characterized by three components:

- **Source:** refers to the origin of the information used by the explainability method. The explainability method transforms the raw information produced by the AI system into an explanation that can be understood by an observer.
- **Form:** refers to the representation through which the explanation is communicated. The form defines how the information is presented to the observer and must allow human interpretation. An explanation can take many different forms, including text, images, computer code, and mathematical equations.
- **Subject:** refers to the aspect of the AI system that is being explained. In XAI, this distinction is commonly associated with the difference between local and global explanations. A local explanation describes why an AI system produced a specific output in a given situation. A global explanation aims to describe the general functioning of the AI system.

These three properties of an explanation are entirely determined by the

1612 explainability method (Definition 4.1.5) used to produce it. In this context,
 1613 an explainability method is a transformation process that maps a source of
 1614 information into an explanation characterized by a specific form and subject.
 1615 Research in XAI focuses on developing new methods capable of performing this
 1616 transformation in order to overcome the lack of interpretability associated with
 1617 DL models.

1618 4.2 Taxonomy of XRL Methods

1619 Now that the concept of XAI has been defined, we turn to its application in RL,
 1620 known as XRL. As in XAI, explanations in XRL are produced by explainability
 1621 methods and can be described in terms of their subject, source, and form. This
 1622 section reviews the possible variations of these three aspects in the specific
 1623 context of RL.

1624 4.2.1 Subject

1625 The first question to address is the subject of explanation in the context of RL.
 1626 To answer this question, it is necessary to identify the different components
 1627 involved in a RL system and determine which of them constitutes the object of
 1628 explanation. A RL system can be described through two main components: the
 1629 environment and the policy.

1630 The environment defines the dynamics in which the agent evolves, including
 1631 the possible states, observations, actions, and transitions. Environments are
 1632 simulations designed by humans, and their dynamics are explicitly defined.
 1633 Therefore, it is assumed that the environment is known to the observer and does
 1634 not require explanation.

1635 The main object of interest is therefore the agent’s policy. The policy
 1636 represents the decision-making mechanism of the agent. It is a function that maps
 1637 observations to actions and is learned through interaction with the environment in
 1638 order to maximize the expected cumulative reward. The policy is the component
 1639 that encodes the learning process and that XRL aims to explain.

1640 Explaining a policy introduces additional challenges compared with classical
 1641 XAI settings. In traditional XAI problems, the objective is to explain an isolated
 1642 decision, such as determining which features of an image led to a classification
 1643 result. In contrast, RL introduces a temporal dimension into the decision-making
 1644 process. An action is not only determined by the current observation but also
 1645 influences future states and rewards. Therefore, understanding the behavior of
 1646 an agent may require considering a sequence of interactions between the agent
 1647 and its environment.

1648 This temporal dimension leads to a distinction between two objectives of
 1649 explanation in XRL:

- 1650 • **Decision explanations:** focus on identifying the elements of the current
 1651 observation that influenced the action selected by the policy. This type of
 1652 explanation is close to classical XAI approaches, as it analyzes the rela-
 1653 tionship between inputs and outputs while abstracting away the sequential
 1654 nature of the decision process.
- 1655 • **Strategy explanations:** aim to characterize the agent’s behavior over
 1656 time. They focus on sequences of decisions and interactions with the
 1657 environment in order to understand how the agent achieves its long-term
 1658 objective. This type of explanation is specific to RL, as it addresses the
 1659 temporal and goal-oriented nature of the learning process.

1660 These two objectives should not be considered as strictly separated categories,
 1661 but rather as two extremes of a continuum. Understanding the elements of an
 1662 observation that influence a specific decision can help the observer generalize
 1663 this information to infer broader aspects of the agent’s behavior. Conversely,
 1664 understanding the agent’s overall strategy can provide insights into the reasoning
 1665 behind individual decisions.

1666 4.2.2 Source

1667 We identify three sources of explanation in RL, all obtained during or after the
 1668 learning process and containing information about the agent:

- 1669 • **Policies:** policies are functions that generate a distribution over the action
 1670 space for a given observation. The objective of RL is to learn a policy that
 1671 maximizes the expected cumulative future rewards. During the learning
 1672 phase, the policy is iteratively improved through interactions with the
 1673 environment.
- 1674 • **Trajectories:** trajectories represent sequences of successive interactions
 1675 between the agent and the environment. At each step, the agent selects
 1676 an action based on its observation, which modifies the environment and
 1677 produces a reward as well as a new observation. This process continues
 1678 until the end of the trajectory.
- 1679 • **Value functions:** value functions provide predictions about the future
 1680 outcome of a trajectory from a given observation. The most commonly
 1681 used value function is the critic function, which estimates the expected
 1682 sum of future rewards. This prediction guides the learning process, as the
 1683 agent updates its policy to maximize the estimated future return. However,
 1684 other types of value functions can also be used as sources of explainability.

4.2.3 Form

We identify six forms of explanation in the RL setting:

- **Policy Model:** an interpretable model of the agent’s policy function.
- **Saliency Map:** a weighting of the input features representing their influence on the output of the policy function.
- **Counterfactual:** a modified observation obtained by perturbing the original observation in order to produce a different action from the agent.
- **Agent Trajectory:** a sequence of successive interactions between the agent and the environment.
- **Trajectory Prediction:** a prediction about the future evolution of the trajectory produced by a value function.
- **Interaction Model:** an interpretable model of the interactions between the agent and the environment.

Among the three components characterizing an explanation, the form is the one that varies most significantly across explainability methods. While subjects and sources tend to be relatively stable in the XRL literature, the diversity in how explanations are presented is considerably greater. The list above is by no means definitive and may evolve as new approaches emerge. For this reason, the following section is organized according to the form of explanations and presents existing explainability methods based on this criterion.

4.3 State-of-the-Art Review of XRL Methods

The purpose of this section is to detail the functioning of existing explainability methods (see Table 4.1). They are presented according to the form of explanation they produce, following the taxonomy introduced in Section 4.2.

A family of methods is defined as a set of methods that share the same source, form, and subject of explanation, as well as a similar procedure for generating explanations. Methods belonging to the same family may differ in their implementation or in the details of the generation process, while relying on the same underlying principle.

The remainder of this review is organized by form of explanation, following the six forms identified in Section 4.2.3. For each form, we first describe its underlying principle and its limitations, then present the families of methods that produce it.

4.3.1 Policy Model

The policy function can be represented by a wide variety of function classes. In modern RL, NN is the most commonly used models for representing policies. Although these models achieve the highest performance, their connectionist nature makes their internal representations non-interpretable.

For this reason, a substitutive approach to XRL aims to replace these models with alternative models considered interpretable. Such models can

Explanation Subject	Explanation Form	Explainability Method Family	Source of Explanation		
			Policy	Trajectories	Value functions
Decision	Policy Model	<i>Interpretable policy learning</i>	✓		
		<i>Policy approximation via interpretable model</i>	✓		
	Saliency Map	<i>Observation perturbation</i>	✓		
		<i>Observation gradient analysis</i>	✓		
		<i>Attention-based architecture</i>	✓		
	Counterfactual	<i>Latent space usage for counterfactual generation</i>	✓		
Strategy	Agent Trajectory	<i>Hierarchical decomposition</i>	✓	✓	
		<i>Extraction via critic analysis</i>		✓	✓
	Trajectory Prediction	<i>Critic enrichment</i>			✓
	Interaction Model	<i>Bayesian network learning</i>		✓	
		<i>Latent space usage for MDP generation</i>	✓		

Table 4.1: Taxonomy of explainability methods in reinforcement learning.

be understood through direct analysis of their representations. The main classes of interpretable models used to represent policy functions include decision trees [Topin et al., 2021], algorithms [Trivedi et al., 2022], formal logic [Payani and Fekri, 2019], and linear models [Garreau and von Luxburg, 2020].

Explanations derived from an interpretable model primarily provide insights into the decision-making process, since what is modeled is the mapping between an observation and an action. If the model is sufficiently simple or well structured, the observer can form a mental representation of a sequence of decisions across the environment and thus obtain information about the agent’s strategy.

The main limitation of these methods is that interpretable models rely on symbolic representations. When the policy to be represented is too complex, achieving comparable performance requires a large number of rules or conditions. The resulting model can quickly become difficult to analyze and lose its interpretability.

For example, representing a policy capable of playing Go using a symbolic model would require considering a very large number of possible situations and specific cases. Even if such a representation were possible, the resulting model would likely become intractable for human analysis due to its complexity. This illustrates why NN-based approaches have become dominant in such domains, as they can learn complex decision strategies without requiring an explicit enumeration of rules.

To obtain an interpretable policy model, two families of explainability methods are available: *interpretable policy learning* and *policy approximation via interpretable model*.

Interpretable Policy Learning

This family of methods aims at learning a policy represented by an interpretable model [Topin et al., 2021, Trivedi et al., 2022]. It is the most direct approach to




Figure 4.3: Policy approximation via an interpretable model [Sieusahai and Guzdial, 2021].

addressing the challenges posed by XRL. This type of training requires modifying the Markov Decision Processes (MDP) in order to adapt it to the search space and to the type of policy representation. Moreover, this search space often needs to be reduced to make learning feasible. Such a reduction requires the integration of expert knowledge into the learning process. With this family of methods, the *interpretable policy function* itself is used to provide explanations to the observer.

Policy Approximation via Interpretable Model

An interpretable model can be trained to approximate a policy learned through deep RL [Sieusahai and Guzdial, 2021]. In this setting, the problem is formulated as a supervised learning task, where the original deep RL policy is used to label each observation with the action it selects, such as the observation from the Atari game “Pong” shown in Figure 4.3a. Using this dataset, the interpretable model is trained to reproduce the behavior of the original policy, yielding for instance the decision tree illustrated in Figure 4.3b, obtained by approximating a deep RL policy trained on that environment. The resulting interpretable model is then used as an explanation for the observer. If this approximation is sufficiently faithful, it can also be deployed in the environment in place of the original policy, thereby providing a high degree of control over the agent.

4.3.2 Saliency Map

The influence of each part of the input on the action selected by the policy can be visualized using heatmaps (Figure 4.4). In the RL setting, the input corresponds to the observation received by the agent at a given time step. Heatmaps therefore highlight the regions of the observation that have the greatest influence on the action selection process. By providing this information, they allow the observer to identify which elements of the environment are considered important by the policy when making a decision. In Figure 4.4, both panels correspond to the same observation of the Atari game “Breakout”, with panel (a) obtained by observation perturbation and panel (b) by observation gradient analysis.




Figure 4.4: Heatmaps for the same observation of the Atari game “Breakout” [Greydanus et al., 2018].

1781 A limitation of heatmap-based methods is that the resulting explanations
 1782 are not always directly meaningful to the observer. As illustrated in Figure 4.4b,
 1783 interpreting which regions are important often requires human reasoning and
 1784 involves a degree of subjectivity. This limitation becomes particularly relevant
 1785 in RL, where the objective is not only to explain a single decision but also to
 1786 understand the factors influencing a sequence of decisions over time.

1787 Three families of explainability methods can be used to obtain this type
 1788 of explanation: *observation perturbation*, *observation gradient analysis*, and
 1789 *attention-based architectures*.

1790 **Observation Perturbation**

1791 Perturbation analysis consists in applying various modifications to the obser-
 1792 vation in order to study their impact on the output of the policy function
 1793 [Greydanus et al., 2018]. The greater the variation in the output caused by
 1794 a perturbation applied to a given component of the input vector, the higher
 1795 the weight assigned to that component. The magnitude of the perturbation
 1796 is measured by computing the distance between the original output and the
 1797 output obtained from the perturbed input. Using this set of weights, a *heatmap*
 1798 is constructed, enabling the observer to visualize which parts of the observation
 1799 vector are most sensitive in the action-selection process (see Figure 4.4a).

1800 **Observation Gradient Analysis**

1801 Observation gradient analysis refers to a set of methods that require a differen-
 1802 tiable policy function in order to extract information from it [Simonyan et al., 2014,
 1803 Weitkamp et al., 2019, Huber et al., 2019]. To this end, the gradient of the pol-
 1804 icy function is computed with respect to all components of the input vector.
 1805 For each input component, the resulting gradient provides a weighting that




Figure 4.5: Attention-based heatmaps for the Atari games “Breakout” and “Ms. Pac-Man” [Mott et al., 2019].

1806 represents its importance in the output of the policy function. Based on these
 1807 weightings, a *heatmap* is constructed to allow the observer to visually identify
 1808 the parts of the input vector that are the most influential in the policy function’s
 1809 output (see Figure 4.4b).

1810 *Attention-Based Architecture*

1811 Attention blocks constitute a NN architecture that makes it possible to weight
 1812 the relative importance of one part of a vector with respect to the other parts
 1813 of the same vector [Vaswani et al., 2023]. These weights are computed as a
 1814 function of the vector itself and are learned by the NN during training. For
 1815 a given observation, each component of the observation vector is assigned a
 1816 scalar value. It can therefore be assumed that the components associated
 1817 with the highest weights are the most influential in determining the output
 1818 [Mott et al., 2019, Itaya et al., 2021]. This attention-based *heatmap* is then
 1819 provided to the observer as the explanation, as illustrated for the Atari game
 1820 “Breakout” in Figure 4.5a, where the highlighted regions are easily interpretable.
 1821 However, as shown for “Ms. Pac-Man” in Figure 4.5b, the attention weights are
 1822 not always concentrated on a small, easily interpretable set of regions, which
 1823 can make this type of explanation difficult to exploit.

1824 4.3.3 Counterfactual

1825 A counterfactual explanation is a new observation, close to the original one, that
 1826 leads the agent to select a different action. Comparing the two observations
 1827 reveals which elements of the environment must change for the agent to behave
 1828 differently, and thus provides insight into its decision-making process. Unlike
 1829 the other forms of explanation, this form is currently associated with a single
 1830 family of methods, based on the *use of the Latent Space (LS) for counterfactual*




Figure 4.6: Counterfactual explanation of the agent’s behavior for the Atari game *Space Invaders* [Olson et al., 2021]. Panel a shows the original observation and panel b the counterfactual observation produced by a generative model.

generation.

Latent Space Usage for Counterfactual Generation

This family of methods uses generative models to build counterfactuals from the Latent Representation (LR) of the observation [Bond-Taylor et al., 2021, Olson et al., 2021]. The observation is encoded into the LS, its representation is modified until the agent selects the desired action, and this modified representation is decoded back into a new observation. Working in the LS rather than at the pixel level keeps the counterfactual close to the original state and consistent with the environment. The fidelity of the generated observations to the real game can then be assessed through user studies [Olson et al., 2021].

Figure 4.6 illustrates this approach on the Atari game *Space Invaders* [Olson et al., 2021]. In the original observation (Figure 4.6a) the agent selects “LEFT”, whereas in the counterfactual observation (Figure 4.6b) it selects “RIGHT_FIRE”. Here, removing an enemy is enough to change the action, which suggests that the position of that enemy influences the agent’s decision. Such observations remain only indicative. Identifying a reliable pattern requires repeating the analysis across many situations.

4.3.4 Agent Trajectory

Until now, the presented methods have focused on explaining individual decisions made by the agent. The methods considered from this point on aim to provide insight into the agent’s strategy by analyzing its behavior over time.

The most direct approach consists in visualizing agent trajectories, most commonly through complete episodes. By analyzing these trajectories, the observer can infer the dynamics underlying the agent’s behavior and how its decisions evolve through time. This type of visualization is not specific to XRL. In practice, visualizing agent episodes is a standard step in the development and evaluation of RL. It provides a simple way to verify that the agent behaves as expected and to detect potential implementation errors.

This simple approach can be further enhanced using families of methods such as *hierarchical decomposition* and *extraction via critic analysis*.

Extraction via Critic Analysis

In this family of methods, the predictions of the critic are used to determine which trajectories to present to the observer [Sequeira and Gervasio, 2020]. A set of states is selected from multiple trajectories. For a given state, all possible actions are examined. The state is assigned a value equal to the difference between the critic’s prediction for the lowest-valued action and the prediction for the highest-valued action. Trajectories containing the states with the highest values according to this procedure are then chosen to be presented to the observer. These trajectories are considered critical moments of the agent and are therefore relevant for analysis. For this family of methods, the *selected trajectories* represent the interpretable information provided to the observer.

Hierarchical Decomposition

Hierarchical agents are composed of a set of sub-policies specialized for specific tasks within subsets of the environment’s states. This family of methods distributes the optimization of the reward function across multiple models, significantly simplifying the learning process. Moreover, this decomposition of the policy function can also provide explainability in the agent’s strategy [Beyret et al., 2019]. If a sub-policy is selected, it indicates that the agent will execute a specialized behavior. This behavioral specialization provides insight into the strategy the agent employs to maximize the reward function. The agent’s *trajectory*, together with the *sub-policy* selected for each action, is used as interpretable information.

4.3.5 Trajectory Prediction

Making predictions about the future of a trajectory allows one to form a representation of what the system expects, based on its training experience. These predictions are produced using value functions. Such functions provide insights into the future evolution of the agent’s trajectory. By their nature, this form of explanation informs the observer about the strategy adopted by the agent. Value functions provide an estimate of the expected cumulative reward from a given state. Consequently, a value function maps an observation to a scalar value in $\mathbb{R}$.

This aggregated prediction provides limited information about the underlying factors that contribute to this value. It does not directly reveal which aspects of the task are being optimized or how the predicted return is composed. Unlike the other forms of explanation, this form is currently associated with a single family of methods, aimed at improving the interpretability of value-based explanations: *critic enrichment*.

Critic Enrichment

In order to obtain interpretable information from the predictions of the critic, it is possible to decompose the critic into multiple sub-functions [Juozapaitis et al., 2019]. The reward function is thus decomposed into multiple sub-reward functions, each




Figure 4.7: Prediction of sub-rewards as a function of the action selected by the agent for a given state [Juozapaitis et al., 2019].

representing a sub-objective of the task to be accomplished. Similarly, for each of these sub-reward functions, corresponding sub-critic functions are defined, along with a function that combines their outputs. During training, the agent selects its actions so as to maximize this combination function of the sub-critics. Using this approach, for each observation and action, a prediction is obtained regarding the sub-objectives the agent is attempting to maximize, so that the action chosen by the agent reveals which sub-reward it aims to maximize (see Figure 4.7). This *prediction* is then used as the explanation for the observer.

4.3.6 Interaction Model

Modeling the interactions between an agent and its environment makes it possible to understand how the agent’s decisions are made and what impact they have on the environment. By incorporating these interactions into an explanatory model, one can better grasp the complex relationships between the agent and its environment, thereby identifying the determining factors in the decision-making process. By explicitly handling the agent–environment interaction, this type of explanation aims to explain the agent’s strategy. The families of methods used to obtain explanations in the form of agent–environment interactions are: *Bayesian network learning* and *LS usage for MDP generation*.

Bayesian Network Learning

The evolution of an agent within an environment can be observed as a set of random variables connected through a Bayesian network [Madumal et al., 2020]. Random variables are defined to represent both external aspects (the environment) and internal aspects of the agent (the actions). By analyzing trajectories, it becomes possible to infer causal relationships among these random variables. At the end of this process, a weighted causal graph is obtained, which serves as an explanation for the observer, as illustrated for the video game *StarCraft II* in Figure 4.8, where nodes represent random variables and edges represent causal




Figure 4.8: Agent-environment interaction model for the video game *StarCraft II* [Madumal et al., 2020].

relationships between them. This *Bayesian network*, representing the interactions between the agent and the environment, is then used as the explanation.

Latent Space Usage for MDP Generation

Each layer of a NN can be viewed as a projection from one space to another [Zahavy et al., 2017]. As information propagates through successive layers of a NN, it becomes increasingly abstract. It can therefore be assumed that two vectors that are close in this projected space are also close in terms of the agent's perception. Based on this observation, it is possible to redefine an MDP by leveraging the abstract representations produced by the final layers of a NN. Once the new MDP has been constructed, a model of the interaction between the agent and its environment is obtained. This model, represented as a *graph*, is then used as the explanation.

Figure 4.9 illustrates this idea using a two-dimensional *t-SNE* projection of the LS of the NN used to represent the policy. Each point is colored according to the value function's prediction for the corresponding state, with warmer colors indicating higher predicted values. An agreement can be observed between the structure of this projected space and the critic's predictions: states that are assigned similar values by the critic also tend to be located close to one another in the projection. This agreement makes it possible to manually identify clusters of states, which likely correspond to states that are perceived as similar by the agent.

4.4 Conclusion

The taxonomy and the state-of-the-art review presented in this chapter make it possible to draw three main observations.

4.4.1 Absence of Explanation of the Learning Process

Regarding the subject of explanation, Table 4.1 shows that none of the reviewed methods addresses the learning process itself. Such an explanation would concern




Figure 4.9: t -SNE projection of a neural network’s latent space [Zahavy et al., 2017].

1956 how the policy is progressively updated throughout training, and why it evolves
 1957 over successive iterations to converge toward a particular solution. Several
 1958 reasons explain this absence.

1959 Among the different machine learning paradigms, RL is arguably the one that
 1960 would benefit the most from explanations of the learning process. Unlike super-
 1961 vised or unsupervised learning, it relies on an explicit exploration–exploitation
 1962 trade-off, making the learning trajectory itself an informative object of explana-
 1963 tion. Most XRL methods, however, are inherited from the XAI literature, where
 1964 explanations focus almost exclusively on the behavior of already-trained models.
 1965 This limitation has naturally carried over to XRL.

1966 From a practical perspective, explaining the learning process is also of limited
 1967 industrial interest. In most real-world applications, practitioners deploy an
 1968 already-trained agent, while training itself takes place offline. Users are generally
 1969 interested in understanding the behavior of the final policy rather than the
 1970 optimization process that produced it.

1971 Explaining the learning process also raises significant methodological chal-
 1972 lenges. Training in RL can be viewed as the iterative improvement of a policy
 1973 through repeated interactions with the environment. In this sense, explaining
 1974 learning would amount to explaining the evolution of an entire sequence of
 1975 policies rather than a single one. Current XRL methods still struggle to provide
 1976 comprehensive explanations of an individual policy, so extending them to a
 1977 full training trajectory appears considerably more difficult. Explanations of

the learning process therefore remain an open research direction rather than a straightforward extension of existing XRL methods.

4.4.2 Lack of Unified Metrics for Explainability

The diversity of explanation forms identified in this review also reveals another limitation shared by nearly all families of XRL methods, which is the absence of reliable metrics for evaluating explainability. In many studies, explanations are not evaluated with human participants. Instead, authors focus on technical properties of the proposed method, such as fidelity, sparsity, or computational efficiency. When human evaluations are conducted, they usually rely on subjective satisfaction questionnaires. Although useful, self-reported satisfaction remains prone to numerous biases and does not necessarily indicate that the explanation has improved the observer’s mental model or understanding of the agent.

This limitation is not specific to XRL, it reflects a broader challenge affecting XAI as a whole. In our view, this challenge stems from a historical development largely driven by the computer science community. Explainability has consequently often been studied from an algorithmic or mathematical perspective, with less attention given to how humans actually benefit from explanations.

Our definition of explanation explicitly places the observer at the center of the explanatory process. From this perspective, the value of an explanation should not only depend on its technical properties but also on its ability to improve the observer’s understanding. However, the field still lacks objective and standardized metrics capable of measuring this aspect.

Addressing this challenge will require stronger interdisciplinary collaboration between computer science and disciplines such as cognitive science, psychology, philosophy, and sociology. These fields provide the theoretical and experimental tools needed to study human understanding and to develop evaluation protocols that measure the actual effectiveness of explanations. The absence of such collaborations is likely a consequence of the relative youth of XAI as a research field, whose theoretical foundations and evaluation methodologies are still being established.

4.4.3 Toward Intrinsic Explainability in XRL

Our review also suggests that current XRL research remains largely dominated by substitutive explainability approaches (Section 4.1). Many existing methods explain RL agents by constructing surrogate models. While these approaches can provide useful explanations for simple policies, their ability to scale to increasingly complex agents is limited. This limitation is particularly important in deep RL, where NN are used to represent complex policies that would be difficult to approximate with simpler surrogate models. Replacing such a model with a simpler one therefore risks losing precisely the complexity that makes deep RL effective.

For this reason, we argue that a promising direction for XRL is to develop explanations directly from the NN underlying the agent. Rather than bypassing the model, such approaches aim to analyze the representations learned by the network itself. From this perspective, the NN should not necessarily be regarded as an inherently opaque black box, but as a complex representation whose apparent opacity may instead reflect the limitations of current analysis tools.

2024 Among the methods reviewed, *latent space usage for MDP generation* (Sec-
2025 tion 4.3.6) appears closest to this intrinsic perspective. Instead of replacing
2026 the learned representation, it directly exploits the LS to identify meaningful
2027 concepts that can support the explanation of the agent’s behavior. However, this
2028 approach still relies on a largely manual and qualitative analysis of the LS, which
2029 limits its reproducibility and its ability to produce systematic explanations. The
2030 next chapter builds upon this idea and proposes a mechanistic XRL approach
2031 based on LS analysis.

Chapter 5

Clustering Agent Representations

As discussed in the previous chapter, current explainability methods for Reinforcement Learning (RL) still lack intrinsic approaches capable of explaining the internal mechanisms of deep neural policies. To the best of our knowledge, only [Zahavy et al., 2017] have proposed exploiting the Latent Space (LS) learned by neural policies in order to identify the concepts used by the agent. Its main idea is to analyze the Latent Representation (LR) produced by the Neural Network (NN) and organize them into clusters that correspond to high-level concepts.

While the core idea of this paper is closely aligned with our vision of explainability in RL, this approach suffers from two important limitations. First, the clusters are manually constructed. Such a procedure heavily depends on the user’s expertise and neither guarantees reproducibility or scales. Second, no evaluation protocol is proposed to assess the quality of the obtained clusters. Their relevance is essentially determined through visual inspection.

The Clustering Agent Representation (CAR) method proposed in this chapter addresses these two limitations. It automates concept discovery by relying on unsupervised clustering algorithms to identify the Concept Structure (CS) of the LS. It introduces an objective evaluation metric based on the stability of the extracted representations across independently trained policy.

This chapter is organized as follows:

- Section 5.1 lays out the theoretical assumptions underlying the proposed CAR method, which we term the *Representation Convergence Hypothesis* (Hyp. 4).
- Section 5.2 then details the CAR method itself, covering both the extraction of LR using surrogate policies and the automatic detection of CS through clustering.
- Section 5.3 introduces the metrics we propose for evaluation.
- Section 5.4 reports the results obtained by applying CAR and these metrics across several Atari environments.
- Section 5.5 discusses these results, situates CAR within the broader landscape of mechanistic interpretability, and outlines directions for future work.

5.1 Representation Convergence Hypothesis

The CAR method relies on a assumption, referred to as the *Representation Convergence Hypothesis* (Hyp. 4). This hypothesis follows the two assumptions established previously throughout this manuscript.

The *Mechanistic Hypothesis* (Hyp. 3) rejects the notion that deep NN constitute irreducible black boxes. Adopting a mechanistic perspective, we consider that every process results from a sequence of causal computations. Although these computations may be complex, they remain explainable through the analysis of the internal mechanisms.

The *Concept Hypothesis* (Hyp. 1) states that NN achieve generalization by learning concepts. During optimization, the network progressively organizes its LS into representations that capture the information required to solve the task.

Following these hypotheses, it becomes possible to explain the NN’s process by understanding the organization of its LS. This family of approaches is referred to as mechanistic interpretability. In mechanistic interpretability, universality refers to the consistency with which a concept is recovered across different models or extraction processes [Templeton et al., 2024].

To provide a explanation for this phenomenon, we developed the *Representation Convergence Hypothesis* (Hyp. 4). This hypothesis states that, when NN are trained to solve the same task, a set of concepts should systematically emerge because these concepts provide a effective way of representing the problem. As a result, independently trained models should develop similar concepts in their learned representations.

Hypothesis 4: Representation Convergence Hypothesis

For a given task, NN tend to converge toward a privileged CS because it provides an efficient representation for solving the task.

It is important to clarify that this hypothesis does not imply convergence at the parameter level. Independently trained models are expected to converge

toward different sets of weights. The proposed hypothesis concerns a abstract level of representation. We assume that the conceptual organization induced by their LR converges toward an equivalent CS, even though the numerical implementation of the networks differ.

The CAR method is designed to investigate this hypothesis by searching for concepts that are consistently recovered across independently trained models. In this sense, the hypothesis serves both as the foundation of the method and as the phenomenon that the method aims to investigate. If independently trained models develop common concepts for a given task, CAR should be able to recover these concepts consistently across models. Conversely, the absence of such convergence would provide evidence against the hypothesis. This also provides a stronger justification for studying these concepts, as concepts consistently recovered across different training runs are more likely to reflect properties of the task than model-specific artifacts.

The remainder of this chapter develops the methodology used to investigate this hypothesis. The *Clustering Universality* metric, introduced in Subsection 5.3.3, provides a quantitative measure of the convergence observed between the concepts identified in independently trained models.

5.2 Method

Following the work of [Zahavy et al., 2017], we analyze the LS through clustering techniques, placing our approach within the broader field of Deep Clustering. This design choice is motivated by our rejection of the *Superposition Hypothesis* (Hyp. 2) and, more generally, of any view that assumes a linearly organized CS, as discussed in Section 2.5.


The CAR methodology is composed of two main stages. The first stage trains multiple independent NN, called surrogate policies, which implement the same decision policy. The second stage clusters their LR in order to identify the underlying conceptual organization.

This section is organized as follows. In the first subsection, we justify the use of multiple surrogate policies in order to study the stability of representations. In the second subsection, we detail how these surrogate policies are trained so as to best match the original policy.

5.2.1 Multiple surrogate policies

The CAR method was developed to test the *Representation Convergence Hypothesis* (Hyp. 4). This requires training several NN independently and comparing their LR. Consistent similarities in the resulting CS would provide empirical support for this hypothesis.

In supervised learning, this protocol is straightforward since several models can be independently trained on the same target function. RL introduces an additional difficulty. Because of stochastic exploration, independently trained agents are not guaranteed to converge toward the same policy. Even if they achieve comparable performance, they may solve the task using different strategies. This makes it difficult to determine whether differences in the LS result from different representations of the same strategy or from differences in the learned strategies themselves.

Figure 5.1: Example of a surrogate policy architecture π' .

2137 To overcome this limitation, we consider an already trained initial policy.
 2138 This policy is then imitated using imitation learning [Ross et al., 2011]. In
 2139 imitation learning, several surrogate NN are independently trained in a supervised
 2140 manner to reproduce the behavior of the initial policy. This makes it possible
 2141 to investigate whether equivalent policies also exhibit equivalent conceptual
 2142 organizations.

2143 It is important to emphasize that the use of surrogate policies does not make
 2144 CAR a surrogate explainability method, as defined in Subsection 4.1.2. The
 2145 surrogate policies remain deep NN. Their role is not to simplify the decision
 2146 process but to provide multiple independent realizations of the same policy
 2147 function, enabling the study of representation convergence.

2148 5.2.2 Training surrogate policies

2149 Throughout the remainder of this chapter, the initial policy we aim to ex-
 2150 plain is denoted as π . This initial policy is obtained without imposing any
 2151 constraints on its origin. It may result from classical or deep reinforcement learn-
 2152 ing [Sutton and Barto, 2018], alpha-beta algorithm [Knuth and Moore, 1975],
 2153 or even from demonstrations provided by humans [Ross and Bagnell, 2010]. The
 2154 crucial requirement is the availability of a dataset containing a sufficient number
 2155 of projections from the observation space $\mathcal{O}$ to the action space $\mathcal{A}$ induced by π
 2156 to train a NN.

2157 The replicated policies are referred to as surrogate policies, denoted by π' .
 2158 For any observation O , the distribution over the action space produced by π'
 2159 should approximate that of π as closely as possible.




Figure 5.2: Illustration of the Clustering Agent Representation method.

2160 The surrogate policy π' is implemented as a NN parameterized by Θ . This
 2161 network is composed of two components: an encoder and a decoder, respectively
 2162 highlighted by the blue and black dashed rectangles in Figure 5.1. The encoder
 2163 f_{enc} maps the observation O to the embedding. The decoder, which reconstructs
 2164 the action distribution parameters from the embedding.

2165 To train π' to accurately replicate the action distribution, we define a loss
 2166 function $\mathcal{L}$ optimized via gradient descent as introduced in Section 1.4. This
 2167 loss, defined in Equation (5.1), is computed as the MSE between the parameters
 2168 of the action distributions, obtained by applying the function p , which maps a
 2169 distribution to its parameterization. In practice, this function does not need to
 2170 be explicitly defined, as neural policies directly output the parameters of the
 2171 action distribution. Once training is complete, this loss ensures that π' closely
 2172 approximates the behavior of the original policy π .

$$\min_{\Theta} \sum_{O \in \mathcal{O}} \mathcal{L}(O) = \min_{\Theta} \sum_{O \in \mathcal{O}} \text{MSE}\left(p(\pi(O)), p(\pi'_\Theta(O))\right) \quad (5.1)$$

2173 Once training is complete, we can identify clusters within the LS. To do so,
 2174 we denote by $\mathcal{E}$ the set of embeddings obtained by applying f_{enc} to a subset
 2175 of observations $\mathbb{O} \subseteq \mathcal{O}$. On the set $\mathcal{E}$, we use a clustering function f_{clust} to
 2176 partition the point cloud of embedded observations into distinct clusters. Each
 2177 observation is assigned a cluster label based on its proximity in the LS, resulting
 2178 in the clustering $\mathcal{C}$.

2179 Figure 5.2 and Algorithm 5 summarize the complete CAR procedure. Fig-
 2180 ure 5.3 illustrates the analysis of a latent space obtained by applying the CAR
 2181 method. The point cloud shows a two-dimensional projection of the LS of a
 2182 surrogate policy trained on Space Invaders. The two-dimensional projection is
 2183 used only for visualization, while the clusters are computed in the original high-
 2184 dimensional LS. Colors indicate the resulting clusters and their corresponding

Algorithm 5: Clustering Agent Representation (CAR)**Input:**

Original policy π
 Observation dataset $\mathbb{O}$
 Number of surrogate policies n
 Surrogate policy architecture π' (encoder f_{enc} , decoder)
 Learning rate η
 Clustering function f_{clust}

Output:

Trained surrogate policies $\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}$
 Clusterings $\mathcal{C}_1, \dots, \mathcal{C}_n$

Behavioral cloning dataset built from the target policy π

1 $\mathcal{D} \leftarrow \{(O, p(\pi(O))) \mid O \in \mathbb{O}\};$

MSE loss between action-distribution parameters, as defined in Equation (??)

2 $\mathcal{L} \leftarrow \text{Mean Squared Error (MSE)};$

Independently train n surrogate policies by behavioral cloning

3 **for** $i = 1, \dots, n$ **do**

4 Initialize Θ_i randomly;

Behavioral cloning of π via SGD, see Algorithm 1

5 $\Theta_i \leftarrow \text{SGD}(\mathcal{D}, \pi', \mathcal{L}, \eta, \Theta_i);$

Cluster the latent representations of each surrogate policy

6 **for** $i = 1, \dots, n$ **do**

Embeddings produced by the encoder of π'_{Θ_i}

7 $\mathcal{E}_i \leftarrow \{f_{enc_i}(O) \mid O \in \mathbb{O}\};$

Partition of the embedded observations into clusters

8 $\mathcal{C}_i \leftarrow f_{clust}(\mathcal{E}_i);$

9 **Return** $(\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}), (\mathcal{C}_1, \dots, \mathcal{C}_n)$

2185 observations. Each cluster corresponds to a distinct, human-interpretable stage
 2186 of the game. The resulting organization suggests a correspondence between the
 2187 position of observations in the LS and the semantic properties of the correspond-
 2188 ing game states. To objectively assess this correspondence we introduce a set of
 2189 evaluation metrics detailed in the following section.




Figure 5.3: Two-dimensional projection of the LS embeddings of a surrogate policy trained on Space Invaders, together with the CAR clusters and their corresponding observations.

5.3 Evaluation Metrics

A key contribution of this work is the development of three evaluation metrics, designed to assess the effectiveness of CAR:

- **Behavior Reconstruction** : it measures how well a surrogate policy faithfully reproduces the behavior of the original policy by comparing their cumulative rewards.
- **Clustering Universality** : it evaluates the stability of clusterings by measuring the similarity of clusters obtained from independently trained surrogate policies.
- **Abstraction Score** : it quantifies the extent to which LS clusters capture abstractions rather than directly reflecting the structure of observations or actions.

The remainder of this section details the motivation for each metric and describes how they are calculated. In the first subsection, we develop the

Behavior Reconstruction. We then introduce the Adjusted Rand Index (ARI) measure, which serves as the foundation for the two other metrics introduced in the subsequent subsections, namely *Clustering Universality* and *Abstraction Score*.

The following notations will be adopted for the rest of the manuscript: we consider n surrogate policies, denoted as $\pi'_1, \dots, \pi'_n$, trained as described in Section 5.2. Each surrogate policy is associated with a projection function $f_{enc_1}, \dots, f_{enc_n}$, which is applied to the set of observations $\mathbb{O}$, yielding embeddings $\mathcal{E}_1, \dots, \mathcal{E}_n$. A clustering function is then applied to each embedding set, producing partitions $\mathcal{C}_1, \dots, \mathcal{C}_n$.

5.3.1 Behavior Reconstruction

We argue that explaining a function b that closely approximates another function a is essentially the same as explaining a . If b accurately captures the key decision patterns of a , understanding b provides meaningful insights into how a works [Guidotti et al., 2018]. In our setting, this means that generating explanations for π' that faithfully reproduce π amounts to explaining π directly. To verify that π' is indeed a faithful copy, we introduce the *Behavior Reconstruction* metric.

This metric is computed by comparing the cumulative rewards obtained by the original policy π and the surrogate policies π' over m evaluation episodes. The behavior reconstruction metric is then defined as:

$$\text{Behavior Reconstruction} = \frac{1}{n \sum_{i=1}^m g_{\pi}^i} \sum_{j=1}^n \sum_{k=1}^m g_{\pi'_j}^k \quad (5.2)$$

In imitation learning, it is unlikely to expect the surrogate policy π' to outperform the initial policy π it aims to approximate through supervised learning [Ross and Bagnell, 2010]. Consequently this metric ranges in $[0, 1]$. Values approaching 1 indicate minimal loss in reproducing π , while values close to 0 indicate poor reproduction.

We deliberately adopt a reward-based definition rather than alternatives based on actions or action distributions (e.g., training loss, accuracy, KL divergence [Shlens, 2014], or optimal transport [Peyré and Cuturi, 2020]). Since surrogate policies are trained in a supervised manner, they cannot develop strategies beyond those demonstrated by the initial policy. This means that if the rewards achieved are comparable, the surrogate policy is effectively copying the original policy’s behavior.

5.3.2 Adjusted Rand Index

Within CAR, concepts are assumed to correspond to clusters in the LS. Combined with the *Representation Convergence Hypothesis* (Hyp. 4), this leads to a direct experimental prediction: if independently trained surrogate policies converge toward the same conceptual organization, then clustering their LR should produce similar partitions.

When a particular observation is processed by the network, it is expected to activate the same concept regardless of the specific neural implementation of the policy. If independently trained policies converge toward the same conceptual

organization, a given observation should activate the same concept across all models. As a consequence, observations grouped together within a cluster for one policy should also be grouped together for another independently trained policy.

In the clustering literature, the most used measures for comparing two partitions is the ARI [Hubert and Arabie, 1985]. This measure quantifies the agreement between cluster assignments independently of the spatial organization of the clusters. Both metrics introduced later in this chapter, *Clustering Universality* and *Abstraction Score*, rely on the ARI. Since the ARI is a complex measure, we review its definition and discuss its main properties.

The ARI is derived from the Rand Index (RI) [Rand, 1971], which measures the agreement between two partitions by considering every possible pair of point cloud. For each pair, four situations are possible: the point may belong to the same cluster in both partitions, to different clusters in both partitions, or they may disagree by being grouped together in only one of the two partitions. The RI measures the proportion of pairs for which the two partitions agree:

$$\text{RI} = \frac{\text{number of agreeing pairs}}{\text{total number of pairs}}.$$

Although the RI is intuitive, it suffers from an important limitation. In most clustering problems, the majority of pairs of points belong to different clusters. Consequently, even two completely independent partitions will agree on a large number of pairs, simply because both assign them to different clusters. This effect becomes more pronounced as the number of points or clusters increases.

The ARI addresses this issue by accounting for the agreement that would be expected purely by chance. It compares the observed agreement RI to the expected agreement under random partitions $\mathbb{E}[\text{RI}]$. The score is then normalized by the maximum agreement achievable beyond this random baseline, $\max(\text{RI}) - \mathbb{E}[\text{RI}]$. As a result, an ARI of 0 corresponds to the level of agreement expected from random clusterings, while an ARI of 1 indicates perfect agreement between the two partitions:

$$\text{ARI} = \frac{\text{RI} - \mathbb{E}[\text{RI}]}{\max(\text{RI}) - \mathbb{E}[\text{RI}]}$$

In practice, the ARI is computed directly from the cluster assignments of the two partitions. Let clustering A be composed of partitions $\{A_i\}$ and clustering B of partitions $\{B_j\}$. Let $a_i = |A_i|$ denote the number of samples in partition i of clustering A , and $b_j = |B_j|$ the number of samples in partition j of clustering B . Furthermore, let $n_{ij} = |A_i \cap B_j|$ be the number of samples shared by partition i of clustering A and partition j of clustering B , and let n be the total number of samples. With these definitions, the ARI can be computed as:

$$\text{ARI} = \frac{\sum_{i,j} \binom{n_{ij}}{2} - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}{\frac{1}{2} \left[\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2} \right] - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}$$

Because it is bounded by 1, a value such as 0.5 can be misleadingly interpreted as “halfway to a perfect match.” This interpretation is incorrect because the ARI is normalized against the agreement expected from two random partitions. As a result, values close to 1 require the two clusterings to be nearly identical, while a value of 0.5 already indicates substantial agreement [Amezquita et al., 2020].




Figure 5.4: Illustration of *Clustering Universality* in the ideal case where projections of the same observations are grouped within the same clusters across all surrogate policies.

5.3.3 Clustering Universality

Having introduced the ARI, we can now define the *Clustering Universality* metric. While the ARI quantifies the similarity between two conceptual partitions, the objective of CAR is to assess the consistency of conceptual organizations across an entire set of independently trained surrogate policies.

To this end, the same set of observations is projected into the LS of each surrogate policy and clustered. Under the *Representation Convergence Hypothesis* (Hyp. 4), these independently obtained partitions are expected to be similar. *Clustering Universality* quantifies this by computing the mean pairwise ARI:

$$\text{Clustering Universality} = \frac{1}{\binom{n}{2}} \sum_{i=1}^n \sum_{j=0}^{i-1} \text{ARI}(\mathcal{C}_i, \mathcal{C}_j). \quad (5.3)$$

This metric takes values in $[-1, 1]$. A value close to 1 indicates that surrogate policies produce highly similar partitions, suggesting that the clusters reflect intrinsic properties of the policy. Conversely, a value near 0 implies that the similarity between partitions is no better than that of two randomly generated clusterings, meaning that the clustering fails to capture any meaningful structure in the policy. A negative value, on the other hand, signals that the agreement between partitions is worse than what would be expected by chance, indicating a systematic disagreement between the clusterings.

Figure 5.4 illustrates *Clustering Universality* in the ideal case. It shows three LS obtained by projecting the same seven observations, labeled 1 to 7, through three independently trained surrogate policies. Although the embeddings occupy different positions, the clusters group exactly the same elements in all three cases. The partitioning is therefore identical, resulting in an ARI of 1.

2307 5.3.4 Abstraction Score

2308 If *Clustering Universality* is high, it becomes important to determine what
2309 underlies this stability. Two possibilities arise:

- 2310 (i) the LS merely reproduces the structure of the observation space or action
2311 distribution, offering no abstraction, as defined in Definition 2.1.1;
- 2312 (ii) the LS captures similarities in decision-making processes, yielding behav-
2313 iorally meaningful clusters.

2314 To distinguish among these cases, we introduce the *Abstraction Score*. This
2315 metric contextualizes latent clusterings by examining how their patterns align
2316 with reference structures derived from the raw observation and action data. We
2317 consider : $\mathcal{C}_{obs}$, the clustering in the observation space $\mathbb{O}$, and $\mathcal{C}_{act}$, the clustering
2318 in the action-space distribution parameters. The clusterings $\mathcal{C}_{obs}$ and $\mathcal{C}_{act}$ are
2319 performed using the same algorithm, distance metric, and hyperparameters as
2320 those used for $\mathcal{C}_i$ in the LS, ensuring that any biases introduced by the clustering
2321 procedure itself are minimized.

2322 For each latent clustering $\mathcal{C}_i$, we define its similarity to these baselines using
2323 the ARI:

$$\begin{aligned} A_i^{obs} &= \text{ARI}(\mathcal{C}_i, \mathcal{C}_{obs}) \\ A_i^{act} &= \text{ARI}(\mathcal{C}_i, \mathcal{C}_{act}) \end{aligned}$$

2324 The *Abstraction Score* is subsequently calculated as one minus the average of
2325 the maximum absolute similarity values across these two reference baselines:

$$\text{Abstraction Score} = 1 - \frac{1}{n} \sum_{i=1}^n \max(|A_i^{obs}|, |A_i^{act}|) \quad (5.4)$$

2326 The *Abstraction Score* ranges from $[0, 1]$. A score around 0 means that the
2327 latent clusters are strongly aligned with either the observation or the action
2328 space, reflecting minimal abstraction, cases (i). A score close to 1 indicates that
2329 the latent clusters are independent of both the observation and action spaces,
2330 capturing abstractions that meaningfully reflect the agent’s behavior, case (ii).

2331 The following section presents the results of applying these metrics to specific
2332 examples.

2333 5.4 Results

2334 This section quantitatively evaluates this behavior across several Atari bench-
2335 mark environments [Bellemare et al., 2013], which are widely used in RL re-
2336 search [Mnih et al., 2013, Schulman et al., 2017, Zahavy et al., 2017]. The envi-

Environment	<i>Behavior Reconstruction</i>	<i>Abstraction Score</i>	<i>Clustering Universality</i>
Breakout	0.83 ± 0.02	0.99 ± 0.00	0.58 ± 0.09
Mario Bros	0.98 ± 0.02	0.99 ± 0.00	0.55 ± 0.10
Pong	0.95 ± 0.00	0.99 ± 0.00	0.48 ± 0.04
Space Invaders	0.97 ± 0.07	0.99 ± 0.00	0.51 ± 0.08
Surround	0.99 ± 0.00	0.99 ± 0.00	0.58 ± 0.07

Table 5.1: Evaluation of *Behavior Reconstruction*, *Abstraction Score* and *Clustering Universality* across the five Atari environments.

ronments considered in this study are:

- **Breakout:** the agent controls a paddle to bounce a ball and break as many bricks as possible, the ball speeding up after each bounce.
- **Mario Bros:** the agent must defeat enemies, avoid fireballs, and collect mushrooms [AtariAge, 2024].
- **Pong:** a two-player game in which the agent tries to prevent the ball from passing its side while attempting to score against the opponent.
- **Space Invaders:** the agent controls a cannon and must destroy waves of invaders before they reach the ground.
- **Surround:** a two-player grid-based game in which each move blocks the traversed cell, the goal being to trap the opponent.

In all these environments, the agent receives visual observations and operates in discrete action spaces. The initial policies π are trained using Proximal Policy Optimization (PPO) [Schulman et al., 2017] with a deep actor-critic architecture. The network consists of four convolutional layers with configurations (16, 4, 2), (32, 4, 2), (64, 4, 2), (128, 4, 2), where each tuple represents (number of channels, kernel size, stride), followed by three fully connected layers with 512, 256, and 128 units, respectively. To ensure that these initial policies effectively learn to maximize rewards, we follow the recommendations of [Machado et al., 2017] and compare the performance of our initial policies against the benchmark scores reported in the same study to verify that they meet the required performance standards.

For each initial policy, three surrogate policies π' sharing the same architecture are trained as described in Section 5.2. These surrogate policies are trained on approximately one million observation-action pairs each.

We first evaluate the ability of the surrogate policies to replicate the behavior of the initial policies (*Behavior Reconstruction*, Table 5.1). Across all environments, the surrogate policies achieve an average *Behavior Reconstruction* of 0.94 ± 0.02 , indicating that they reliably reproduce the decisions of the initial policies over 100 evaluation episodes.




Figure 5.5: Evolution of *Clustering Universality* with respect to the number of clusters, for each of the five Atari environments.

For the remainder of the analysis, a random subset of 50,000 observations is drawn from the training dataset and embedded into the LS of each surrogate policy, using the representation extracted directly after the last convolutional layer. This projection results in a point cloud $\mathcal{E}$. For clustering, we opted for the simplest approach and applied **K-means clustering** [MacQueen, 1967] to the point cloud, using the Euclidean distance between the projected points to form the clusters.

Across all environments, policies, and numbers of clusters, the average *Abstraction Score* is 0.99 ± 0.00 . This indicates that the clusters formed in the LS are not mere reflections of the raw pixel inputs or of the action distributions.

Figure 5.5 reports the evolution of *Clustering Universality* with respect to the number of clusters. From three clusters onward, all curves stabilize around 0.5 in every environment. Since *Clustering Universality* is computed from the ARI, this result indicates a strong similarity in the clustering of the LS across surrogate policies within each environment [Amezquita et al., 2020], regardless of the number of clusters chosen. This stability indicates that there is no single optimal number of clusters for explaining an agent’s policy. This raises the question of how many clusters should be selected to perform an analysis of LS. Our analysis suggests that this depends on the desired level of granularity.

With fewer clusters, CAR highlights broad strategic patterns. For instance, in the Mario Bros environment, the two main groups correspond to whether the POW block is present (highlighted in red in Figure 5.6A) or not. This is particularly interesting, as the use of the POW block is a key determinant of strategy in this environment [AtariAge, 2024].

With a larger number of clusters, the analysis captures more fine-grained tactical behaviors, such as ‘Attacking the left.’ and ‘Navigating through the level.’ (see Figure 5.6B).

Both perspectives are valid, depending on whether the goal is to study strategic or tactical behaviors. Similar interpretable structures can be observed across all evaluated environments.




Figure 5.6: Effect of the number of clusters on the groupings identified by CAR in the Mario Bros environment.

5.5 Conclusion

The three hypotheses investigated in this manuscript, *Mechanistic Hypothesis* (Hyp. 3), *Concept Hypothesis* (Hyp. 1), and *Representation Convergence Hypothesis* (Hyp. 4), formed the conceptual basis for the development of CAR. The experiments conducted with CAR provide empirical support for these hypotheses, with the high *Abstraction Score* and stable *Clustering Universality* obtained across all environments. These findings suggest that this hypothesis-driven approach provides a promising direction for understanding the structure of learned representations in RL agents.

Beyond its role as an experimental framework, CAR provides a general and automated method for analysing agent representations. Both the analysis procedure and the metrics are fully automated, requiring no manual intervention or environment-specific adjustments. This makes CAR a general method that can be applied to a wide range of RL settings.

However, several implementation choices were made for have not yet been evaluated. These include the NN architectures, the NN layer selected for analysis, the distance used to compare embeddings, and the clustering algorithm. Their impact should therefore be studied jointly to determine how changes in the analysis procedure affect the discovered CS.

A more fundamental limitation concerns the semantic interpretation of the clusters. While CAR provides a quantitative description of the structure present

2418 in the agent’s representations, the semantics assigned to individual clusters
2419 currently rely on manual interpretation. The labels associated with the clusters
2420 in Figure 5.6 represent our interpretation of the discovered structure. This
2421 process is time-consuming and subjective. Developing methods to automatically
2422 extract the semantics of each cluster is therefore an important direction for
2423 future work.

2424 More broadly, when considering CAR within the field of explainability, it is
2425 important to emphasize that substantial work remains to be done. The results
2426 presented in this chapter show that CAR can provide information about the
2427 concepts represented in the agent’s LS, which can therefore be considered an
2428 explanation in the sense defined in Section 4.1. How these explanations can be
2429 used to support validation, diagnosis, or certification remains to be established
2430 ??.

2431 These limitations directly motivate several avenues for future research. Study-
2432 ing the impact of the selected LS layer, distance measures, and clustering al-
2433 gorithms would clarify how much the discovered structure depends on these
2434 methodological choices. Automating cluster interpretation and developing human-
2435 centered evaluation protocols would strengthen semantic validity of the resulting
2436 explanations. Finally, connecting CAR explanations more directly to the agent’s
2437 decisions and failures would help determine how they can support the broader
2438 objectives of explainability.

Acronyms

2439	
2440	AI Artificial Intelligence 51, 55
2441	ARI Adjusted Rand Index 77–80, 82
2442	CAR Clustering Agent Representation 70–74, 76, 77, 79, 82–84
2443	CNN Convolutional Neural Network 2, 15
2444	CS Concept Structure 20, 24, 25, 29–33, 70–72, 83
2445	DDPG Deep Deterministic Policy Gradient 37
2446	DL Deep Learning 1–7, 9, 15, 17, 37, 49, 51, 55, 56
2447	DQN Deep Q-Networks 37
2448	LLM Large Language Model 3, 17, 21, 23, 29, 38, 51
2449	LR Latent Representation 17, 18, 21–33, 63, 70–72, 77
2450	LS Latent Space 15–21, 23–25, 27, 28, 30–32, 62, 63, 65, 66, 69–72, 74–77, 79,
2451	80, 82, 84
2452	LSTM Long Short-Term Memory 2, 15
2453	MDP Markov Decision Processes 36–38, 40–42, 44–50, 60, 65, 66
2454	MLP MultiLayer Perceptron 5–7, 15
2455	MSE Mean Squared Error 4, 32, 48, 50, 74, 75
2456	NN Neural Network 1–10, 13, 15–18, 20, 21, 25–27, 31, 33, 34, 37, 38, 45–47,
2457	55, 58, 59, 62, 66, 68, 70–74, 83
2458	PPO Proximal Policy Optimization 37, 81
2459	RI Rand Index 78
2460	RL Reinforcement Learning 17, 34–47, 49, 51, 52, 56–58, 60, 61, 63, 67, 68, 70,
2461	72, 80, 83
2462	SAE Sparse Autoencoder 26, 28–30, 32
2463	SGD Stochastic Gradient Descent 10, 13, 15, 22, 29, 31
2464	XAI eXplainable Artificial Intelligence 51, 52, 55–57, 68
2465	XRL eXplainable Reinforcement Learning 51, 52, 56–58, 60, 63, 67–69

Bibliography

- [Alaarabiou et al., 2024] Alaarabiou, M., Delestre, N., and Vercouter, L. (2024). Explicabilité en Apprentissage par Renforcement : vers une Taxinomie Unifiée. In 18èmes Journées d’Intelligence Artificielle Fondamentale, La rochelle, France.
- [Amezquita et al., 2020] Amezcquita, R. A., Lun, A. T., Becht, E., Carey, V. J., Carpp, L. N., Geistlinger, L., Marini, F., Rue-Albrecht, K., Risso, D., Sone-son, C., et al. (2020). Orchestrating single-cell analysis with Bioconductor, volume 17. Nature Publishing Group US New York. **Book**.
- [AtariAge, 2024] AtariAge (2024). Mario bros. atari 2600. **Web**.
- [Barredo Arrieta et al., 2020] Barredo Arrieta, A., Díaz-Rodríguez, N., Del Ser, J., Bennetot, A., Tabik, S., Barbado, A., García, S., Gil-López, S., Molina, D., Benjamins, R., Chatila, R., and Herrera, F. (2020). Explainable artificial intelligence (xai): Concepts, taxonomies, opportunities and challenges toward responsible ai. Information Fusion, 58:82–115. **Article**.
- [Barto et al., 1983] Barto, A. G., Sutton, R. S., and Anderson, C. W. (1983). Neuronlike adaptive elements that can solve difficult learning control problems. IEEE Transactions on Systems, Man, and Cybernetics, 13(5):835–846.
- [Bellemare et al., 2013] Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. Journal of Artificial Intelligence Research, 47:253–279. **Article**.
- [Bellman, 1957] Bellman, R. E. (1957). Dynamic programming.
- [Bellucci et al., 2021] Bellucci, M., Delestre, N., Malandain, N., and Zanni-Merk, C. (2021). Towards a terminology for a fully contextualized xai. Procedia Computer Science, 192:241–250. **Article**.
- [Bengio et al., 2014] Bengio, Y., Courville, A., and Vincent, P. (2014). Representation learning: A review and new perspectives.
- [Beyret et al., 2019] Beyret, B., Shafti, A., and Faisal, A. A. (2019). Dot-to-dot: Explainable hierarchical reinforcement learning for robotic manipulation. In 2019 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS). IEEE. **Article**.

- [Bond-Taylor et al., 2021] Bond-Taylor, S., Leach, A., Long, Y., and Willcocks, C. G. (2021). Deep generative modelling: A comparative review of vaes, gans, normalizing flows, energy-based and autoregressive models. **Article**.
- [Bricken et al., 2023] Bricken, T., Templeton, A., Batson, J., Chen, B., Jermyn, A., Conerly, T., et al. (2023). Towards monosemanticity: Decomposing language models with dictionary learning. Transformer Circuits Thread. **Web**.
- [Brown et al., 2020] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., et al. (2020). Language models are few-shot learners. Advances in neural information processing systems, 33:1877–1901.
- [Burda et al., 2018] Burda, Y., Edwards, H., Pathak, D., Storkey, A., Darrell, T., and Efros, A. A. (2018). Large-scale study of curiosity-driven learning.
- [Chen et al., 2021] Chen, M., Tworek, J., Jun, H., Yuan, Q., de Oliveira Pinto, H. P., Kaplan, J., et al. (2021). Evaluating large language models trained on code. **Article**.
- [Coates et al., 2011] Coates, A., Ng, A. Y., and Lee, H. (2011). An analysis of single-layer networks in unsupervised feature learning. In Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics (AISTATS), pages 215–223.
- [Covington et al., 2016] Covington, P., Adams, J., and Sargin, E. (2016). Deep neural networks for youtube recommendations. In Proceedings of the 10th ACM conference on recommender systems, pages 191–198.
- [Doshi-Velez and Kim, 2017] Doshi-Velez, F. and Kim, B. (2017). Towards a rigorous science of interpretable machine learning. **Article**.
- [Elhage et al., 2022] Elhage, N., Hume, T., Olsson, C., Schiefer, N., Henighan, T., Kravec, S., et al. (2022). Toy models of superposition. Transformer Circuits Thread. **Web**.
- [Espeholt et al., 2018] Espeholt, L., Soyer, H., Munos, R., Simonyan, K., Mnih, V., Ward, T., Doron, Y., Firoiu, V., Harley, T., Dunning, I., Legg, S., and Kavukcuoglu, K. (2018). IMPALA: Scalable distributed deep-RL with importance weighted actor-learner architectures. In Proceedings of the 35th International Conference on Machine Learning (ICML).
- [Garreau and von Luxburg, 2020] Garreau, D. and von Luxburg, U. (2020). Explaining the explainer: A first theoretical analysis of lime. In Chiappa, S. and Calandra, R., editors, Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics, volume 108 of Proceedings of Machine Learning Research, pages 1287–1296. PMLR. **Article**.
- [Goodfellow et al., 2016] Goodfellow, I., Bengio, Y., Courville, A., and Bengio, Y. (2016). Deep learning, volume 1. MIT press Cambridge.
- [Greydanus et al., 2018] Greydanus, S., Koul, A., Dodge, J., and Fern, A. (2018). Visualizing and understanding atari agents. **Article**. **Code**.

- [Guidotti et al., 2018] Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Pedreschi, D., and Giannotti, F. (2018). A survey of methods for explaining black box models. **Article**.
- [Gunning and Aha, 2019] Gunning, D. and Aha, D. (2019). Darpa’s explainable ai (xai) program. *AI Magazine*, 40(2):44–58. **Article**.
- [Guo et al.,] Guo, X., Gao, L., Liu, X., and Yin, J. Improved deep embedded clustering with local structure preservation. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence*, pages 1753–1759. International Joint Conferences on Artificial Intelligence Organization.
- [Haas et al., 2024] Haas, R., Huberman-Spiegelglas, I., Mulayoff, R., Graßhof, S., Brandt, S. S., and Michaeli, T. (2024). Discovering interpretable directions in the semantic latent space of diffusion models. In *2024 IEEE 18th International Conference on Automatic Face and Gesture Recognition (FG)*, pages 1–9. IEEE.
- [He et al., 2017] He, X., Liao, L., Zhang, H., Nie, L., Hu, X., and Chua, T.-S. (2017). Neural collaborative filtering.
- [Hessel et al., 2017] Hessel, M., Modayil, J., van Hasselt, H., Schaul, T., Ostrovski, G., Dabney, W., Horgan, D., Piot, B., Azar, M., and Silver, D. (2017). Rainbow: Combining improvements in deep reinforcement learning.
- [Hilbert and López, 2011] Hilbert, M. and López, P. (2011). The world’s technological capacity to store, communicate, and compute information. *science*, 332(6025):60–65.
- [Hochreiter and Schmidhuber, 1997] Hochreiter, S. and Schmidhuber, J. (1997). Long short-term memory. *Neural computation*, 9(8):1735–1780.
- [Hornik et al., 1989] Hornik, K., Stinchcombe, M., and White, H. (1989). Multilayer feedforward networks are universal approximators. *Neural networks*, 2(5):359–366.
- [Howard, 1960] Howard, R. A. (1960). Dynamic programming and markov processes. MIT Press.
- [Huber et al., 2019] Huber, T., Schiller, D., and André, E. (2019). Enhancing explainability of deep reinforcement learning through selective layer-wise relevance propagation. In *KI 2019: Advances in Artificial Intelligence: 42nd German Conference on AI, Kassel, Germany, September 23–26, 2019, Proceedings 42*, pages 188–202. Springer. **Article**.
- [Hubert and Arabie, 1985] Hubert, L. and Arabie, P. (1985). Comparing partitions. *Journal of classification*, 2(1):193–218. **Article**.
- [Itaya et al., 2021] Itaya, H., Hirakawa, T., Yamashita, T., Fujiyoshi, H., and Sugiura, K. (2021). Visual explanation using attention mechanism in actor-critic-based deep reinforcement learning. **Article**.
- [Jermyn et al., 2024] Jermyn, A., Templeton, A., Batson, J., and Bricken, T. (2024). Tanh penalty in dictionary learning. **Web**.

- [Jumper et al., 2021] Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Žídek, A., Potapenko, A., et al. (2021). Highly accurate protein structure prediction with alphafold. *nature*, 596(7873):583–589.
- [Juozapaitis et al., 2019] Juozapaitis, Z., Koul, A., Fern, A., Erwig, M., and Doshi-Velez, F. (2019). Explainable reinforcement learning via reward decomposition. In *IJCAI/ECAI Workshop on explainable artificial intelligence*. **Article**.
- [Karvonen, 2024] Karvonen, A. (2024). Emergent world models and latent variable estimation in chess-playing language models.
- [Kawaguchi et al., 2022] Kawaguchi, K., Bengio, Y., and Kaelbling, L. (2022). *Generalization in Deep Learning*, page 112–148. Cambridge University Press.
- [Knuth and Moore, 1975] Knuth, D. E. and Moore, R. W. (1975). An analysis of alpha-beta pruning. *Artificial intelligence*, 6(4):293–326. **Article**.
- [Krizhevsky et al., 2012] Krizhevsky, A., Sutskever, I., and Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. *Advances in neural information processing systems*, 25.
- [LeCun et al., 1989] LeCun, Y., Boser, B., Denker, J. S., Henderson, D., Howard, R. E., Hubbard, W., and Jackel, L. D. (1989). Backpropagation applied to handwritten zip code recognition. *Neural computation*, 1(4):541–551.
- [Lewis et al., 2021] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., tau Yih, W., Rocktäschel, T., Riedel, S., and Kiela, D. (2021). Retrieval-augmented generation for knowledge-intensive nlp tasks.
- [Li et al., 2024] Li, K., Hopkins, A. K., Bau, D., Viégas, F., Pfister, H., and Wattenberg, M. (2024). Emergent world representations: Exploring a sequence model trained on a synthetic task.
- [Liang et al., 2018] Liang, E., Liaw, R., Nishihara, R., Moritz, P., Fox, R., Goldberg, K., Gonzalez, J. E., Jordan, M. I., and Stoica, I. (2018). RLlib: Abstractions for distributed reinforcement learning. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*.
- [Lillicrap et al., 2015] Lillicrap, T. P., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., and Wierstra, D. (2015). Continuous control with deep reinforcement learning.
- [Lu et al., 2019] Lu, G., Ouyang, W., Xu, D., Zhang, X., Cai, C., and Gao, Z. (2019). Dvc: An end-to-end deep video compression framework.
- [Machado et al., 2017] Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., and Bowling, M. (2017). Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. **Article. Web**.

- [MacQueen, 1967] MacQueen, J. (1967). Multivariate observations. In Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability, volume 1, pages 281–297. **Article**.
- [Madumal et al., 2020] Madumal, P., Miller, T., Sonenberg, L., and Vetere, F. (2020). Explainable reinforcement learning through a causal lens. In Proceedings of the AAAI conference on artificial intelligence, volume 34, pages 2493–2500. **Article**.
- [Marconato et al., 2023] Marconato, E., Passerini, A., and Teso, S. (2023). Interpretability is in the mind of the beholder: A causal framework for human-interpretable representation learning.
- [McCulloch and Pitts, 1943] McCulloch, W. S. and Pitts, W. (1943). A logical calculus of the ideas immanent in nervous activity. The bulletin of mathematical biophysics, 5(4):115–133.
- [Miklautz et al.,] Miklautz, L., Klein, T., Sidak, K., Leiber, C., Lang, T., Shkabrii, A., Tschischek, S., and Plant, C. Breaking the reclustering barrier in centroid-based deep clustering.
- [Miller, 2019] Miller, T. (2019). Explanation in artificial intelligence: Insights from the social sciences. Artificial Intelligence, 267:1–38. **Article**.
- [Minsky and Papert, 1969] Minsky, M. and Papert, S. (1969). An introduction to computational geometry. Cambridge tiass., HIT, 479(480):104.
- [Misak, 2013] Misak, C. (2013). The American Pragmatists. Oxford University Press.
- [Mitchell, 1997] Mitchell, T. M. (1997). Machine Learning. McGraw-Hill, New York.
- [Mnih et al., 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.
- [Mott et al., 2019] Mott, A., Zoran, D., Chrzanowski, M., Wierstra, D., and Jimenez Rezende, D. (2019). Towards interpretable reinforcement learning using attention augmented agents. Advances in neural information processing systems, 32. **Article**.
- [Newell, 1980] Newell, A. (1980). Physical symbol systems. Cognitive Science, 4(2):135–183.
- [Olah et al., 2020] Olah, C., Cammarata, N., Schubert, L., Goh, G., Petrov, M., and Carter, S. (2020). Zoom in: An introduction to circuits. Distill. **Article**.
- [Olson et al., 2021] Olson, M. L., Khanna, R., Neal, L., Li, F., and Wong, W.-K. (2021). Counterfactual state explanations for reinforcement learning agents via generative deep learning. Artificial Intelligence, 295:103455. **Article**. **Code**.

- [OpenAI et al., 2019] OpenAI, :, Berner, C., Brockman, G., Chan, B., Cheung, V., Dębiak, P., Dennison, C., Farhi, D., Fischer, Q., Hashme, S., Hesse, C., Józefowicz, R., Gray, S., Olsson, C., Pachocki, J., Petrov, M., d. O. Pinto, H. P., Raiman, J., Salimans, T., Schlatter, J., Schneider, J., Sidor, S., Sutskever, I., Tang, J., Wolski, F., and Zhang, S. (2019). Dota 2 with large scale deep reinforcement learning.
- [Ouyang et al., 2022] Ouyang, L., Wu, J., Jiang, X., Almeida, D., Wainwright, C. L., Mishkin, P., Zhang, C., Agarwal, S., Slama, K., Ray, A., Schulman, J., Hilton, J., Kelton, F., Miller, L., Simens, M., Askell, A., Welinder, P., Christiano, P., Leike, J., and Lowe, R. (2022). Training language models to follow instructions with human feedback.
- [Payani and Fekri, 2019] Payani, A. and Fekri, F. (2019). Inductive logic programming via differentiable deep neural logic networks. **Article. Code.**
- [Peirce, 1878] Peirce, C. S. (1878). How to make our ideas clear. Popular Science Monthly, 12:286–302.
- [Peyré and Cuturi, 2020] Peyré, G. and Cuturi, M. (2020). Computational optimal transport. **Article.**
- [Radford et al., 2021] Radford, A., Kim, J. W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., Sastry, G., Askell, A., Mishkin, P., Clark, J., Krueger, G., and Sutskever, I. (2021). Learning transferable visual models from natural language supervision.
- [Raina et al., 2009] Raina, R., Madhavan, A., Ng, A. Y., et al. (2009). Large-scale deep unsupervised learning using graphics processors. In Icml, volume 9, pages 873–880.
- [Rand, 1971] Rand, W. M. (1971). Objective criteria for the evaluation of clustering methods. Journal of the American Statistical association, 66(336):846–850.
- [Ren et al.,] Ren, Y., Pu, J., Yang, Z., Xu, J., Li, G., Pu, X., Yu, P. S., and He, L. Deep clustering: A comprehensive survey.
- [Rosenblatt, 1958] Rosenblatt, F. (1958). The perceptron: a probabilistic model for information storage and organization in the brain. Psychological review, 65(6):386.
- [Ross and Bagnell, 2010] Ross, S. and Bagnell, D. (2010). Efficient reductions for imitation learning. In Teh, Y. W. and Titterton, M., editors, Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics, volume 9 of Proceedings of Machine Learning Research, pages 661–668, Chia Laguna Resort, Sardinia, Italy. PMLR. **Article.**
- [Ross et al., 2011] Ross, S., Gordon, G., and Bagnell, D. (2011). A reduction of imitation learning and structured prediction to no-regret online learning. In Gordon, G., Dunson, D., and Dudík, M., editors, Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, volume 15 of Proceedings of Machine Learning Research, pages 627–635, Fort Lauderdale, FL, USA. PMLR. **Article.**

- [Rudin, 2019] Rudin, C. (2019). Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. Nature Machine Intelligence, 1(5):206–215. **Article**.
- [Rumelhart et al., 1986] Rumelhart, D. E., Hinton, G. E., and Williams, R. J. (1986). Learning representations by back-propagating errors. nature, 323(6088):533–536.
- [Russell and Norvig, 2021] Russell, S. J. and Norvig, P. (2021). Artificial Intelligence: A Modern Approach. Pearson, 4th edition.
- [Samvelyan et al., 2019] Samvelyan, M., Rashid, T., de Witt, C. S., Farquhar, G., Nardelli, N., Rudner, T. G. J., Hung, C.-M., Torr, P. H. S., Foerster, J., and Whiteson, S. (2019). The starcraft multi-agent challenge.
- [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms.
- [Sequeira and Gervasio, 2020] Sequeira, P. and Gervasio, M. (2020). Interestingness elements for explainable reinforcement learning: Understanding agents’ capabilities and limitations. Artificial Intelligence, 288:103367. **Article**.
- [Seshia et al., 2022] Seshia, S. A., Sadigh, D., and Sastry, S. S. (2022). Toward verified artificial intelligence. Communications of the ACM, 65(7):46–55. **Article**.
- [Shao et al., 2024] Shao, Z., Wang, P., Zhu, Q., Xu, R., Song, J., Bi, X., Zhang, H., Zhang, M., Li, Y. K., Wu, Y., and Guo, D. (2024). Deepseekmath: Pushing the limits of mathematical reasoning in open language models.
- [Shlens, 2014] Shlens, J. (2014). Notes on kullback-leibler divergence and likelihood. **Article**.
- [Siegel and Craver, 2024] Siegel, G. and Craver, C. F. (2024). Phenomenological laws and mechanistic explanations. Philosophy of Science, 91(1):132–150.
- [Sieusahai and Guzdial, 2021] Sieusahai, A. and Guzdial, M. (2021). Explaining deep reinforcement learning agents in the atari domain through a surrogate model. **Article**.
- [Silver et al., 2017a] Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., et al. (2017a). Mastering chess and shogi by self-play with a general reinforcement learning algorithm. arXiv preprint arXiv:1712.01815.
- [Silver et al., 2017b] Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., Lillicrap, T., Simonyan, K., and Hassabis, D. (2017b). Mastering chess and shogi by self-play with a general reinforcement learning algorithm.
- [Silver et al., 2021] Silver, D., Singh, S., Precup, D., and Sutton, R. S. (2021). Reward is enough. Artificial intelligence, 299:103535.
- [Simon, 1996] Simon, H. A. (1996). The Sciences of the Artificial. MIT Press, 3 edition.

- [Simonyan et al., 2014] Simonyan, K., Vedaldi, A., and Zisserman, A. (2014). Deep inside convolutional networks: Visualising image classification models and saliency maps. **Article. Code.**
- [Skinner, 1938] Skinner, B. F. (1938). The behavior of organisms: An experimental analysis. Appleton-Century.
- [Sutton, 1988] Sutton, R. S. (1988). Learning to predict by the methods of temporal differences. Machine Learning, 3:9–44.
- [Sutton and Barto, 1998] Sutton, R. S. and Barto, A. G. (1998). Reinforcement Learning: An Introduction. MIT Press, 1st edition.
- [Sutton and Barto, 2018] Sutton, R. S. and Barto, A. G. (2018). Reinforcement Learning: An Introduction. MIT Press. **Book.**
- [Telgarsky, 2016] Telgarsky, M. (2016). Benefits of depth in neural networks. In Conference on learning theory, pages 1517–1539. PMLR.
- [Templeton et al., 2024] Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., et al. (2024). Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet. Transformer Circuits Thread. **Web.**
- [Thorndike, 1911] Thorndike, E. L. (1911). Animal intelligence: Experimental studies. The Psychological Review Monograph Supplement, 8(4).
- [Topin et al., 2021] Topin, N., Milani, S., Fang, F., and Veloso, M. (2021). Iterative bounding mdps: Learning interpretable policies via non-interpretable methods. In Proceedings of the AAAI Conference on Artificial Intelligence, volume 35, pages 9923–9931. **Article.**
- [Trivedi et al., 2022] Trivedi, D., Zhang, J., Sun, S.-H., and Lim, J. J. (2022). Learning to synthesize programs as interpretable and generalizable policies. **Article. Code.**
- [Tsitsiklis and Van Roy, 1996] Tsitsiklis, J. and Van Roy, B. (1996). Analysis of temporal-difference learning with function approximation. Advances in neural information processing systems, 9.
- [Turing, 1948] Turing, A. (1948). Intelligent machinery. National Physical Laboratory Report.
- [Vaswani et al., 2017] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2017). Attention is all you need. Advances in neural information processing systems, 30.
- [Vaswani et al., 2023] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2023). Attention is all you need. **Article.**
- [Voynov and Babenko, 2020] Voynov, A. and Babenko, A. (2020). Unsupervised discovery of interpretable directions in the gan latent space.

- 2781 [Watkins, 1989] Watkins, C. J. C. H. (1989). Learning from delayed rewards.
2782 PhD Thesis, University of Cambridge.
- 2783 [Watkins and Dayan, 1992] Watkins, C. J. C. H. and Dayan, P. (1992). Q-
2784 learning. Machine Learning, 8:279–292.
- 2785 [Wei et al.,] Wei, X., Zhang, Z., Huang, H., and Zhou, Y. An overview on deep
2786 clustering. 590:127761.
- 2787 [Weitkamp et al., 2019] Weitkamp, L., van der Pol, E., and Akata, Z. (2019).
2788 Visual rationalizations in deep reinforcement learning for atari games. **Article.**
- 2789 [Xie et al.,] Xie, J., Girshick, R., and Farhadi, A. Unsupervised deep embedding
2790 for clustering analysis.
- 2791 [Zahavy et al., 2017] Zahavy, T., Zrihem, N. B., and Mannor, S. (2017). Graying
2792 the black box: Understanding dqns. **Article.**
- 2793 [Zednik, 2019] Zednik, C. (2019). Solving the black box problem: A nor-
2794 mative framework for explainable artificial intelligence. arXiv preprint
2795 arXiv:1903.04361.