

1 Explainability in Reinforcement Learning

2 Maxime Alaarabiou

3 Friday 14th August, 2026

Chapter 1

Clustering Agent Representations

As discussed in the previous chapter, current explainability methods for Reinforcement Learning (RL) still lack intrinsic approaches capable of explaining the internal mechanisms of deep neural policies. To the best of our knowledge, only [Zahavy et al., 2017] have proposed exploiting the Latent Space (LS) learned by neural policies in order to identify the concepts used by the agent. Its main idea is to analyze the Latent Representations (LRs) produced by the Neural Network (NN) and organize them into clusters that correspond to high-level concepts.

While the core idea of this paper is closely aligned with our vision of explainability in RL, this approach suffers from two important limitations. First, the clusters are manually constructed. Such a procedure heavily depends on the user’s expertise and neither guarantees reproducibility or scales. Second, no evaluation protocol is proposed to assess the quality of the obtained clusters. Their relevance is essentially determined through visual inspection.

The Clustering Agent Representation (CAR) method proposed in this chapter addresses these two limitations. It automates concept discovery by relying on unsupervised clustering algorithms to identify the Concept Structure (CS) of the LS. It introduces an objective evaluation metric based on the stability of the extracted representations across independently trained policy.

This chapter is organized as follows:

- Section 1.1 lays out the theoretical assumptions underlying the proposed CAR method, which we term the *Representation Convergence Hypothesis* (Hyp. 1).
- Section 1.2 then details the CAR method itself, covering both the extraction of LRs using surrogate policies and the automatic detection of CSs through clustering.
- Section 1.3 introduces the metrics we propose for evaluation.
- Section 1.4 reports the results obtained by applying CAR and these metrics across several Atari environments.
- Section 1.5 discusses these results and situates CAR within the broader landscape of mechanistic interpretability.

- Section 1.6 concludes the chapter and outlines directions for future work.

1.1 Representation Convergence Hypothesis

The CAR method relies on a assumption, referred to as the *Representation Convergence Hypothesis* (Hyp. 1). This hypothesis follows the two assumptions established previously throughout this manuscript.

The *Mechanistic Hypothesis* (Hyp. ??) rejects the notion that deep NNs constitute irreducible black boxes. Adopting a mechanistic perspective, we consider that every process results from a sequence of causal computations. Although these computations may be highly complex, they remain explainable through the analysis of the internal mechanisms.

The *Concept Hypothesis* (Hyp. ??) states that NNs achieve generalization by learning concepts. During optimization, the network progressively organizes its LS into representations that capture the information required to solve the task.

Following these hypotheses, it becomes possible to explain the NN’s process by understanding the organization of its LS. This family of approaches is commonly referred to as mechanistic interpretability.

In mechanistic interpretability, universality refers to the consistency with which a concept is recovered across different models or independent extraction processes [Templeton et al., 2024]. We develop the *Representation Convergence Hypothesis* (Hyp. 1) provides a theoretical explanation for this phenomenon. Under this assumption, the universality of a concept is not merely an empirical observation but the expected consequence of a common underlying CS shared by independently learned representations.

Hypothesis 1: Representation Convergence Hypothesis

For a given task and a given NN architecture, there exists a privileged CS, which emerges because it constitutes an efficient representation required to solve the task.

It is important to clarify that this hypothesis does not imply convergence at the parameter level. Independently trained models are expected to converge toward different sets of weights. The proposed hypothesis concerns a abstract level of representation. We assume that the conceptual organization induced by their LR converges toward an equivalent CS, even though the numerical implementation of the networks differ.

The remainder of this chapter builds upon this hypothesis. This principle forms the theoretical foundation of the *Clustering Universality* metric, introduced in Subsection 1.3.3.

1.2 Method

Following the work of [Zahavy et al., 2017], we analyze the LS through clustering techniques, placing our approach within the broader field of Deep Clustering. This design choice is motivated by our rejection of the *Superposition Hypothesis* (Hyp. ??) and, more generally, of any view that assumes a linearly organized CS, as discussed in Section ??.

The CAR methodology is composed of two main stages. The first stage trains multiple independent NNs, called surrogate policies, which implement the same decision policy. The second stage clusters their LRs in order to identify the underlying conceptual organization.

This section is organized as follows. In the first subsection, we justify the use of multiple surrogate policies in order to study the stability of representations. In the second subsection, we detail how these surrogate policies are trained so as to best match the original policy.

1.2.1 Multiple surrogate policies

The CAR method was developed with the aim of verifying the *Representation Convergence Hypothesis* (Hyp. 1). To test this hypothesis, several NNs must be trained independently and their LRs compared. If similar CSs consistently emerge, this provides empirical evidence supporting the Representation Convergence Hypothesis.

In supervised learning, this protocol is straightforward since several models can be independently trained on the same target function. RL introduces an additional difficulty. Because of stochastic exploration, independently trained agents are not guaranteed to converge toward the same policy. Even if they achieve comparable performance, they may solve the task using different strategies. Making it impossible to determine whether differences observed in the LS originate from the learning dynamics or from genuine differences between the learned policies.

To overcome this limitation, we consider an already trained initial policy. This policy is then imitated using imitation learning [Ross et al., 2011]. To this end, we independently train several surrogate NNs in a supervised manner. Since all surrogate models approximate the same target function, they converge toward equivalent decision policies. This makes it possible to investigate whether equivalent policies also exhibit equivalent conceptual organizations.


It is important to emphasize that the use of surrogate policies does not make CAR a surrogate explainability method, as defined in Subsection ?? . The surrogate policies remain deep NNs. Their role is not to simplify the decision process but to provide multiple independent realizations of the same policy function, enabling the study of representation convergence.

1.2.2 Training surrogate policies

Throughout the remainder of this manuscript, the initial policy we aim to explain is denoted as π . This initial policy is obtained without imposing any constraints on its origin. It may result from classical or deep reinforcement learning [Sutton and Barto, 2018], alpha-beta algorithm [Knuth and Moore, 1975], or even from demonstrations provided by humans [Ross and Bagnell, 2010]. The crucial requirement is the availability of a dataset containing a sufficient number of projections from the observation space $\mathcal{O}$ to the action space $\mathcal{A}$ induced by π .

The replicated policies are referred to as surrogate policies, denoted by π' . For any observation O , the distribution over the action space produced by π' should approximate that of π as closely as possible.

The surrogate policy π' is implemented as a NN parameterized by Θ . This network is composed of two components: an encoder and a decoder, respectively

Figure 1.1: Example of a surrogate policy architecture π' .

122 highlighted by the blue and black dashed rectangles in Figure 1.1. The encoder
 123 f_{enc} maps the observation O to the embedding. The decoder, which reconstructs
 124 the action distribution parameters from the embedding.

125 To train π' to accurately replicate the action distribution, we define a global
 126 loss function $\mathcal{L}$ optimized via gradient descent as introduced in Section ??.
 127 This loss, defined in Equation (1.1), is computed as the Mean Squared Error
 128 (MSE) between the parameters of the action distributions, obtained by applying
 129 the function p , which maps a distribution to its parameterization. In practice,
 130 this function does not need to be explicitly defined, as neural policies directly
 131 output the parameters of the action distribution. Once training is complete, this
 132 loss ensures that π' closely approximates the behavior of the original policy π .
 133 Figure 1.2 summarizes this overall procedure.

$$\min_{\Theta} \sum_{O \in \mathcal{O}} \mathcal{L}(O) = \min_{\Theta} \sum_{O \in \mathcal{O}} \text{MSE}(p(\pi(O)), p(\pi'_{\Theta}(O))) \quad (1.1)$$

134 We denote by $\mathcal{E}$ the set of embeddings obtained by applying f_{enc} to a subset
 135 of observations $\mathbb{O} \subseteq \mathcal{O}$.

136 On the set $\mathcal{E}$, we use a clustering function f_{clust} to partition the point cloud
 137 of embedded observations into distinct clusters. Each observation is assigned
 138 a cluster label based on its proximity in the LS, resulting in the clustering $\mathcal{C}$.
 139 Algorithm 1 summarizes the complete CAR procedure.

140 To demonstrate this clustering faithfully reflects the agent's internal decision-
 141 making process, we introduce a set of evaluation metrics, which will be detailed
 142 in the following section.




Figure 1.2: Illustration of the Clustering Agent Representation method.

1.3 Evaluation Metrics

A key contribution of this work is the development of three evaluation metrics, designed to assess the effectiveness of CAR:

- **Behavior Reconstruction** : it measures how well a surrogate policy faithfully reproduces the behavior of the original policy by comparing their cumulative rewards.
- **Clustering Universality** : it evaluates the stability of clusterings by measuring the similarity of clusters obtained from independently trained surrogate policies.
- **Abstraction Score** : it quantifies the extent to which LS clusters capture abstractions rather than directly reflecting the structure of observations or actions.

The remainder of this section details the motivation for each metric and describes how they are calculated. In the first subsection, we develop the *Behavior Reconstruction*. We then introduce the Adjusted Rand Index (ARI) measure, which serves as the foundation for the two other metrics introduced in the subsequent subsections, namely *Clustering Universality* and *Abstraction Score*.

The following notations will be adopted for the rest of the manuscript: we consider n surrogate policies, denoted as $\pi'_1, \dots, \pi'_n$, trained as described in Section 1.2. Each surrogate policy is associated with a projection function $f_{enc_1}, \dots, f_{enc_n}$, which is applied to the set of observations $\mathcal{O}$, yielding embeddings $\mathcal{E}_1, \dots, \mathcal{E}_n$. A clustering function is then applied to each embedding set, producing partitions $\mathcal{C}_1, \dots, \mathcal{C}_n$.

Algorithm 1: Clustering Agent Representation (CAR)**Input:**

Original policy π
 Observation dataset $\mathbb{O}$
 Number of surrogate policies n
 Surrogate policy architecture π' (encoder f_{enc} , decoder)
 Learning rate η
 Clustering function f_{clust}

Output:

Trained surrogate policies $\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}$
 Clusterings $\mathcal{C}_1, \dots, \mathcal{C}_n$

Behavioral cloning dataset built from the target policy π

1 $\mathcal{D} \leftarrow \{(O, p(\pi(O))) \mid O \in \mathbb{O}\};$

MSE loss between action-distribution parameters, as defined in Equation (??)

2 $\mathcal{L} \leftarrow \text{MSE};$

Independently train n surrogate policies by behavioral cloning

3 **for** $i = 1, \dots, n$ **do**

4 Initialize Θ_i randomly;

Behavioral cloning of π via SGD, see Algorithm ??

5 $\Theta_i \leftarrow \text{SGD}(\mathcal{D}, \pi', \mathcal{L}, \eta, \Theta_i);$

Cluster the latent representations of each surrogate policy

6 **for** $i = 1, \dots, n$ **do**

Embeddings produced by the encoder of π'_{Θ_i}

7 $\mathcal{E}_i \leftarrow \{f_{enc_i}(O) \mid O \in \mathbb{O}\};$

Partition of the embedded observations into clusters

8 $\mathcal{C}_i \leftarrow f_{clust}(\mathcal{E}_i);$

9 **Return** $(\pi'_{\Theta_1}, \dots, \pi'_{\Theta_n}), (\mathcal{C}_1, \dots, \mathcal{C}_n)$

1.3.1 Behavior Reconstruction

We argue that explaining a function b that closely approximates another function a is essentially the same as explaining a . If b accurately captures the key decision patterns of a , understanding b provides meaningful insights into how a works [Guidotti et al., 2018]. In our setting, this means that generating explanations for π' that faithfully reproduce π amounts to explaining π directly. To verify that π' is indeed a faithful copy, we introduce the *Behavior Reconstruction* metric.

This metric is computed by comparing the cumulative rewards obtained by the original policy π and the surrogate policies π' over m evaluation episodes. The behavior reconstruction metric is then defined as:

$$\text{Behavior Reconstruction} = \frac{1}{n \sum_{i=1}^m g_{\pi}^i} \sum_{j=1}^n \sum_{k=1}^m g_{\pi'_j}^k \quad (1.2)$$

In imitation learning, it is unlikely to expect the surrogate policy π' to

outperform the initial policy π it aims to approximate through supervised learning [Ross and Bagnell, 2010]. Consequently this metric ranges in $[0, 1]$. Values approaching 1 indicate minimal loss in reproducing π , while values close to 0 indicate poor reproduction.

We deliberately adopt a reward-based definition rather than alternatives based on actions or action distributions (e.g., training loss, accuracy, KL divergence [Shlens, 2014], or optimal transport [Peyré and Cuturi, 2020]). Since surrogate policies are trained in a supervised manner, they cannot develop strategies beyond those demonstrated by the initial policy. This means that if the rewards achieved are comparable, the surrogate policy is effectively copying the original policy’s behavior.

1.3.2 Adjusted Rand Index

Within CAR, concepts are assumed to correspond to clusters in the LS. Combined with the *Representation Convergence Hypothesis* (Hyp. 1), this leads to a direct experimental prediction: if independently trained surrogate policies converge toward the same conceptual organization, then clustering their LRs should produce similar partitions.

When a particular observation is processed by the network, it is expected to activate the same concept regardless of the specific neural implementation of the policy. If independently trained policies converge toward the same conceptual organization, a given observation should activate the same concept across all models. As a consequence, observations grouped together within a cluster for one policy should also be grouped together for another independently trained policy.

In the clustering literature, the most used measures for comparing two partitions is the ARI [Hubert and Arabie, 1985]. This measure quantifies the agreement between cluster assignments independently of the spatial organization of the clusters. Both metrics introduced later in this chapter, namely *Clustering Universality* and *Abstraction Score*, rely on the ARI. Since the ARI is a complex measure, we review its definition and discuss its main properties.

The ARI is derived from the Rand Index (RI) [Rand, 1971], which measures the agreement between two partitions by considering every possible pair of point cloud. For each pair, four situations are possible: the point may belong to the same cluster in both partitions, to different clusters in both partitions, or they may disagree by being grouped together in only one of the two partitions. The RI measures the proportion of pairs for which the two partitions agree:

$$\text{RI} = \frac{\text{number of agreeing pairs}}{\text{total number of pairs}}.$$

Although the RI is intuitive, it suffers from an important limitation. In most clustering problems, the majority of pairs of points belong to different clusters. Consequently, even two completely independent partitions will agree on a large number of pairs, simply because both assign them to different clusters. This effect becomes more pronounced as the number of points or clusters increases.

The ARI addresses this issue by accounting for the agreement that would be expected purely by chance. It compares the observed agreement RI to the expected agreement under random partitions $\mathbb{E}[\text{RI}]$. The score is then normalized by the maximum agreement achievable beyond this random baseline,

223 $\max(\text{RI}) - \mathbb{E}[\text{RI}]$. As a result, an ARI of 0 corresponds to the level of agreement
 224 expected from random clusterings, while an ARI of 1 indicates perfect agreement
 225 between the two partitions:

$$\text{ARI} = \frac{\text{RI} - \mathbb{E}[\text{RI}]}{\max(\text{RI}) - \mathbb{E}[\text{RI}]}$$

226 In practice, the ARI is computed directly from the cluster assignments of the
 227 two partitions. Let clustering A be composed of partitions $\{A_i\}$ and clustering B
 228 of partitions $\{B_j\}$. Let $a_i = |A_i|$ denote the number of samples in partition i of
 229 clustering A , and $b_j = |B_j|$ the number of samples in partition j of clustering B .
 230 Furthermore, let $n_{ij} = |A_i \cap B_j|$ be the number of samples shared by partition i
 231 of clustering A and partition j of clustering B , and let n be the total number of
 232 samples. With these definitions, the ARI can be computed as:

$$\text{ARI} = \frac{\sum_{i,j} \binom{n_{ij}}{2} - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}{\frac{1}{2} \left[\sum_i \binom{a_i}{2} + \sum_j \binom{b_j}{2} \right] - \left[\sum_i \binom{a_i}{2} \sum_j \binom{b_j}{2} \right] / \binom{n}{2}}$$

233 1.3.3 Clustering Universality

234 Having introduced the ARI, we can now define the *Clustering Universality* metric.
 235 While the ARI quantifies the similarity between two conceptual partitions, the
 236 objective of CAR is to assess the consistency of conceptual organizations across
 237 an entire set of independently trained surrogate policies.

238 To this end, the same set of observations is projected into the LS of each sur-
 239rogate policy and clustered. Under the Representation Convergence Hypothesis,
 240 these independently obtained partitions are expected to be similar. *Clustering*
 241 *Universality* quantifies this by computing the mean pairwise ARI:

$$\text{Clustering Universality} = \frac{1}{\binom{n}{2}} \sum_{i=1}^n \sum_{j=0}^{i-1} \text{ARI}(\mathcal{C}_i, \mathcal{C}_j). \quad (1.3)$$

242 This metric takes values in $[-1, 1]$. A value close to 1 indicates that surrogate
 243 policies produce highly similar partitions, suggesting that the clusters reflect
 244 intrinsic properties of the policy. Conversely, a value near 0 implies that the
 245 similarity between partitions is no better than that of two randomly generated
 246 clusterings, meaning that the clustering fails to capture any meaningful structure
 247 in the policy. A negative value, on the other hand, signals that the agreement
 248 between partitions is worse than what would be expected by chance, indicating
 249 a systematic disagreement between the clusterings.

250 Figure 1.3 illustrates *Clustering Universality* in the ideal case. It shows three
 251 LS obtained by projecting the same seven observations, labeled 1 to 7, through
 252 three independently trained surrogate policies. Although the embeddings occupy
 253 different positions, the clusters group exactly the same elements in all three
 254 cases: the partitioning is therefore identical.

255 1.3.4 Abstraction Score

256 If *Clustering Universality* is high, it becomes important to determine what
 257 underlies this stability. Two possibilities arise:




Figure 1.3: Illustration of *Clustering Universality* in the ideal case where projections of the same observations are grouped within the same clusters across all surrogate policies.

- (i) the LS merely reproduces the structure of the observation space or action distribution, offering no abstraction, as defined in Definition ??;
- (ii) the LS captures similarities in decision-making processes, yielding behaviorally meaningful clusters.

To distinguish among these cases, we introduce the *Abstraction Score*. This metric contextualizes latent clusterings by examining how their patterns align with reference structures derived from the raw observation and action data. We consider : $\mathcal{C}_{obs}$, the clustering in the observation space $\mathbb{O}$, and $\mathcal{C}_{act}$, the clustering in the action-space distribution parameters. The clusterings $\mathcal{C}_{obs}$ and $\mathcal{C}_{act}$ are performed using the same algorithm, distance metric, and hyperparameters as those used for $\mathcal{C}_i$ in the LS, ensuring that any biases introduced by the clustering procedure itself are minimized.

For each latent clustering $\mathcal{C}_i$, we define its similarity to these baselines using the ARI:

$$A_i^{obs} = \text{ARI}(\mathcal{C}_i, \mathcal{C}_{obs})$$

$$A_i^{act} = \text{ARI}(\mathcal{C}_i, \mathcal{C}_{act})$$

The *Abstraction Score* is subsequently calculated as one minus the average of the maximum absolute similarity values across these two reference baselines:

$$\text{Abstraction Score} = 1 - \frac{1}{n} \sum_{i=1}^n \max(|A_i^{obs}|, |A_i^{act}|) \quad (1.4)$$

The *Abstraction Score* ranges from $[0, 1]$. A score around 0 means that the latent clusters are strongly aligned with either the observation or the action space, reflecting minimal abstraction, cases (i). A score close to 1 indicates that the latent clusters are independent of both the observation and action spaces, capturing abstractions that meaningfully reflect the agent’s behavior, case (ii).

The following section presents the results of applying these metrics to specific examples.




Figure 1.4: Two-dimensional projection of the LS embeddings of a surrogate policy trained on Space Invaders, together with the CAR clusters and their corresponding observations.

1.4 Results

Figure 1.4 previews the kind of explanation produced by CAR: a two-dimensional projection of the LS of a surrogate policy trained on Space Invaders, together with the resulting clusters and their corresponding observations. Each cluster corresponds to a distinct, human-interpretable stage of the game, from “few or no invaders have been neutralized” to “invaders descend, but their numbers are reduced.” The remainder of this section quantitatively evaluates this behavior across several Atari benchmark environments [Bellemare et al., 2013], which are widely used in RL research [Mnih et al., 2013, Schulman et al., 2017, Zahavy et al., 2017]. The environments considered in this study are:

- **Breakout:** the agent controls a paddle to bounce a ball and break as many bricks as possible, the ball speeding up after each bounce.
- **Mario Bros:** the agent must defeat enemies, avoid fireballs, and collect mushrooms [AtariAge, 2024].
- **Pong:** a two-player game in which the agent tries to prevent the ball from passing its side while attempting to score against the opponent.
- **Space Invaders:** the agent controls a cannon and must destroy waves of invaders before they reach the ground.

Table 1.1: Evaluation of *Behavior Reconstruction*, *Abstraction Score* and *Clustering Universality* across the five Atari environments.

Environment	<i>Behavior Reconstruction</i>	<i>Abstraction Score</i>	<i>Clustering Universality</i>
Breakout	0.83 ± 0.02	0.99 ± 0.00	0.58 ± 0.09
Mario Bros	0.98 ± 0.02	0.99 ± 0.00	0.55 ± 0.10
Pong	0.95 ± 0.00	0.99 ± 0.00	0.48 ± 0.04
Space Invaders	0.97 ± 0.07	0.99 ± 0.00	0.51 ± 0.08
Surround	0.99 ± 0.00	0.99 ± 0.00	0.58 ± 0.07

- **Surround:** a two-player grid-based game in which each move blocks the traversed cell, the goal being to trap the opponent.

In all these environments, the agent receives visual observations and operates in discrete action spaces. The initial policies π are trained using Proximal Policy Optimization (PPO) [Schulman et al., 2017] with a deep actor-critic architecture. The network consists of four convolutional layers with configurations (16, 4, 2), (32, 4, 2), (64, 4, 2), (128, 4, 2), where each tuple represents (number of channels, kernel size, stride), followed by three fully connected layers with 512, 256, and 128 units, respectively. To ensure that these initial policies effectively learn to maximize rewards, we follow the recommendations of [Machado et al., 2017] and compare the performance of our initial policies against the benchmark scores reported in the same study to verify that they meet the required performance standards.

For each initial policy, three surrogate policies π' sharing the same architecture are trained as described in Section 1.2. These surrogate policies are randomly initialized and trained on approximately one million observation-action pairs each.

We first evaluate the ability of the surrogate policies to replicate the behavior of the initial policies (*Behavior Reconstruction*, Table 1.1). Across all environments, the surrogate policies achieve an average *Behavior Reconstruction* of 0.94 ± 0.02 , indicating that they reliably reproduce the decisions of the initial policies over 100 evaluation episodes.

For the remainder of the analysis, a random subset of 50,000 observations is drawn from the training dataset and embedded into the LS of each surrogate policy, using the representation extracted directly after the last convolutional layer. This projection results in a point cloud $\mathcal{E}$. For clustering, we opted for the simplest approach and applied **K-means clustering** [MacQueen, 1967] to the point cloud, using the Euclidean distance between the projected points to form the clusters.

Across all environments, policies, and numbers of clusters, the average *Abstraction Score* is 0.99 ± 0.00 . This indicates that the clusters formed in the LS are not mere reflections of the raw pixel inputs or of the action distributions, and therefore capture a genuine abstraction of the agent’s behavior, as defined in Definition ??.




Figure 1.5: Evolution of *Clustering Universality* with respect to the number of clusters, for each of the five Atari environments.

Figure 1.5 reports the evolution of *Clustering Universality* with respect to the number of clusters. From three clusters onward, all curves stabilize around 0.5 in every environment. Since *Clustering Universality* is computed from the ARI, this result indicates a strong similarity in the clustering of the LS across surrogate policies within each environment [Amezquita et al., 2020], regardless of the number of clusters chosen.

The combination of a high *Abstraction Score* and a strong *Clustering Universality* indicates that the latent clusters are both highly similar across independently trained surrogate policies and distinct from those obtained in the raw observation or action space. These results support the *Representation Convergence Hypothesis* (Hyp. 1) and confirm that the LS captures high-level concepts, making it a reliable representation for analyzing agent behavior.

1.5 Discussion

Across the five Atari environments, our results demonstrate strong stability, as illustrated in Figure 1.5. This stability indicates that there is no single optimal number of clusters to explain an agent’s policy. Instead, the level of abstraction depends on the desired granularity. With fewer clusters, CAR highlights broad strategic patterns. For instance, in the Mario Bros environment, the two main groups correspond to whether the POW block is present (highlighted in red in Figure 1.6A) or not. This is particularly interesting, as the use of the POW block is a key determinant of strategy in this environment [AtariAge, 2024]. Conversely, with a larger number of clusters, the analysis captures more fine-grained tactical behaviors, such as “Attacking the left.” versus “Navigating through the level.” (see Figure 1.6B). Both perspectives are valid, depending on whether one seeks strategic or tactical insights. Similar interpretable structures can be observed in all evaluated environments.

A major strength of CAR lies in its general applicability. We validated it across multiple Atari environments, and it remained effective regardless of the




Figure 1.6: Effect of the number of clusters on the groupings identified by CAR in the Mario Bros environment.

environment considered. This suggests that clustering the LS is a robust tool for extracting meaningful information about agent behavior.

As emphasized throughout this manuscript, this work aligns with the perspective of mechanistic interpretability, in which NNs are not treated as black boxes but as structured systems whose LS can be studied and understood [Bereska and Gavves, 2024]. Research teams at Anthropic [Elhage et al., 2022, Bricken et al., 2023, Templeton et al., 2024] were among the first to propose and exploit this mechanistic approach to explainability through the use of Sparse Autoencoders (SAEs). However, they rely on the *Superposition Hypothesis* (Hyp. ??), unlike CAR, which demonstrates that meaningful structures can be uncovered through clustering without assuming a linearly organized CS.

1.6 Conclusion

In this chapter, we introduced **Clustering Agent Representation** (CAR), a method that extracts behaviorally meaningful clusters from an agent’s LS to enhance explainability. To enable a rigorous evaluation, we proposed three metrics, *Behavior Reconstruction*, *Clustering Universality*, and *Abstraction Score*, which respectively confirm that CAR achieves faithful policy reproduction, strong abstraction capabilities, and robust universality across Atari environments. This demonstrates the relevance and effectiveness of our method in uncovering stable structures underlying agent behavior, and provides empirical support for the *Representation Convergence Hypothesis* (Hyp. 1). By providing a scalable, data-

382 driven, and quantitatively grounded framework, CAR contributes a significant
383 step toward explainability in RL.

384 Future work will aim to extend CAR to continuous-control and real-world
385 domains to evaluate its generality beyond visual Atari tasks. Exploring the
386 impact of the selected LS layer on the learned representations could provide
387 further insight into how the model organizes its internal knowledge. Adopting
388 alternative distance measures and clustering paradigms, such as density-based
389 or hierarchical approaches, may offer a more comprehensive understanding of
390 how the LS is organized. In parallel, advancing automated cluster interpretation
391 and developing human-centered evaluation protocols would strengthen both the
392 explainability and the practical relevance of the approach.

Acronyms

394	ARI Adjusted Rand Index 5, 7–9, 12
395	CAR Clustering Agent Representation 1–5, 7, 8, 10, 12–14
396	CS Concept Structure 1–3, 13
397	LR Latent Representation 1–3, 7
398	LS Latent Space 1–5, 7–14
399	MSE Mean Squared Error 4, 6
400	NN Neural Network 1–3, 13
401	PPO Proximal Policy Optimization 11
402	RI Rand Index 7
403	RL Reinforcement Learning 1, 3, 10, 14
404	SAE Sparse Autoencoder 13

Bibliography

- [Amezquita et al., 2020] Amezquita, R. A., Lun, A. T., Becht, E., Carey, V. J., Carpp, L. N., Geistlinger, L., Marini, F., Rue-Albrecht, K., Risso, D., Soneson, C., et al. (2020). Orchestrating single-cell analysis with Bioconductor, volume 17. Nature Publishing Group US New York. **Book**.
- [AtariAge, 2024] AtariAge (2024). Mario bros. atari 2600. **Web**.
- [Bellemare et al., 2013] Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. Journal of Artificial Intelligence Research, 47:253–279. **Article**.
- [Bereska and Gavves, 2024] Bereska, L. and Gavves, E. (2024). Mechanistic interpretability for ai safety – a review. **Article**.
- [Bricken et al., 2023] Bricken, T., Templeton, A., Batson, J., Chen, B., Jermyn, A., Conerly, T., et al. (2023). Towards monosemanticity: Decomposing language models with dictionary learning. Transformer Circuits Thread. **Web**.
- [Elhage et al., 2022] Elhage, N., Hume, T., Olsson, C., Schiefer, N., Henighan, T., Kravec, S., et al. (2022). Toy models of superposition. Transformer Circuits Thread. **Web**.
- [Guidotti et al., 2018] Guidotti, R., Monreale, A., Ruggieri, S., Turini, F., Pedreschi, D., and Giannotti, F. (2018). A survey of methods for explaining black box models. **Article**.
- [Hubert and Arabie, 1985] Hubert, L. and Arabie, P. (1985). Comparing partitions. Journal of classification, 2(1):193–218. **Article**.
- [Knuth and Moore, 1975] Knuth, D. E. and Moore, R. W. (1975). An analysis of alpha-beta pruning. Artificial intelligence, 6(4):293–326. **Article**.
- [Machado et al., 2017] Machado, M. C., Bellemare, M. G., Talvitie, E., Veness, J., Hausknecht, M., and Bowling, M. (2017). Revisiting the arcade learning environment: Evaluation protocols and open problems for general agents. **Article**. **Web**.
- [MacQueen, 1967] MacQueen, J. (1967). Multivariate observations. In Proceedings of the 5th Berkeley Symposium on Mathematical Statistics and Probability, volume 1, pages 281–297. **Article**.

- [Mnih et al., 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. (2013). Playing atari with deep reinforcement learning.
- [Peyré and Cuturi, 2020] Peyré, G. and Cuturi, M. (2020). Computational optimal transport. **Article**.
- [Rand, 1971] Rand, W. M. (1971). Objective criteria for the evaluation of clustering methods. Journal of the American Statistical association, 66(336):846–850.
- [Ross and Bagnell, 2010] Ross, S. and Bagnell, D. (2010). Efficient reductions for imitation learning. In Teh, Y. W. and Titterton, M., editors, Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics, volume 9 of Proceedings of Machine Learning Research, pages 661–668, Chia Laguna Resort, Sardinia, Italy. PMLR. **Article**.
- [Ross et al., 2011] Ross, S., Gordon, G., and Bagnell, D. (2011). A reduction of imitation learning and structured prediction to no-regret online learning. In Gordon, G., Dunson, D., and Dudík, M., editors, Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, volume 15 of Proceedings of Machine Learning Research, pages 627–635, Fort Lauderdale, FL, USA. PMLR. **Article**.
- [Schulman et al., 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., and Klimov, O. (2017). Proximal policy optimization algorithms.
- [Shlens, 2014] Shlens, J. (2014). Notes on kullback-leibler divergence and likelihood. **Article**.
- [Sutton and Barto, 2018] Sutton, R. S. and Barto, A. G. (2018). Reinforcement Learning: An Introduction. MIT Press. **Book**.
- [Templeton et al., 2024] Templeton, A., Conerly, T., Marcus, J., Lindsey, J., Bricken, T., Chen, B., et al. (2024). Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet. Transformer Circuits Thread. **Web**.
- [Zahavy et al., 2017] Zahavy, T., Zrihem, N. B., and Mannor, S. (2017). Graying the black box: Understanding dqns. **Article**.